

Traditional Language Models



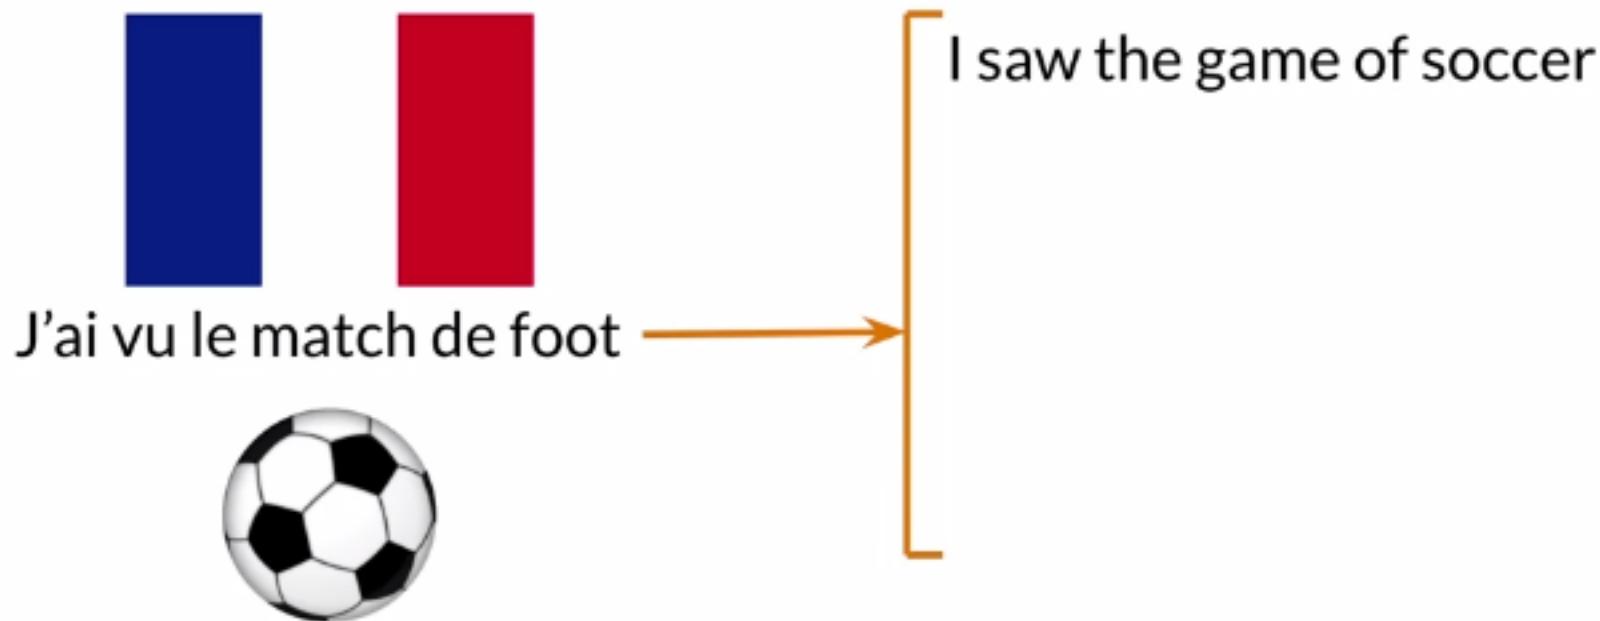
J'ai vu le match de foot



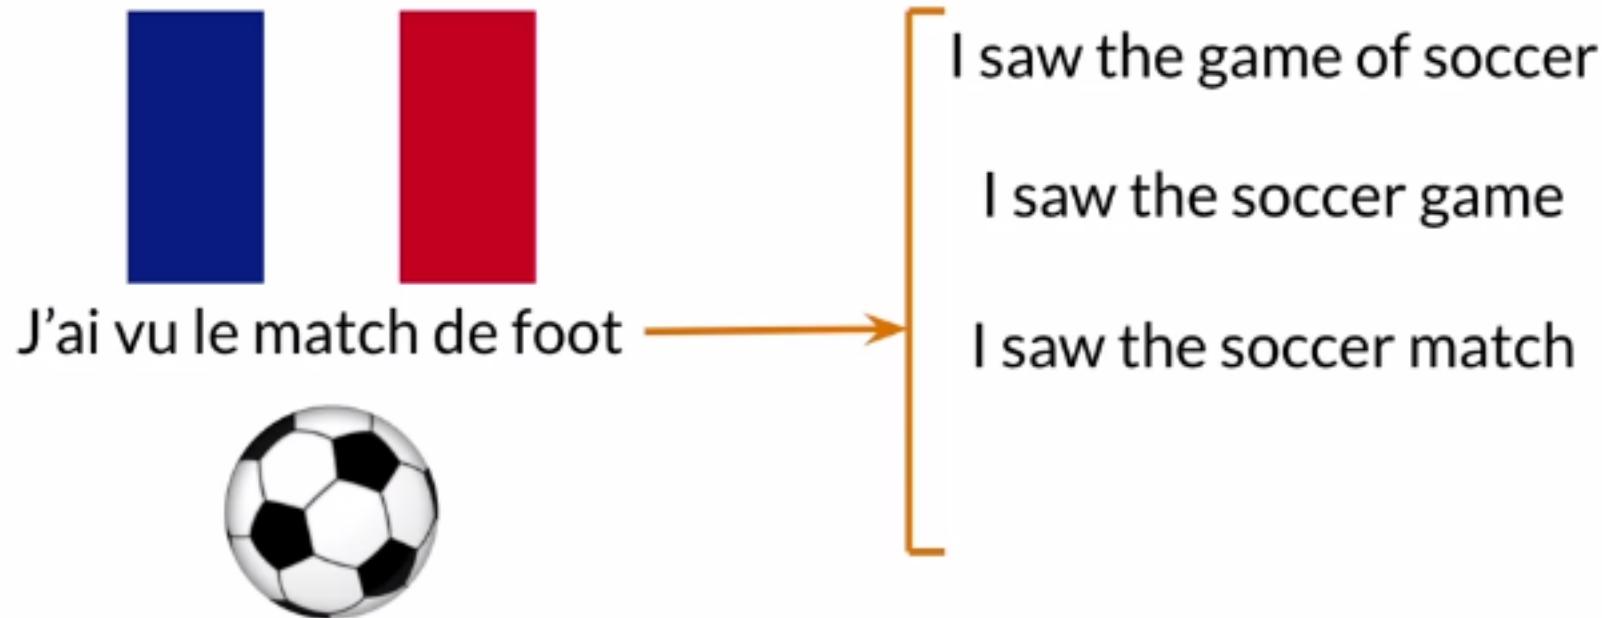
Traditional Language Models



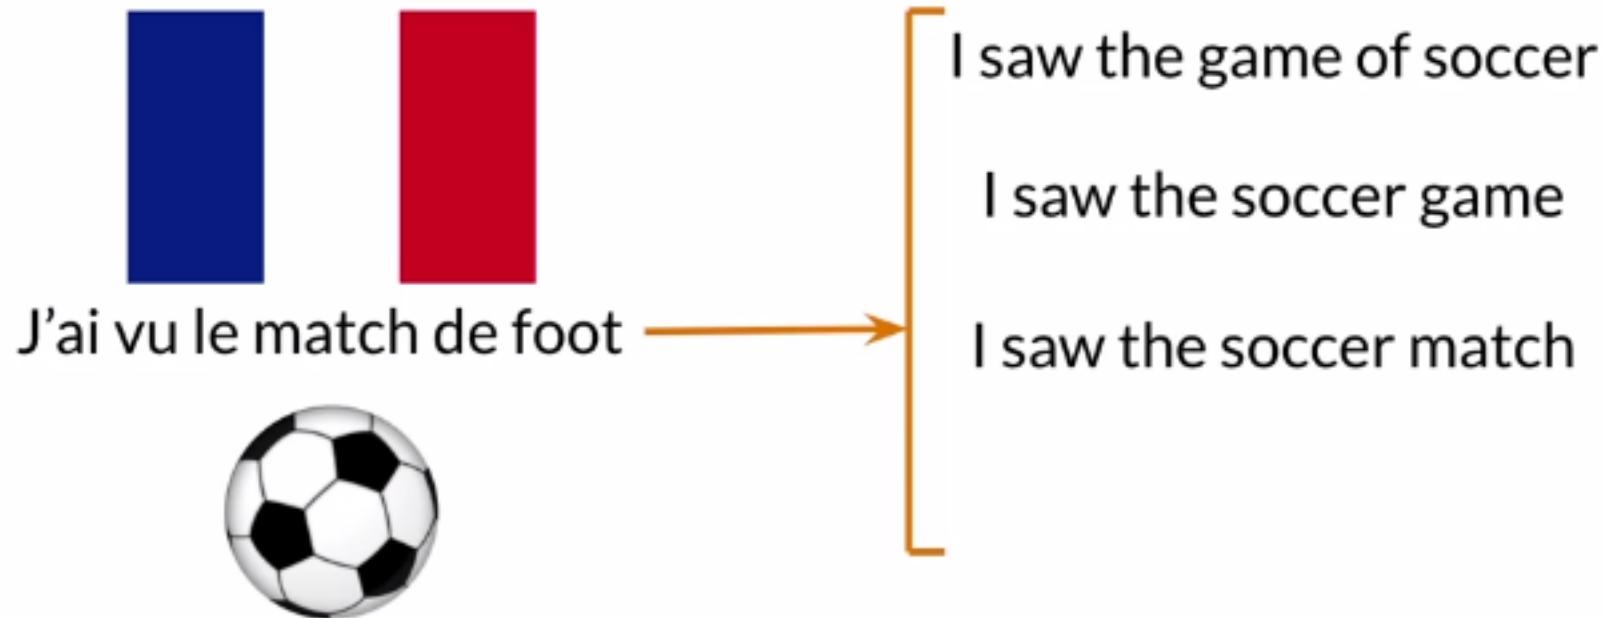
Traditional Language Models



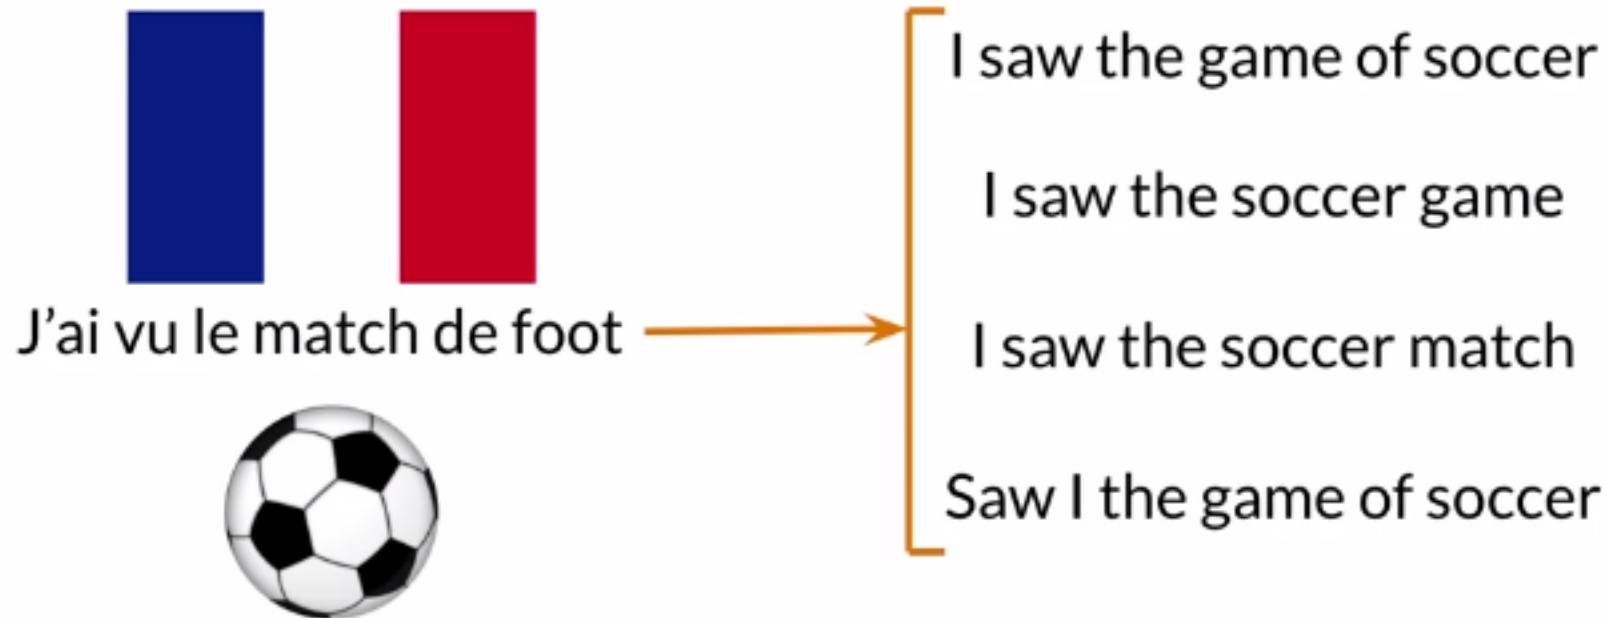
Traditional Language Models



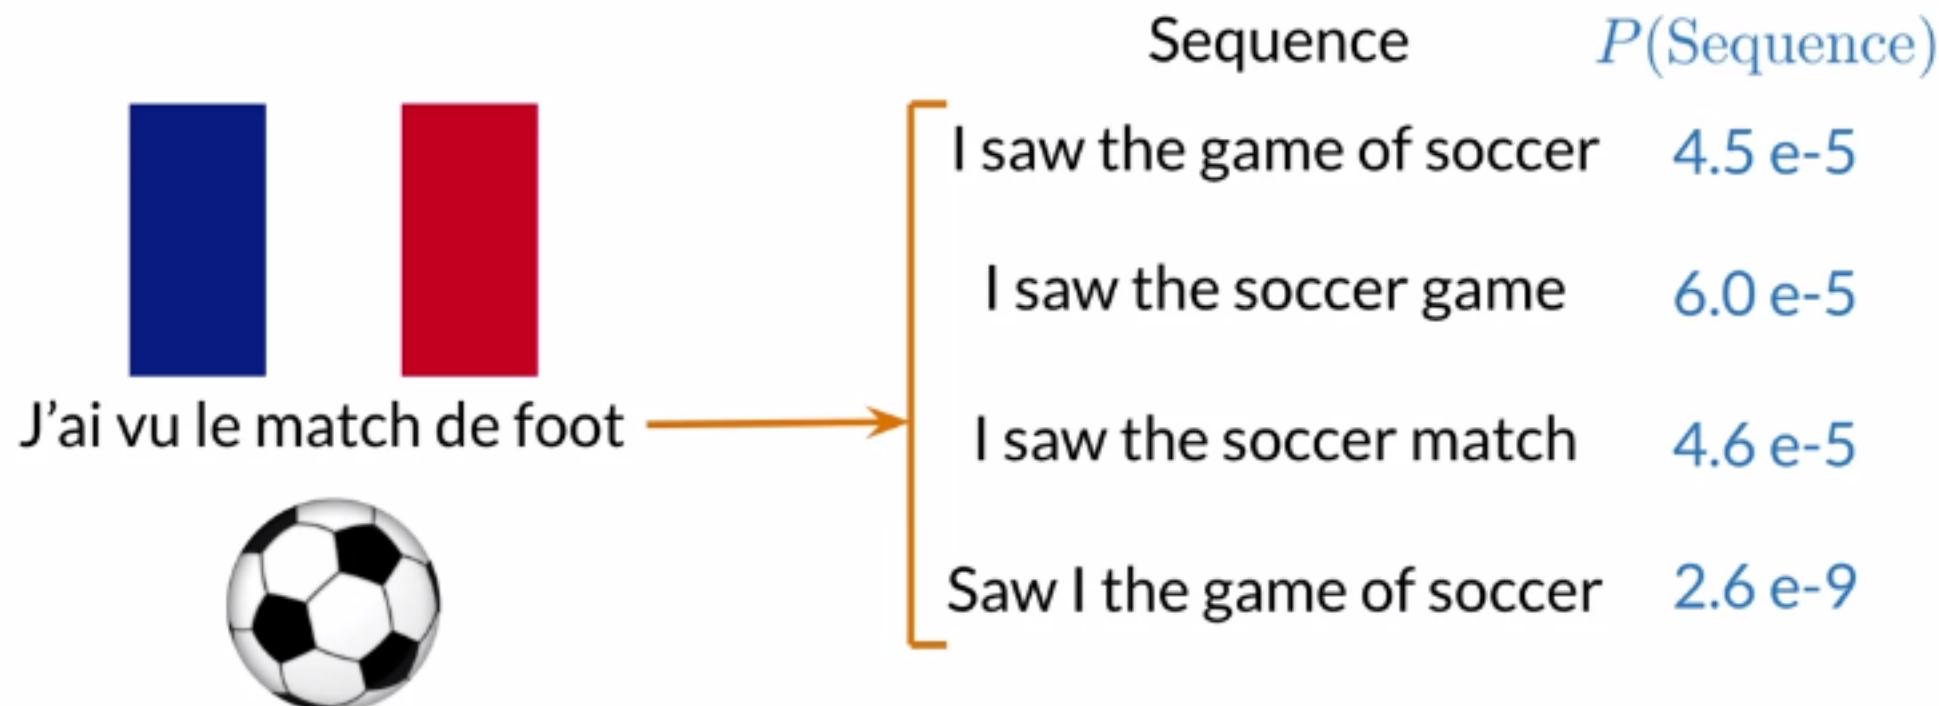
Traditional Language Models



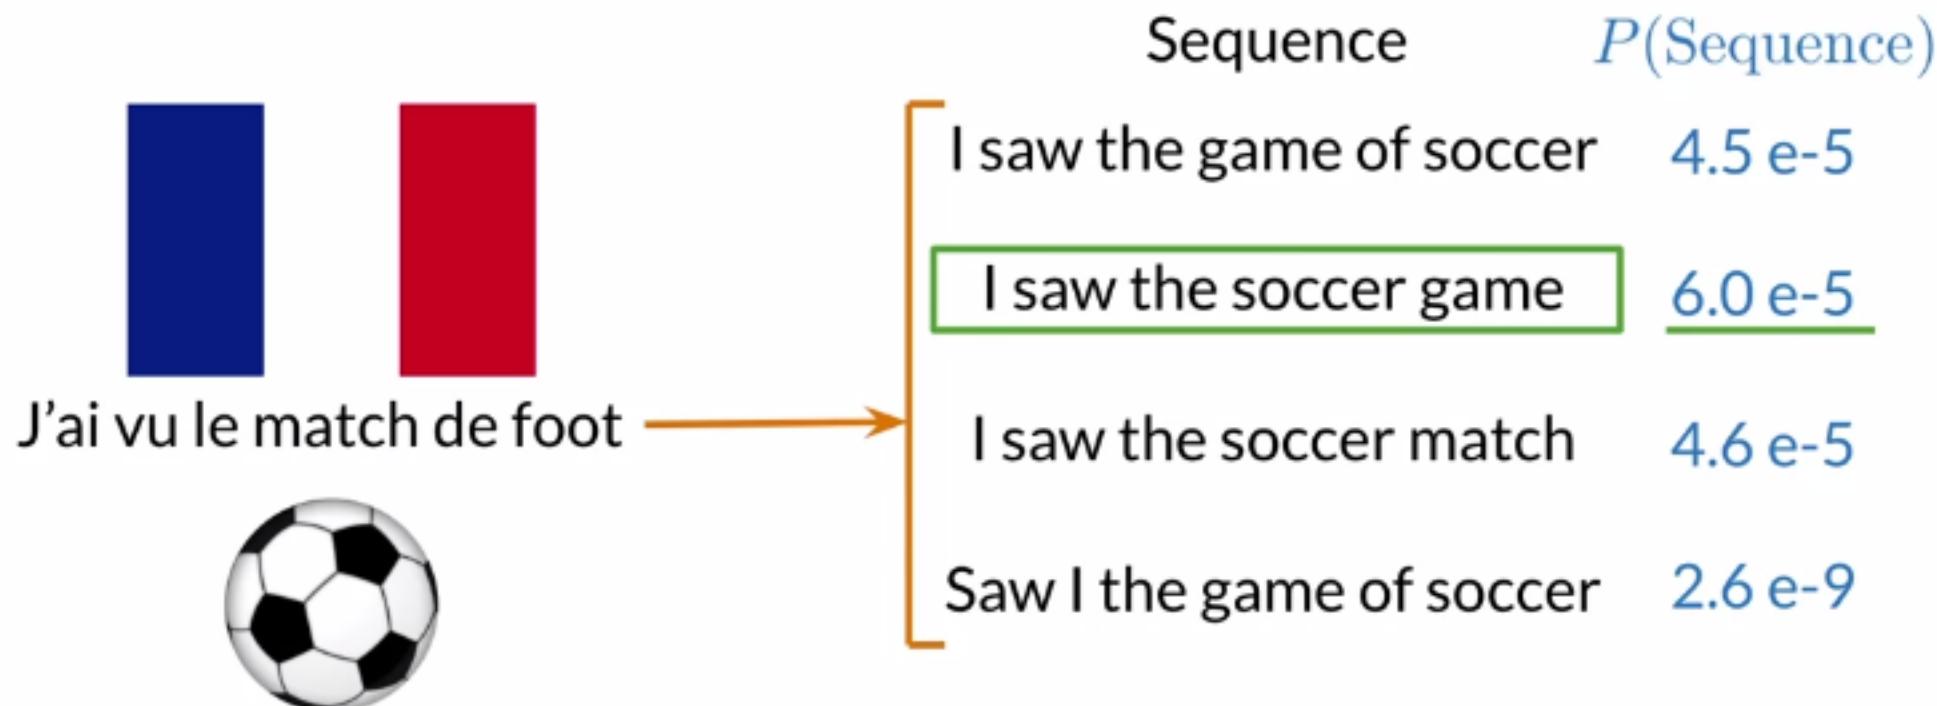
Traditional Language Models



Traditional Language Models



Traditional Language Models



N-grams

N-grams

$$P(w_2|w_1) = \frac{\text{count}(w_1, w_2)}{\text{count}(w_1)}$$

N-grams

$$P(w_2|w_1) = \frac{\text{count}(w_1, w_2)}{\text{count}(w_1)} \longrightarrow \text{Bigrams}$$

N-grams

$$P(w_2|w_1) = \frac{\text{count}(w_1, w_2)}{\text{count}(w_1)} \longrightarrow \text{Bigrams}$$

$$P(w_3|w_1, w_2) = \frac{\text{count}(w_1, w_2, w_3)}{\text{count}(w_1, w_2)} \longrightarrow \text{Trigrams}$$

N-grams

$$P(w_2|w_1) = \frac{\text{count}(w_1, w_2)}{\text{count}(w_1)} \longrightarrow \text{Bigrams}$$

$$P(w_3|w_1, w_2) = \frac{\text{count}(w_1, w_2, w_3)}{\text{count}(w_1, w_2)} \longrightarrow \text{Trigrams}$$

$$P(w_1, w_2, w_3) = P(w_1) \times P(w_2|w_1) \times P(w_3|w_2)$$

N-grams

$$P(w_2|w_1) = \frac{\text{count}(w_1, w_2)}{\text{count}(w_1)} \longrightarrow \text{Bigrams}$$

$$P(w_3|w_1, w_2) = \frac{\text{count}(w_1, w_2, w_3)}{\text{count}(w_1, w_2)} \longrightarrow \text{Trigrams}$$

$$P(w_1, w_2, w_3) = P(w_1) \times P(w_2|w_1) \times P(w_3|w_2)$$

- Large N-grams to capture dependencies between distant words

N-grams

$$P(w_2|w_1) = \frac{\text{count}(w_1, w_2)}{\text{count}(w_1)} \longrightarrow \text{Bigrams}$$

$$P(w_3|w_1, w_2) = \frac{\text{count}(w_1, w_2, w_3)}{\text{count}(w_1, w_2)} \longrightarrow \text{Trigrams}$$

$$P(w_1, w_2, w_3) = P(w_1) \times P(w_2|w_1) \times P(w_3|w_2)$$

- Large N-grams to capture dependencies between distant words
- Need a lot of space and RAM

Advantages of RNNs

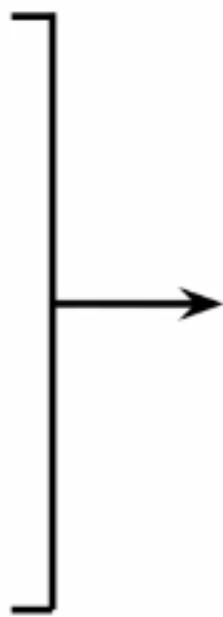
Advantages of RNNs

Nour was supposed to study with me. I called her but she did not _____

Advantages of RNNs

Nour was supposed to study with me. I called her but she did not _____

want
respond
choose
want
have
ask
attempt
answer
know

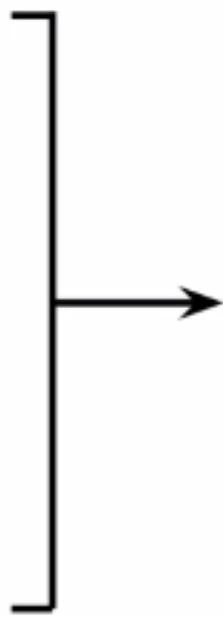


Similar probabilities with trigram

Advantages of RNNs

Nour was supposed to study with me. I called her but she did not have

want
respond
choose
want
have
ask
attempt
answer
know



Similar probabilities with trigram

Advantages of RNNs

Nour was supposed to study with me. I called her but she did not

want
respond
choose
want
have
ask
attempt
answer
know



RNNs look at every previous word

Similar probabilities with trigram

Advantages of RNNs

Nour was supposed to study with me. I called her but she did not

want
respond
choose
want
have
ask
attempt
answer
know



RNNs look at every previous word

Similar probabilities with trigram

Advantages of RNNs

Nour was supposed to study with me. I called her but she did not answer

want
respond
choose
want
have
ask
attempt
answer
know



RNNs look at every previous word

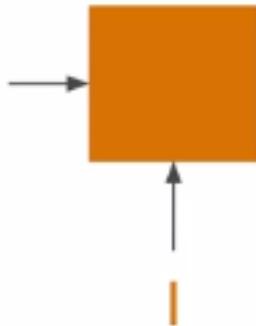
Similar probabilities with trigram

RNNs Basic Structure

I called her but she did not _____

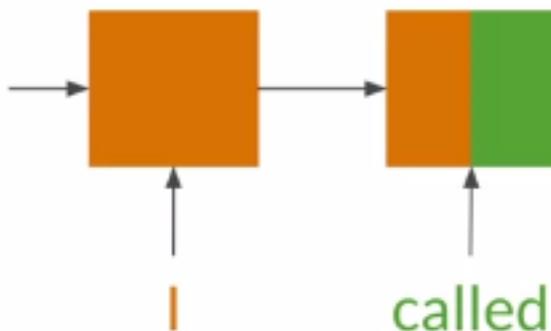
RNNs Basic Structure

I called her but she did not _____



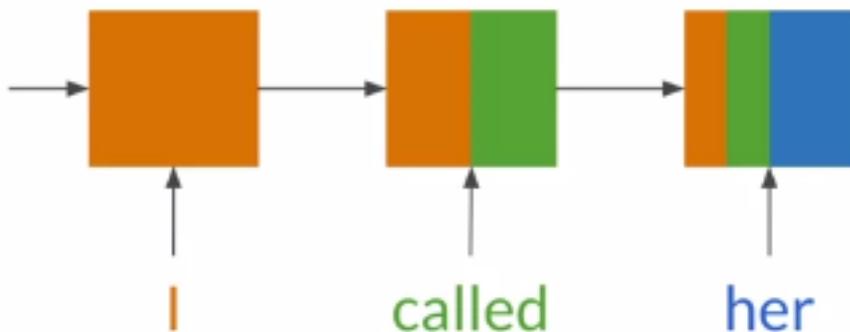
RNNs Basic Structure

I called her but she did not _____



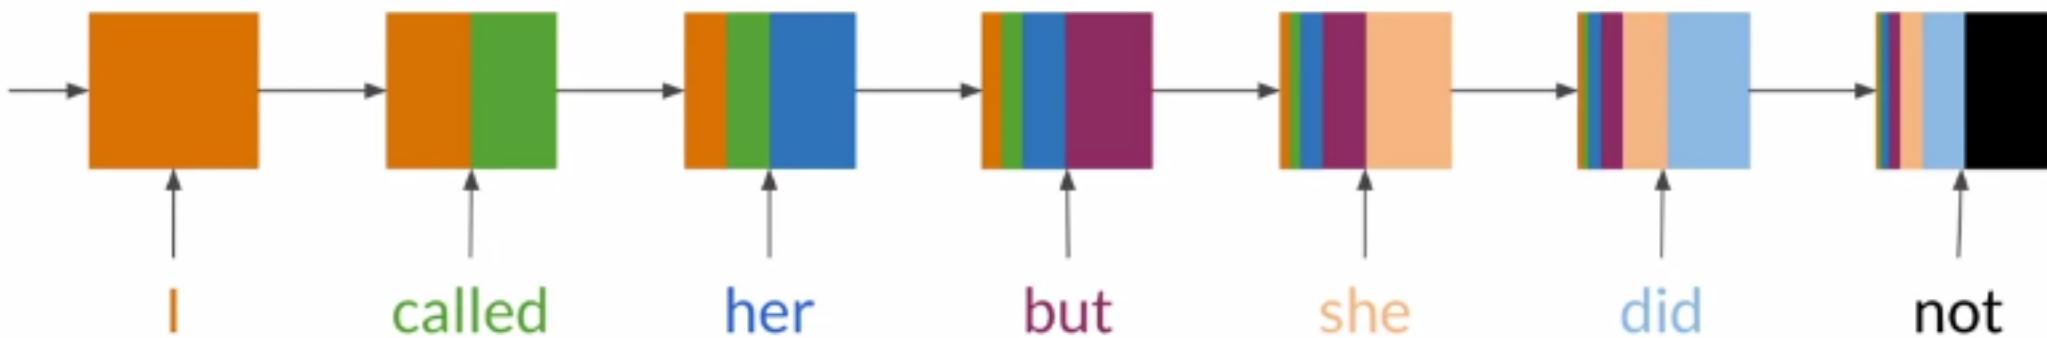
RNNs Basic Structure

I called her but she did not _____



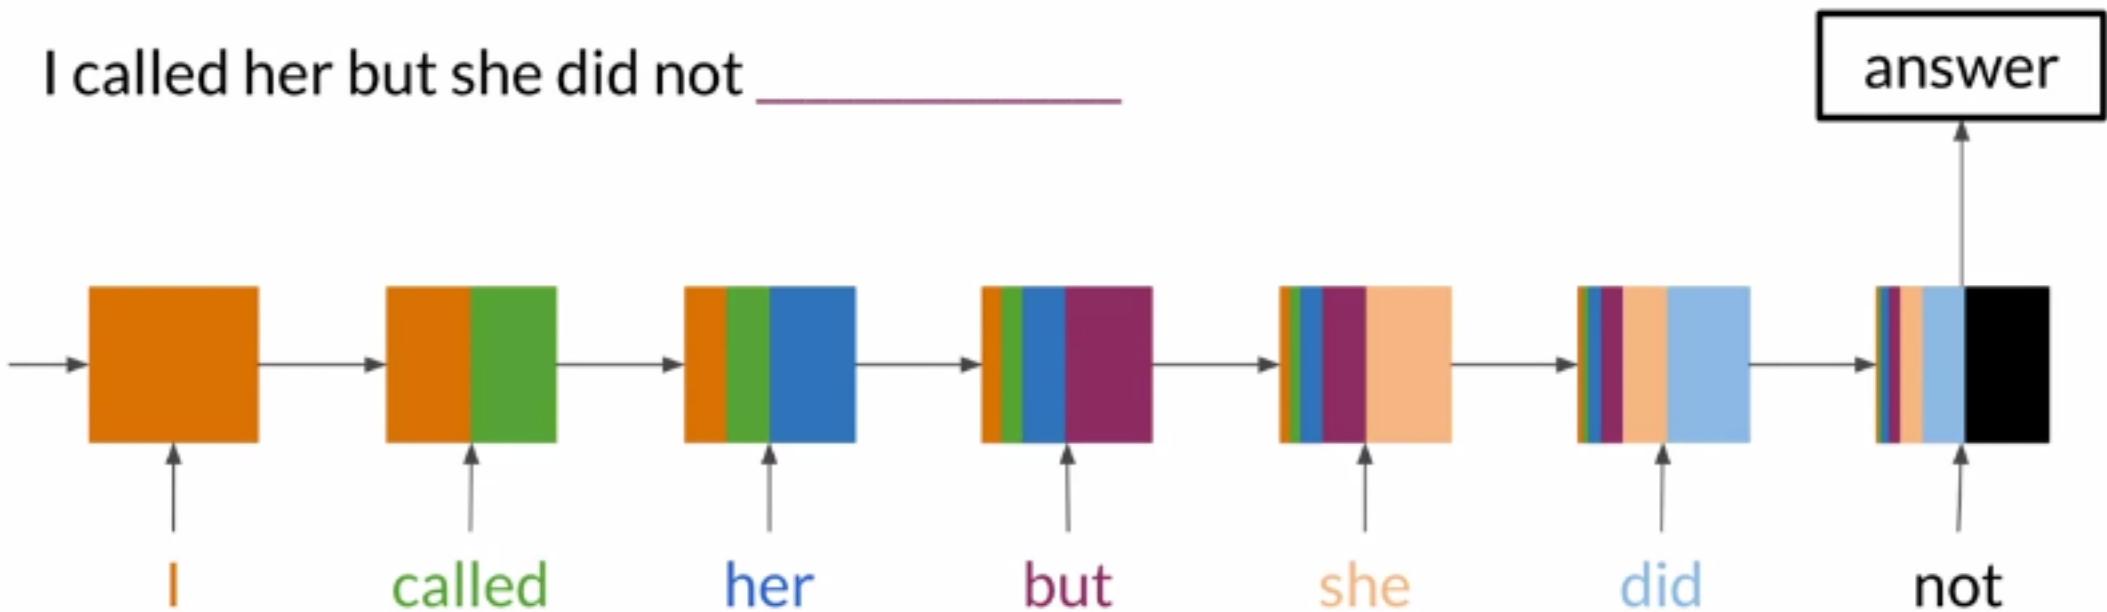
RNNs Basic Structure

I called her but she did not _____



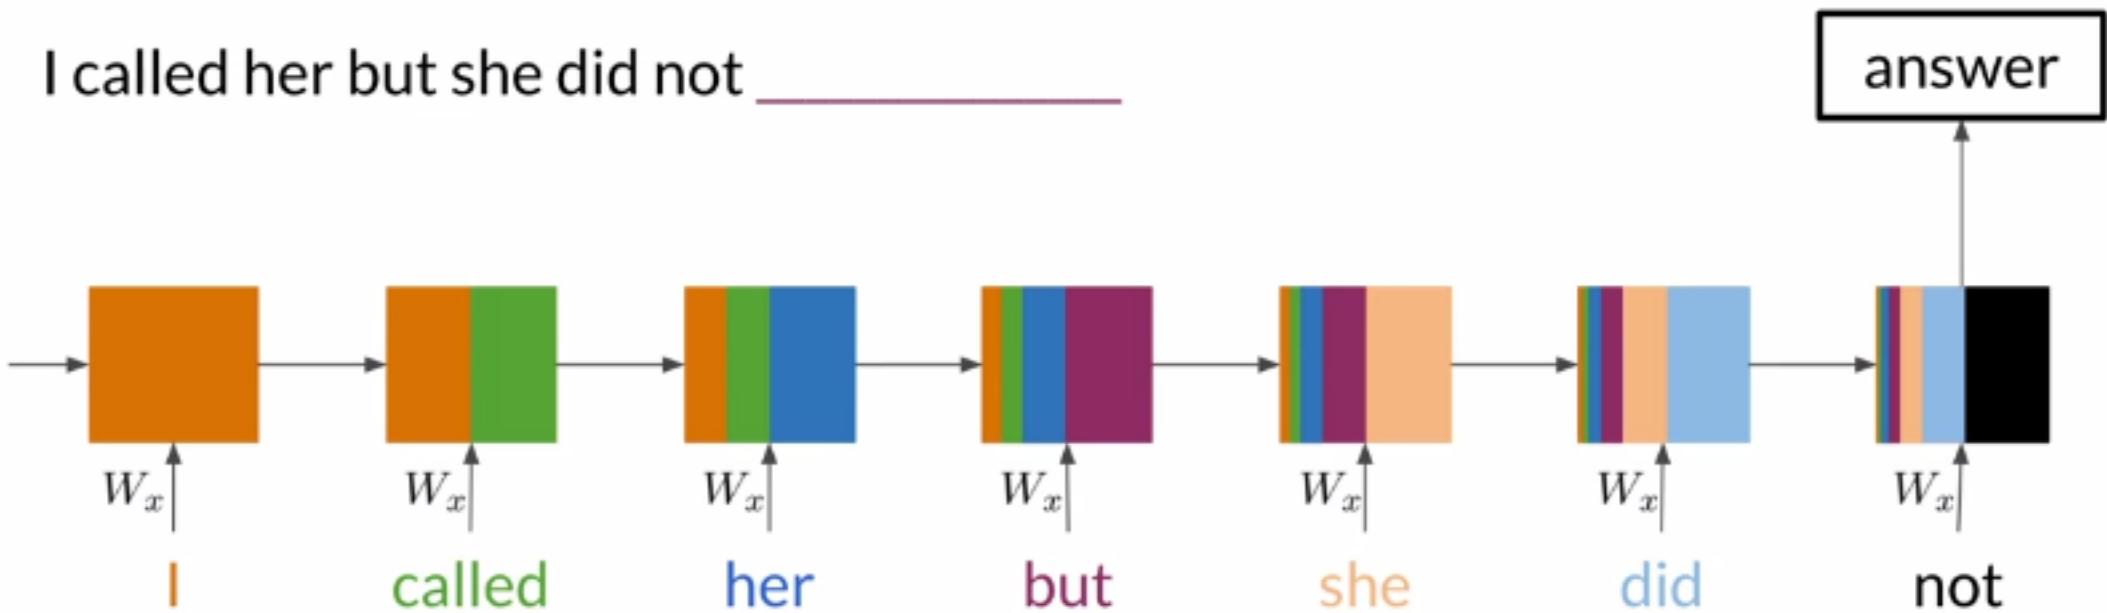
RNNs Basic Structure

I called her but she did not _____



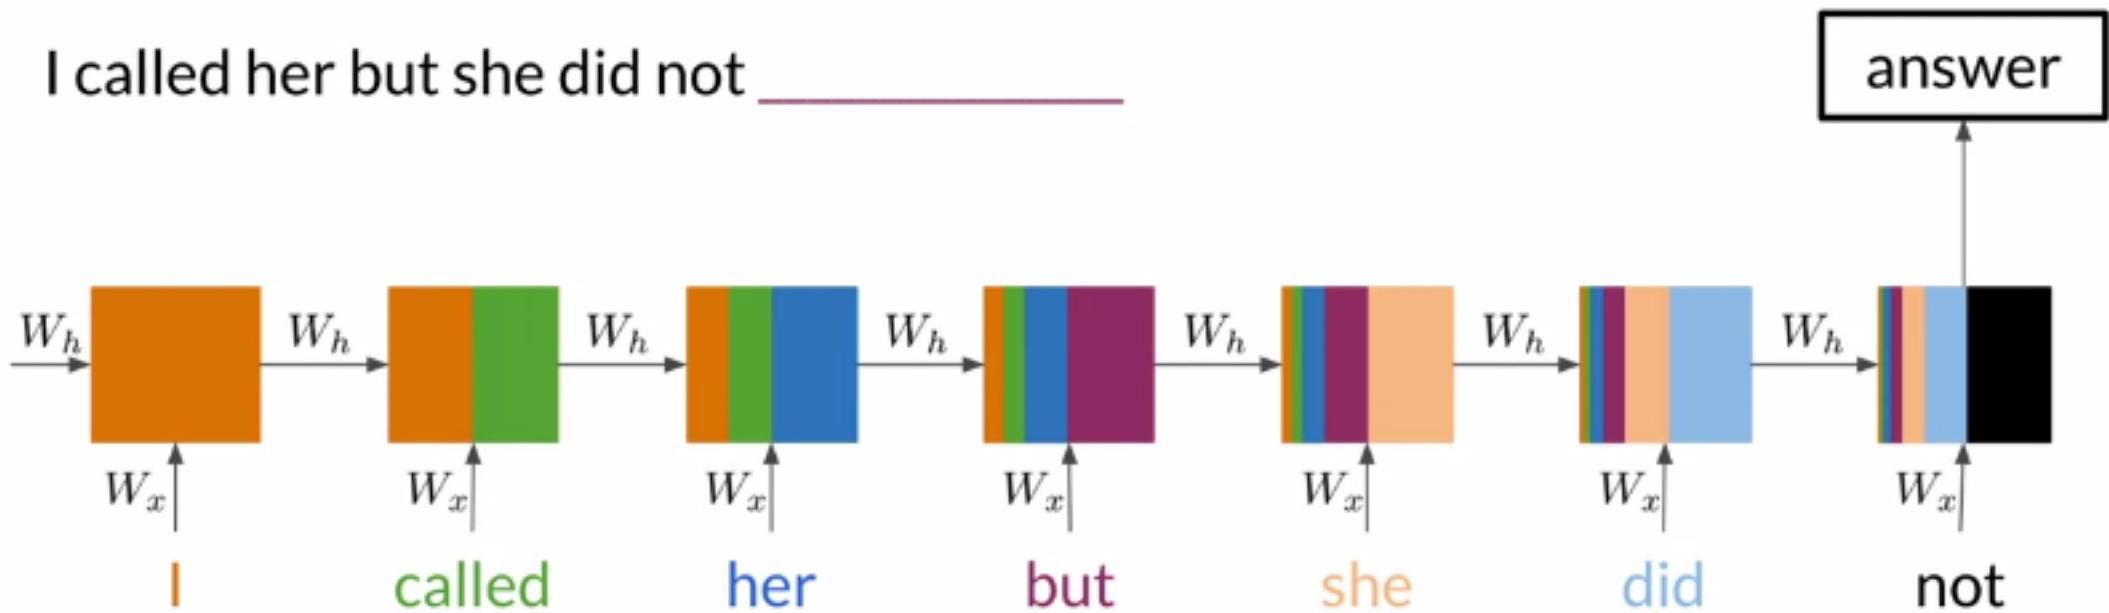
RNNs Basic Structure

I called her but she did not _____



RNNs Basic Structure

I called her but she did not _____



RNNs Basic Structure

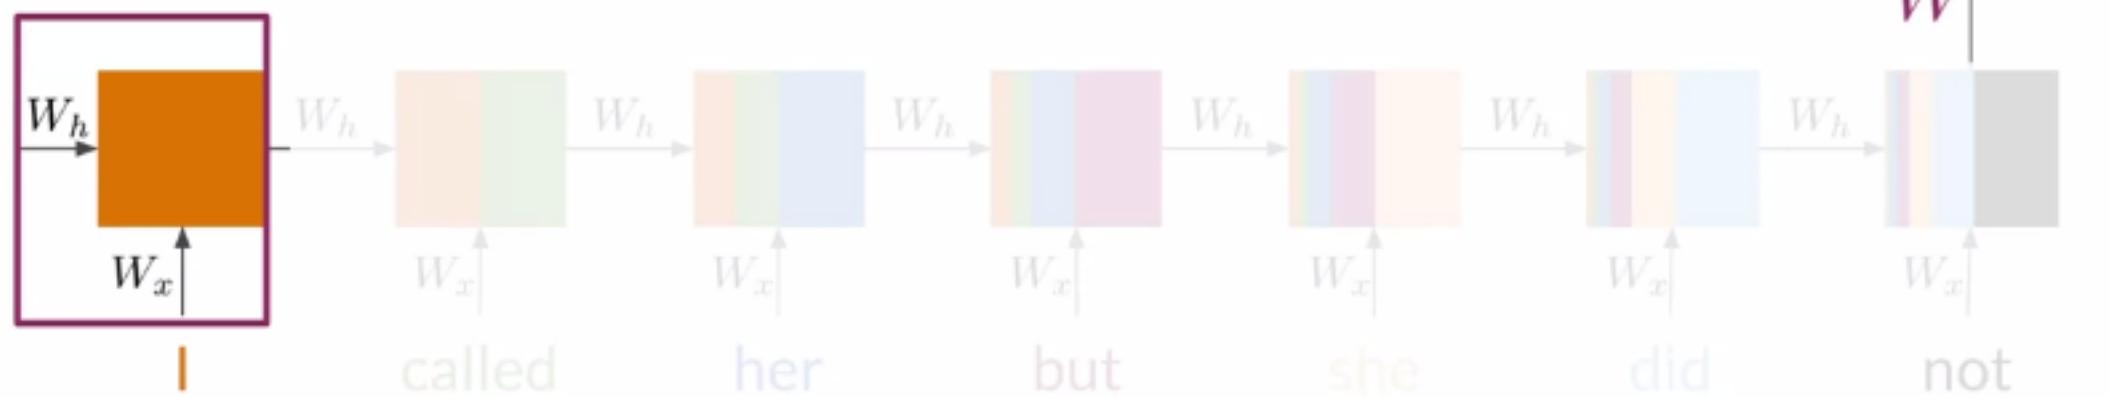
I called her but she did not _____



Learnable parameters

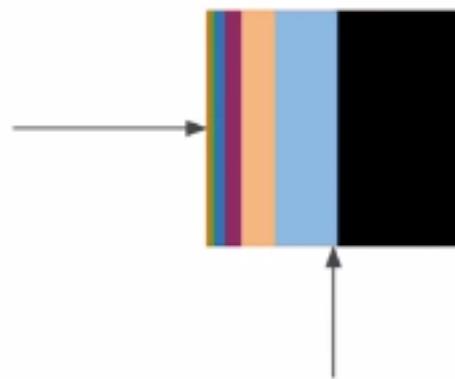
RNNs Basic Structure

I called her but she did not _____



Summary

- RNNs model relationships among distant words
- In RNNs a lot of computations share parameters



One to One

One to One



Real Valladolid	0	1	Real Madrid
Real Zaragoza	0	4	Real Madrid
Atletico Madrid	0	1	Real Madrid

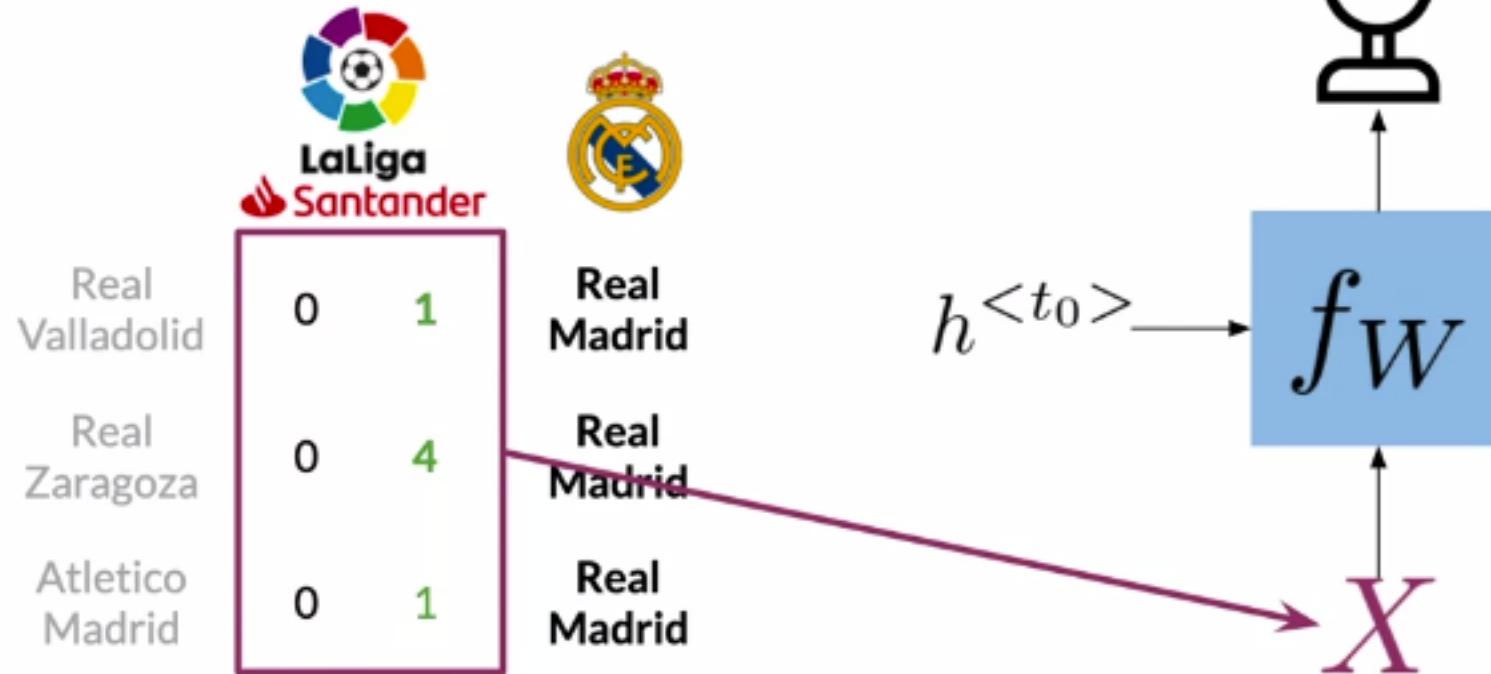
One to One



One to One



One to One



One to Many

One to Many

Caption
generation

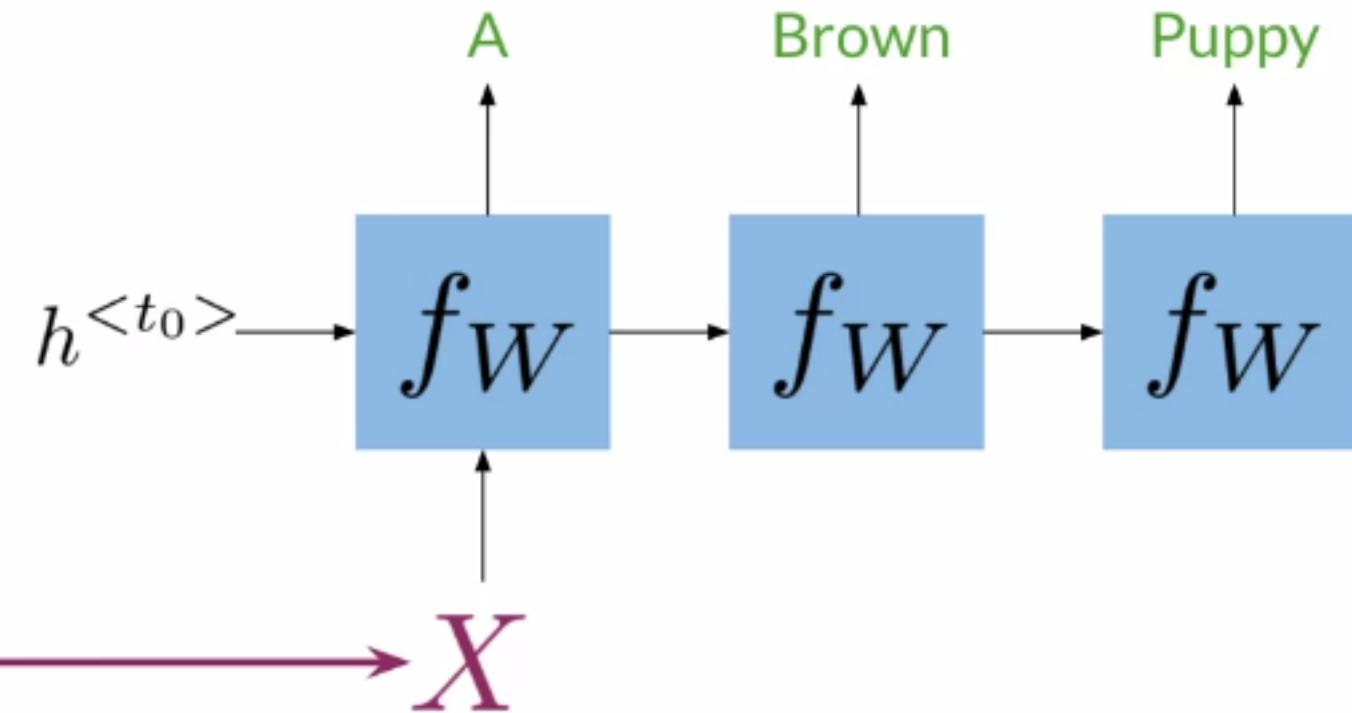
A Brown Puppy



$\rightarrow X$

One to Many

Caption
generation



Many to One

Many to One

Sentiment
analysis

Tweet:

I am very happy !

Many to One

Sentiment
analysis

Tweet:

I am very happy !

Many to One

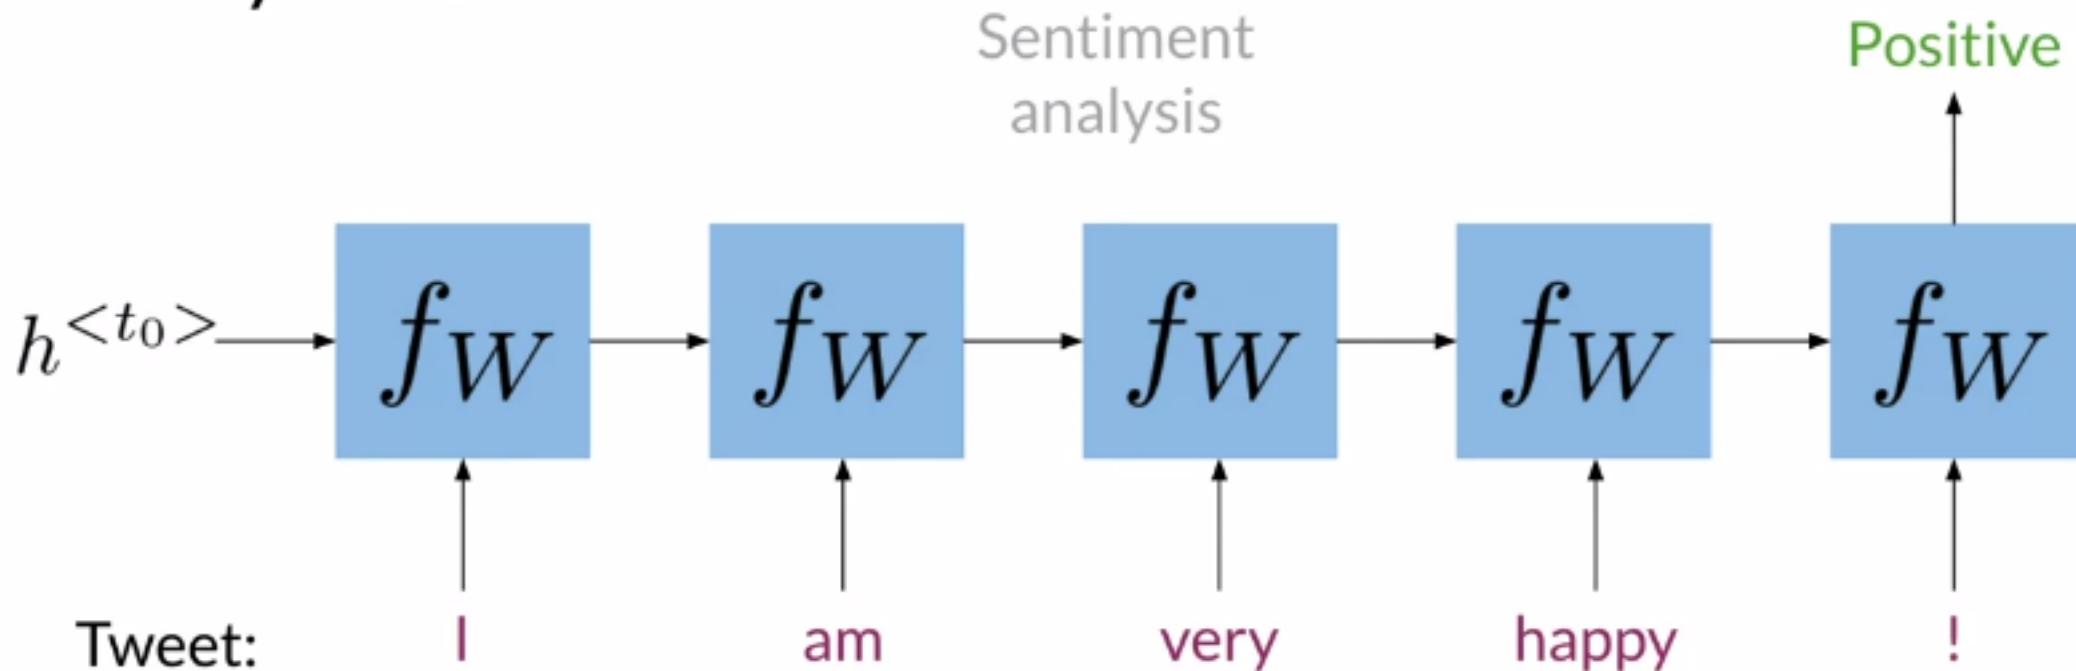
Sentiment
analysis

Positive

Tweet:

I am very happy !

Many to One



Many to Many

Many to Many



J'
ai
faim

Machine
Translation

Many to Many



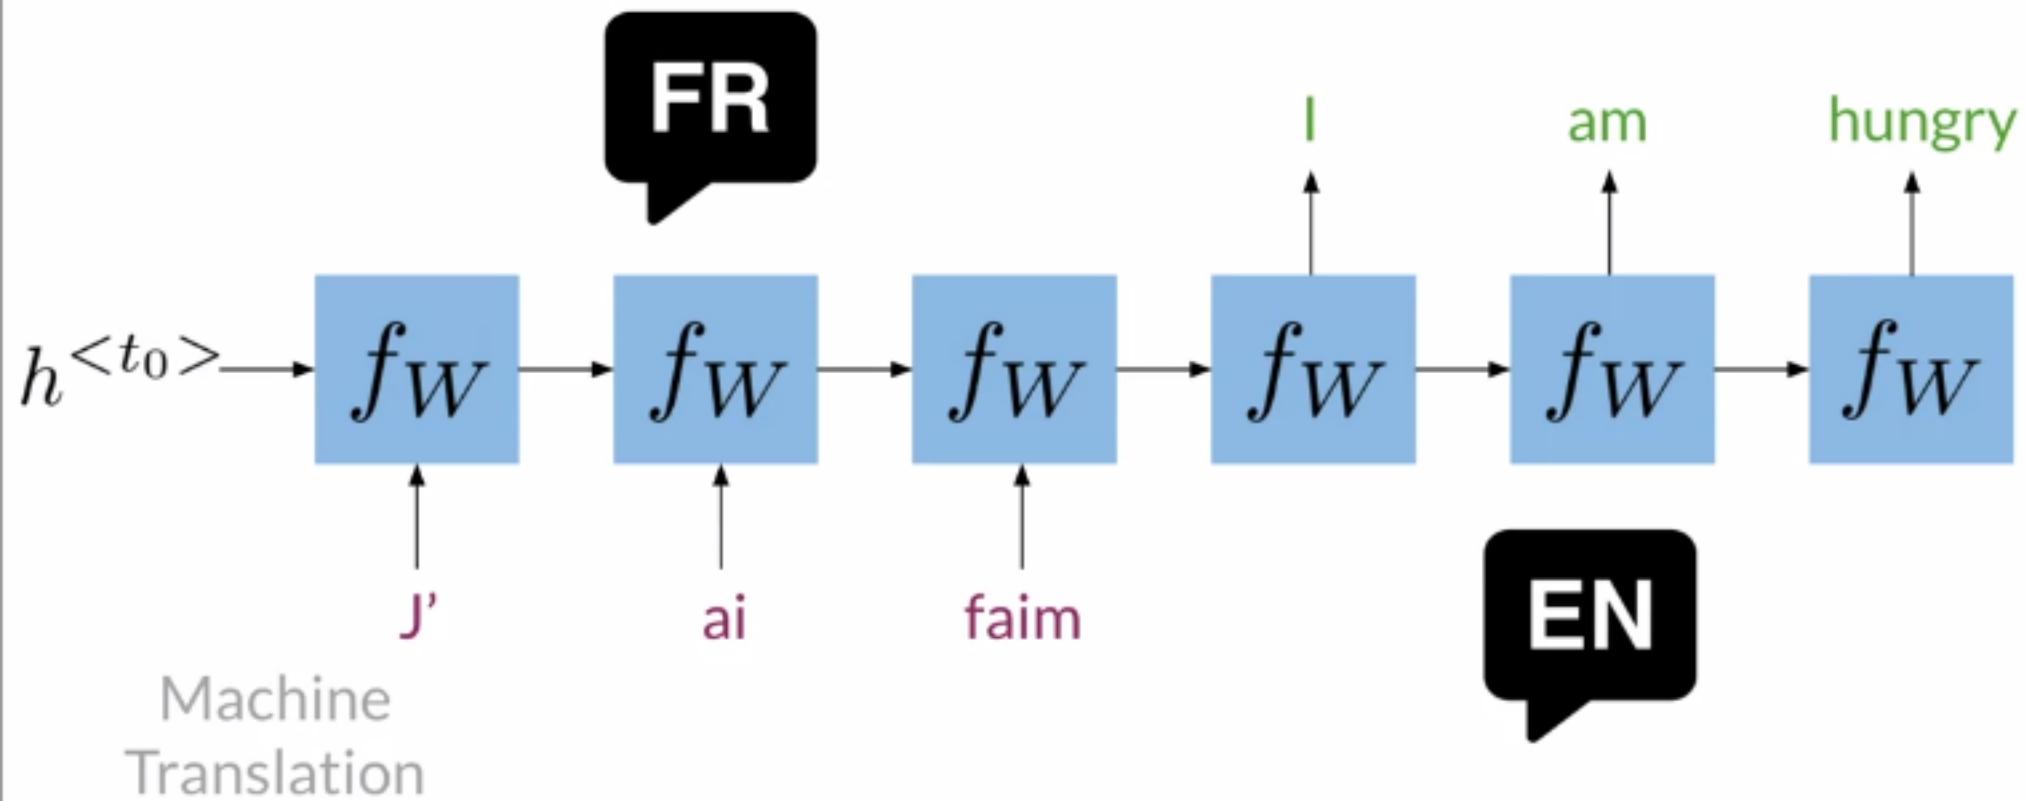
I am hungry

J' ai faim

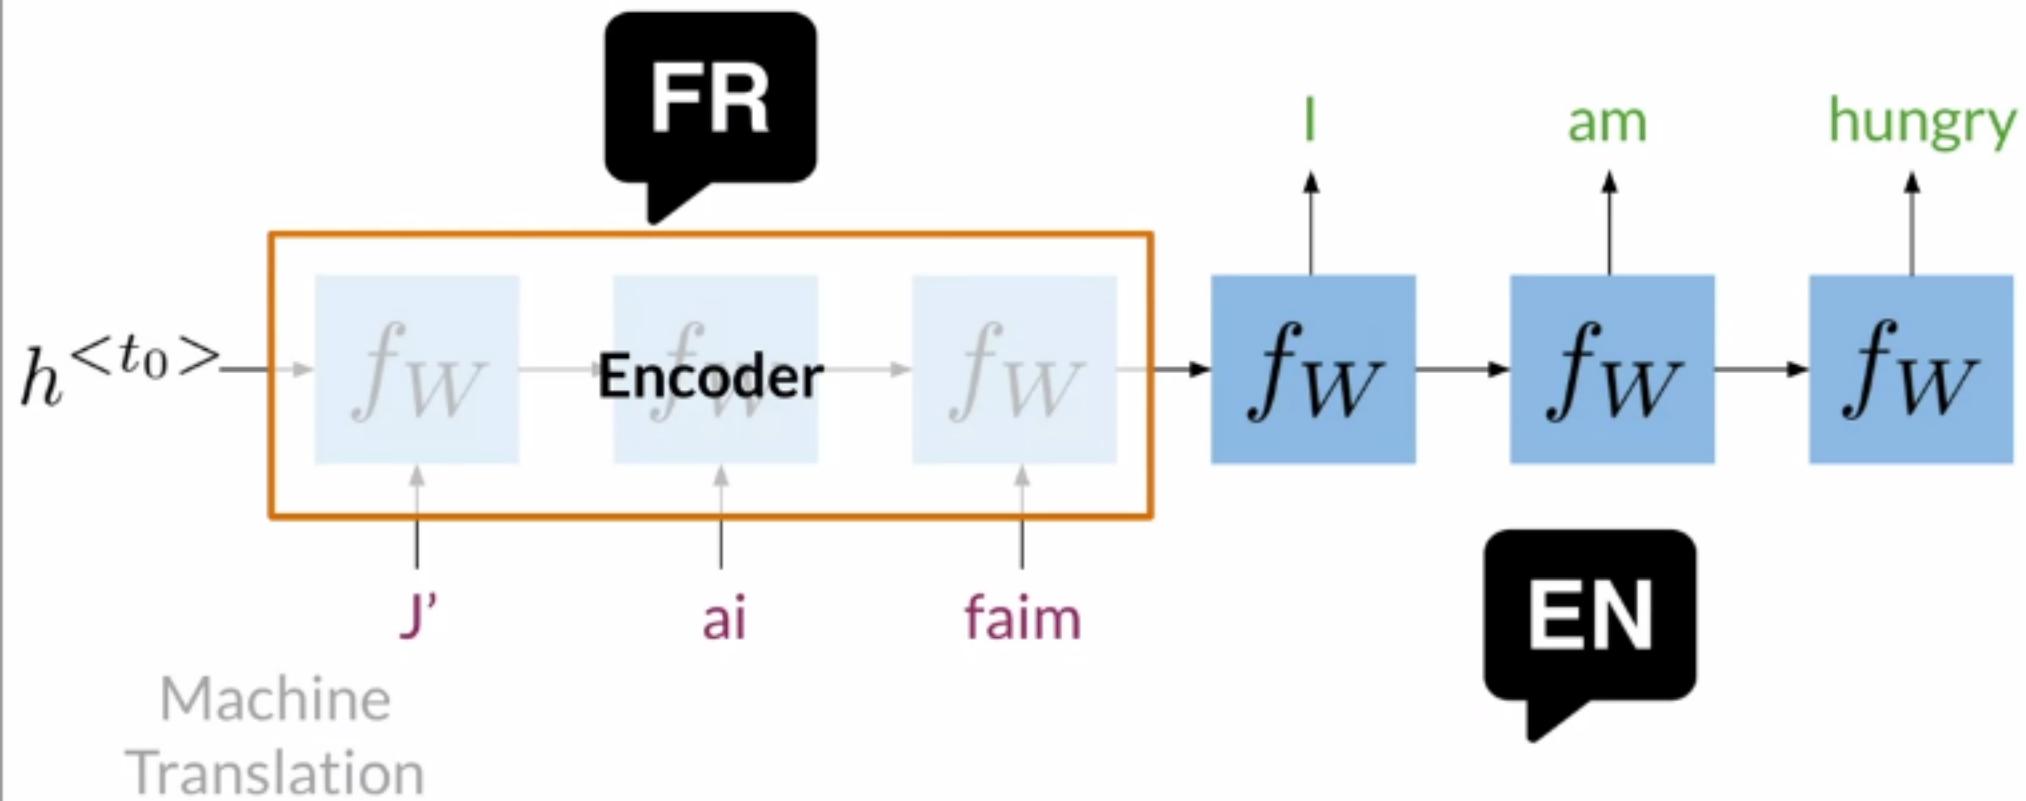
Machine
Translation



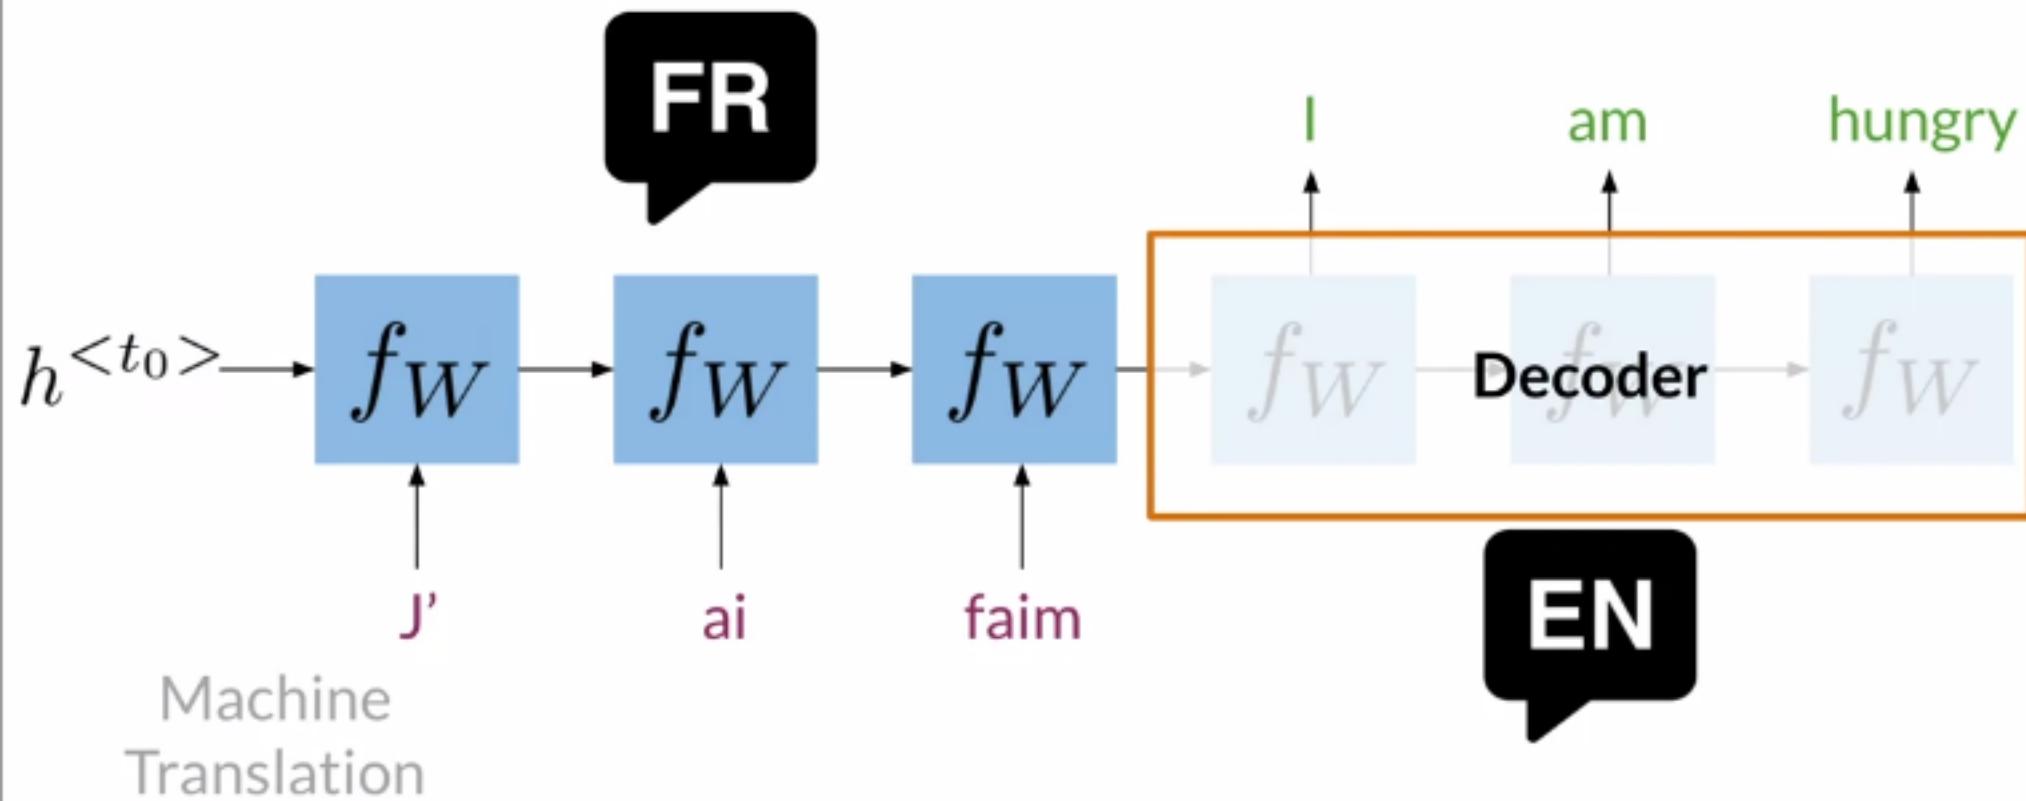
Many to Many



Many to Many



Many to Many



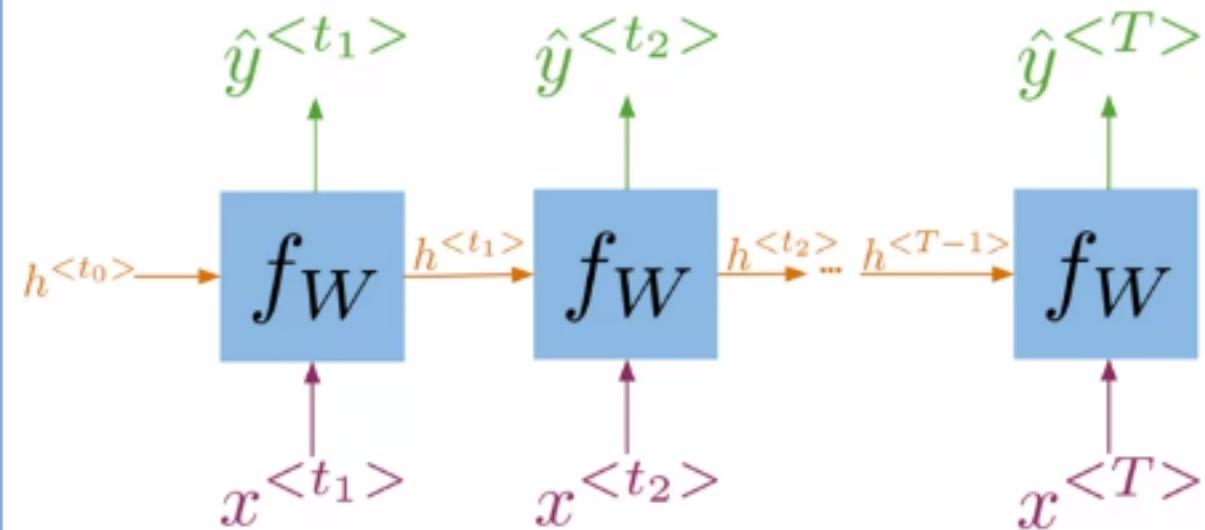
Summary

- RNNs can be implemented for a variety of NLP tasks
- Applications include Machine translation and caption generation

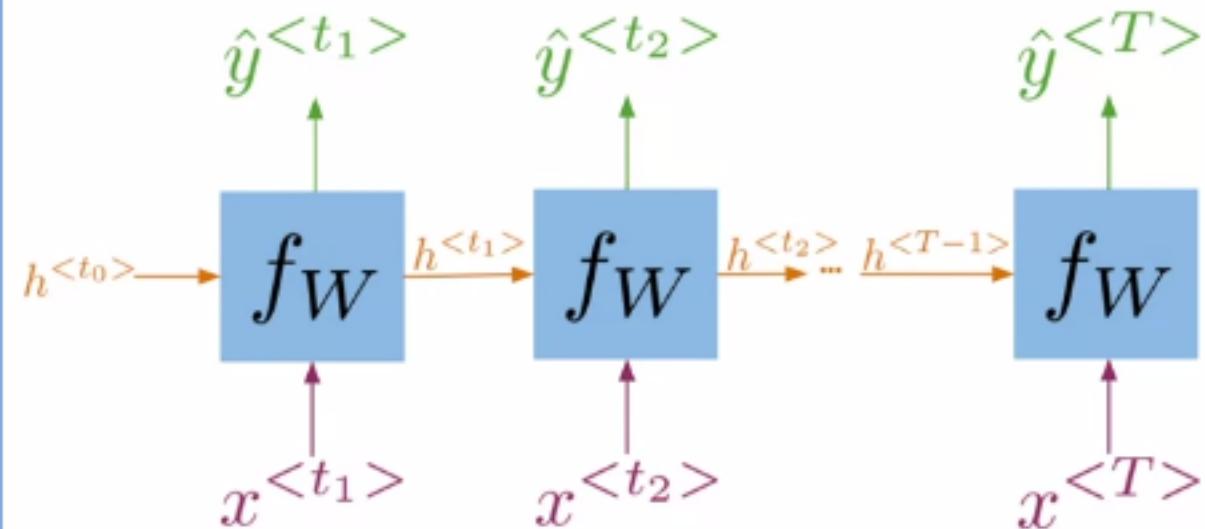


A Vanilla RNN

A Vanilla RNN

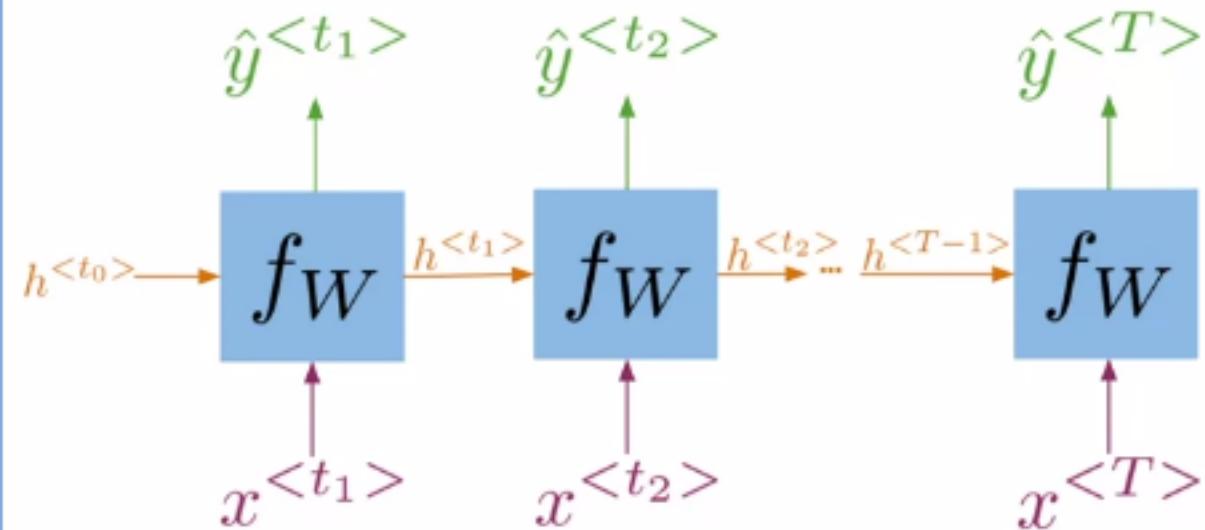


A Vanilla RNN



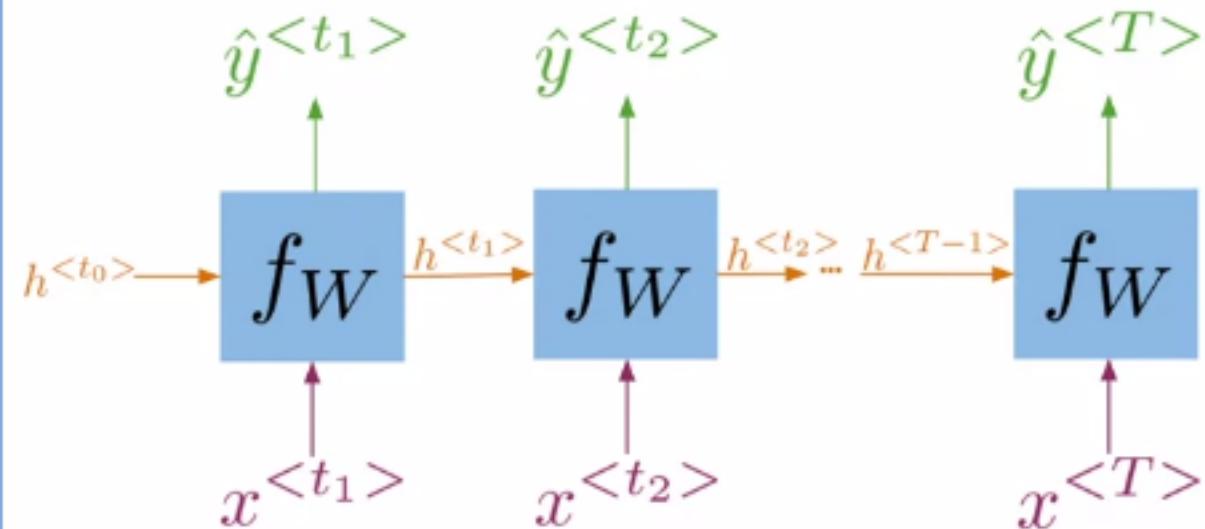
$$h^{<t>} = g(W_h[h^{<t-1>}, x^{<t>}] + b_h)$$

A Vanilla RNN



$$h^{<t>} = g(W_h[h^{<t-1>}, x^{<t>}] + b_h)$$

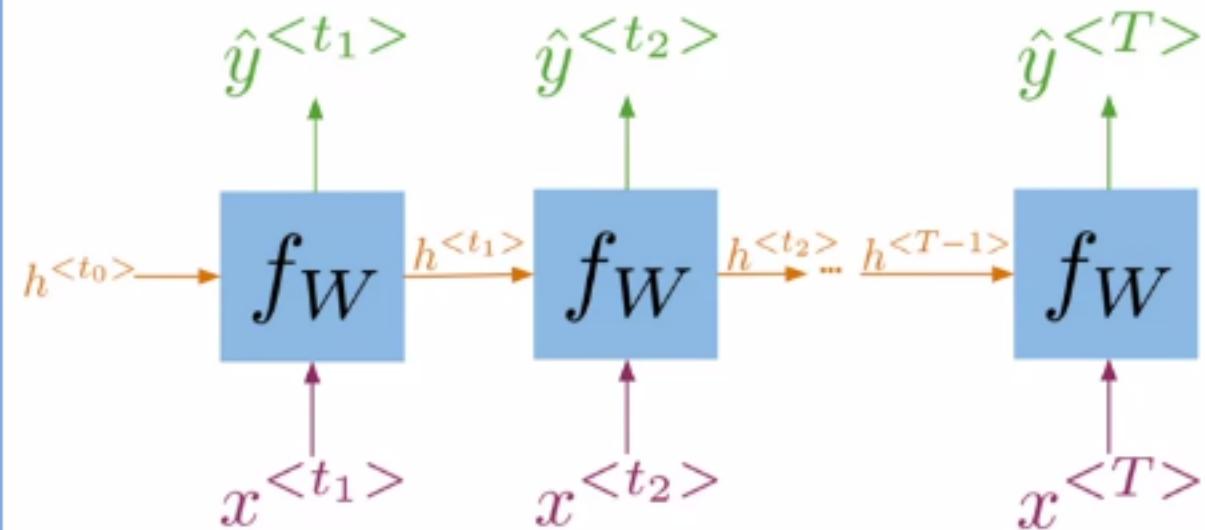
A Vanilla RNN



$$h^{<t>} = g(W_h[h^{<t-1>}, x^{<t>}] + b_h)$$

$$h^{<t>} = g(W_{hh}h^{<t-1>} \oplus W_{hx}x^{<t>} + b_h)$$

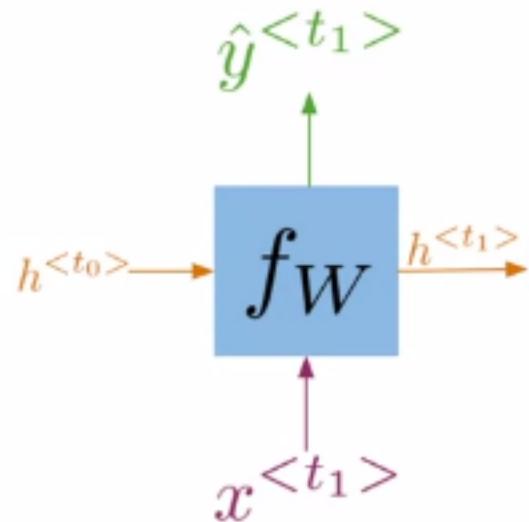
A Vanilla RNN



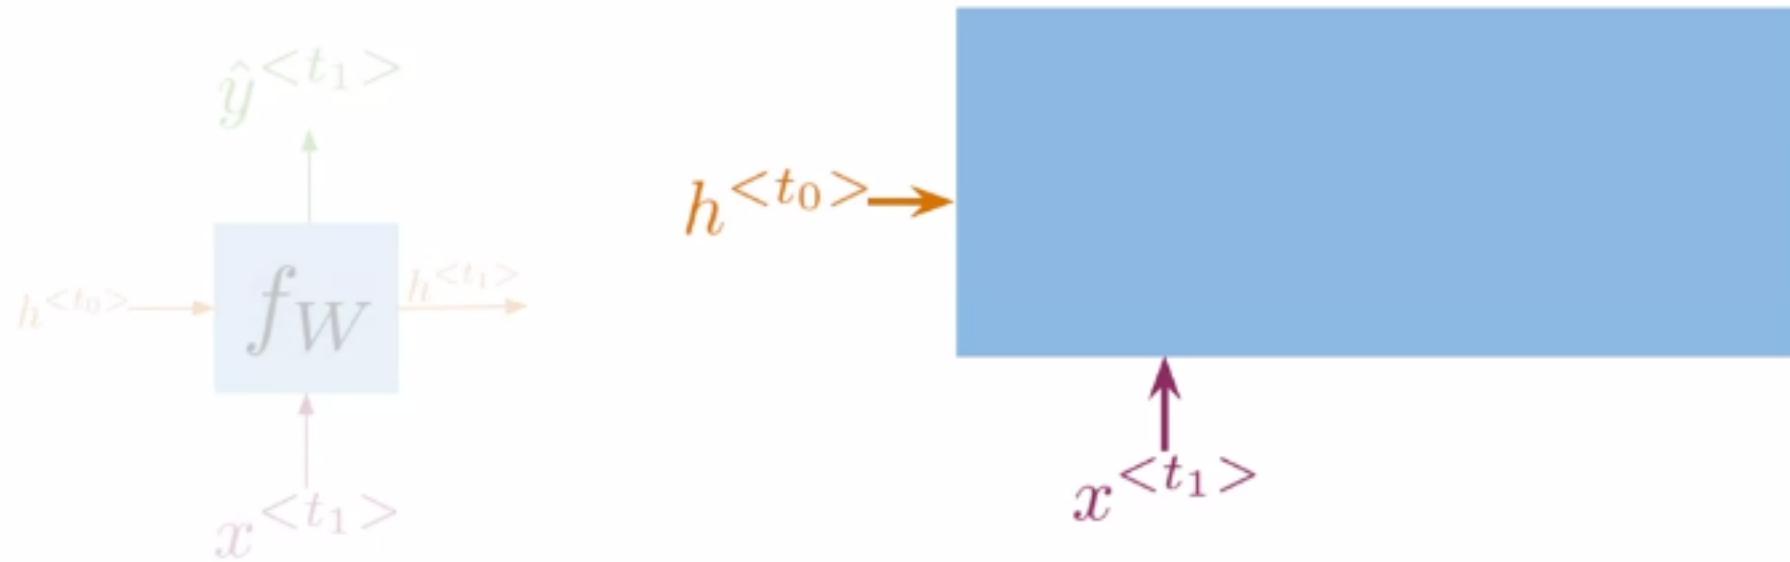
$$h^{<t>} = g(W_h[h^{<t-1>}, x^{<t>}] + b_h)$$

$$h^{<t>} = g(W_{hh}h^{<t-1>} \oplus W_{hx}x^{<t>} + b_h)$$

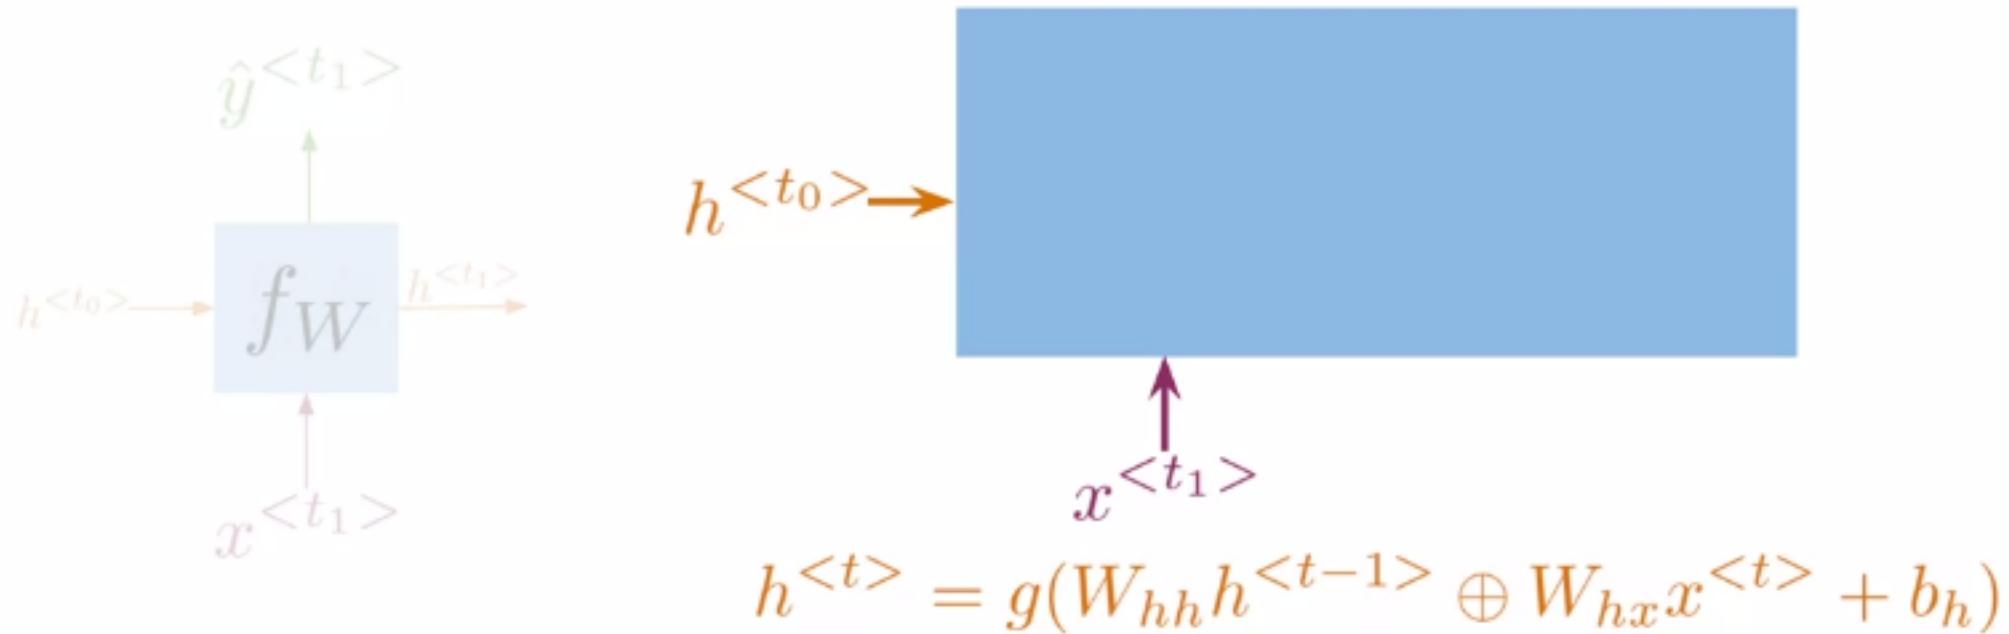
A Vanilla RNN



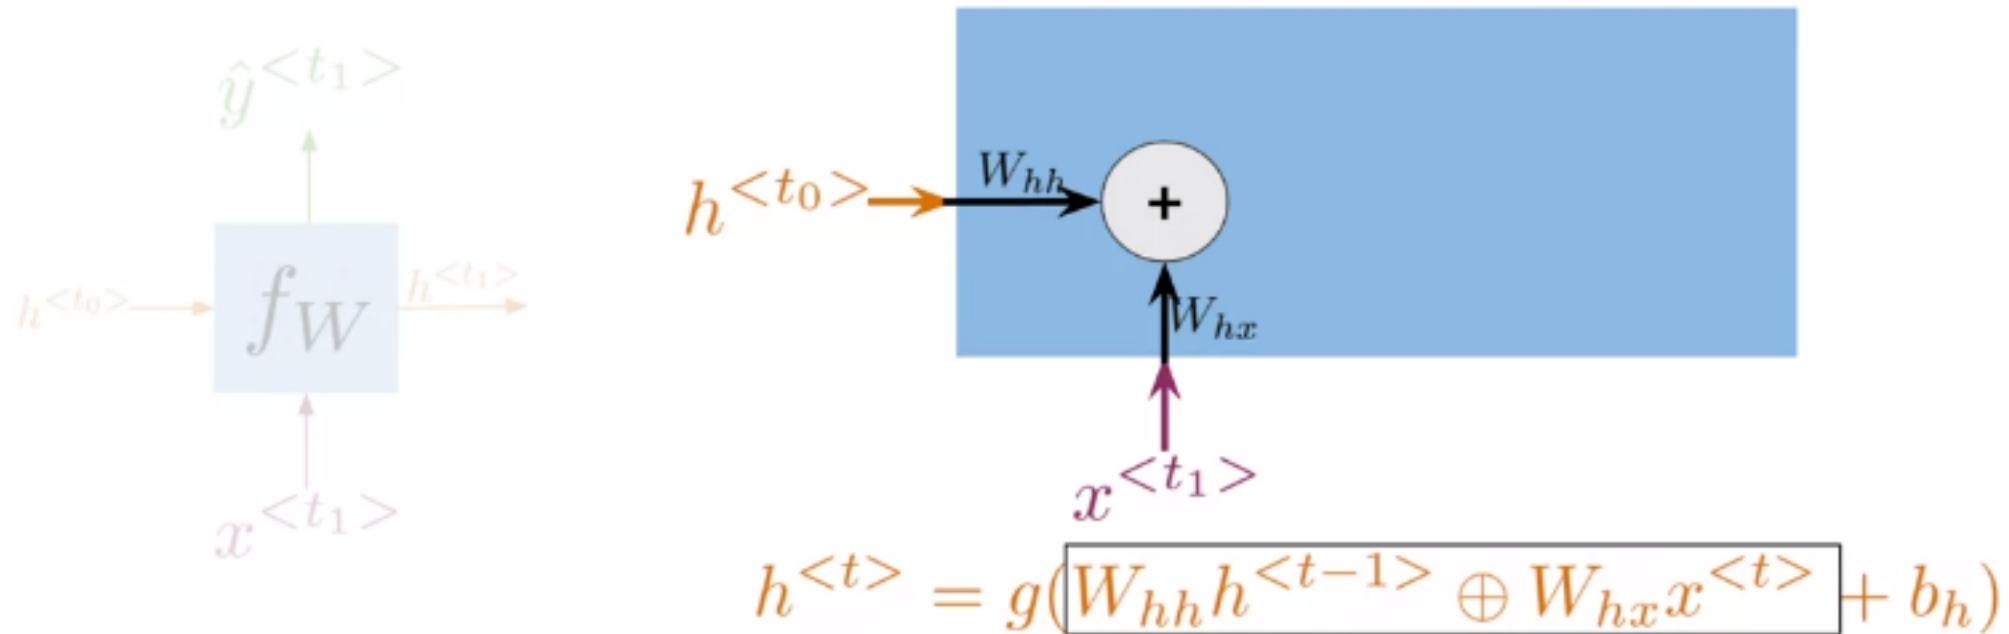
A Vanilla RNN



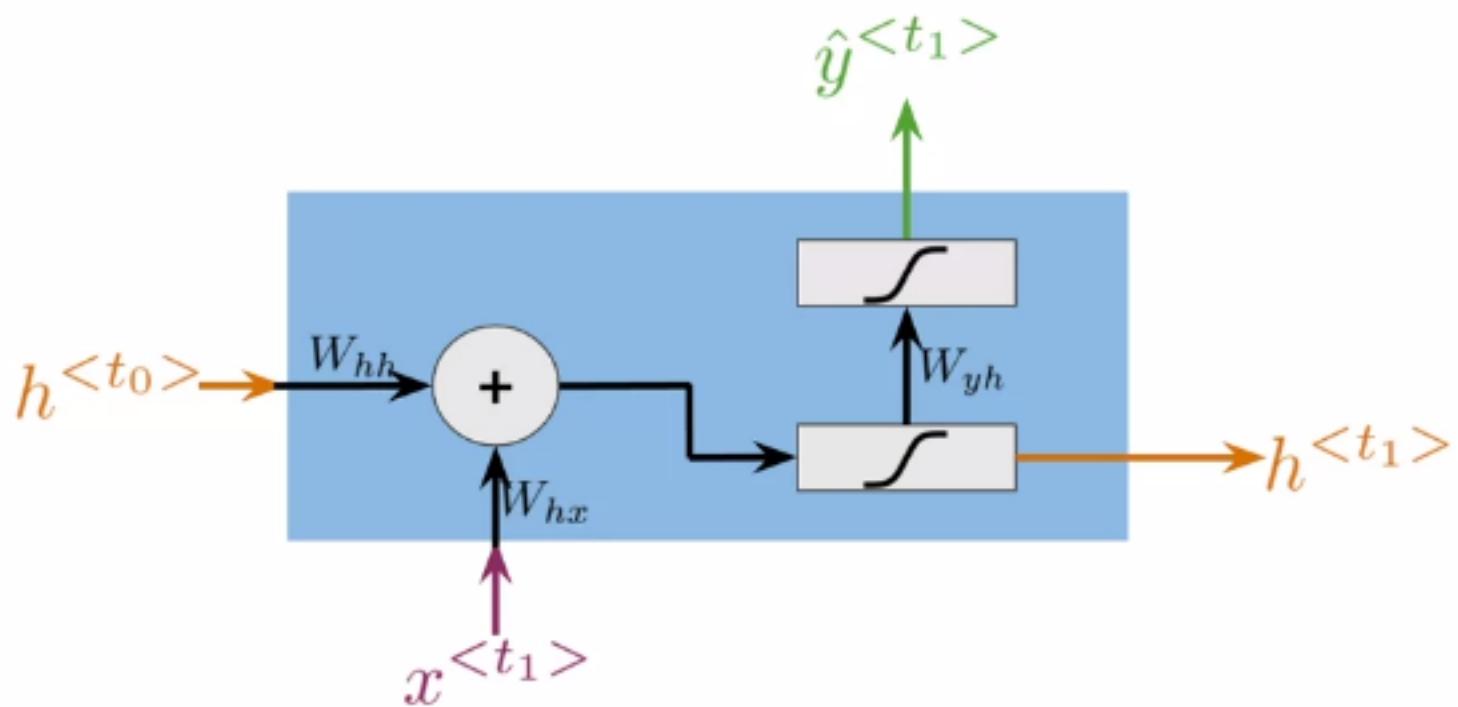
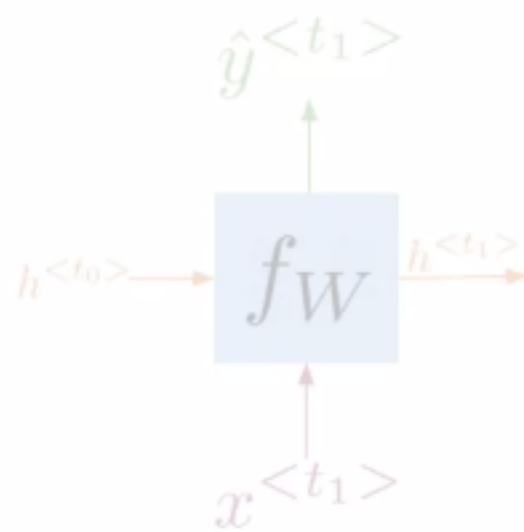
A Vanilla RNN



A Vanilla RNN



A Vanilla RNN



$$h^{<t>} = g(W_{hh}h^{<t-1>} \oplus W_{hx}x^{<t>} + b_h)$$

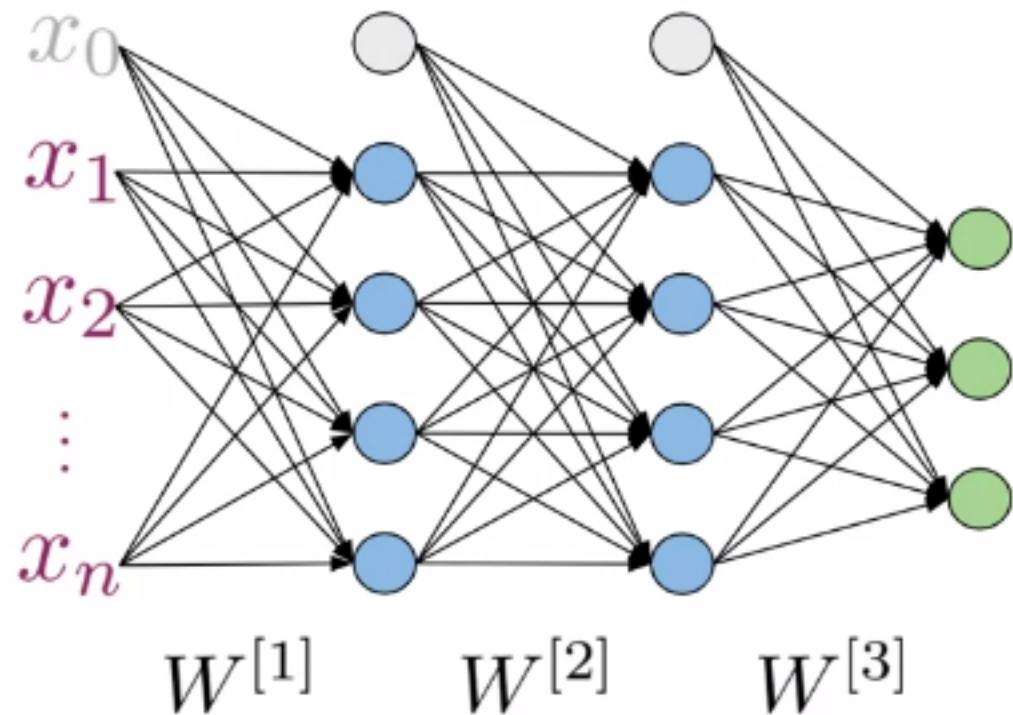
$$\hat{y}^{<t>} = g(W_{yh}h^{<t>} + b_y)$$

Summary

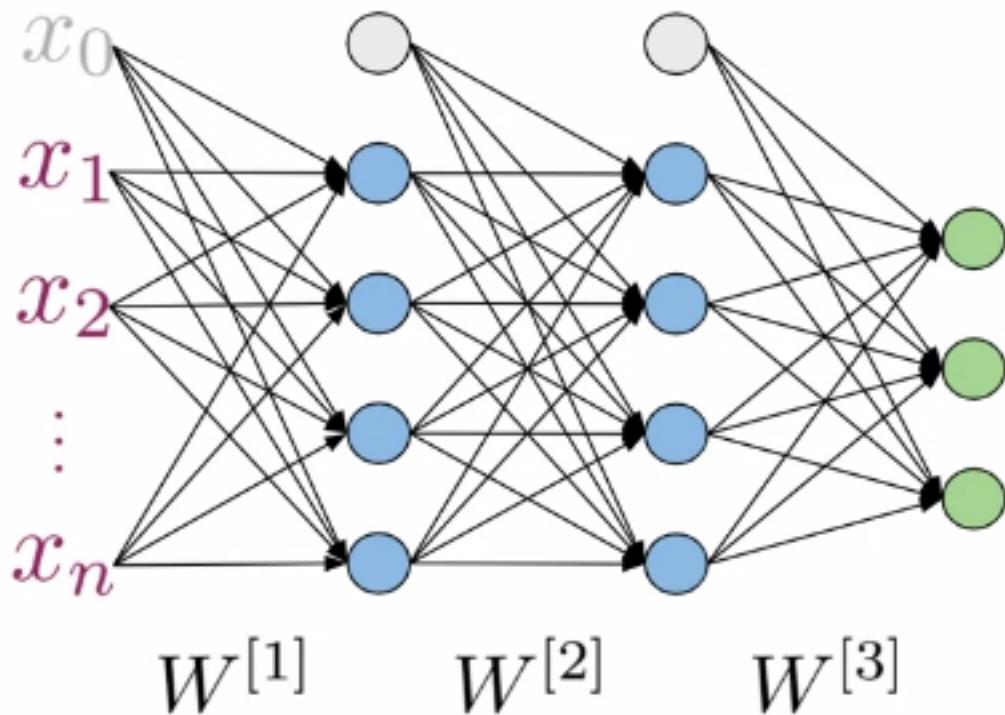
- Hidden states propagate information through time
- Basic recurrent units have two inputs at each time: $h^{}$, $x^{}$



Cross Entropy Loss



Cross Entropy Loss

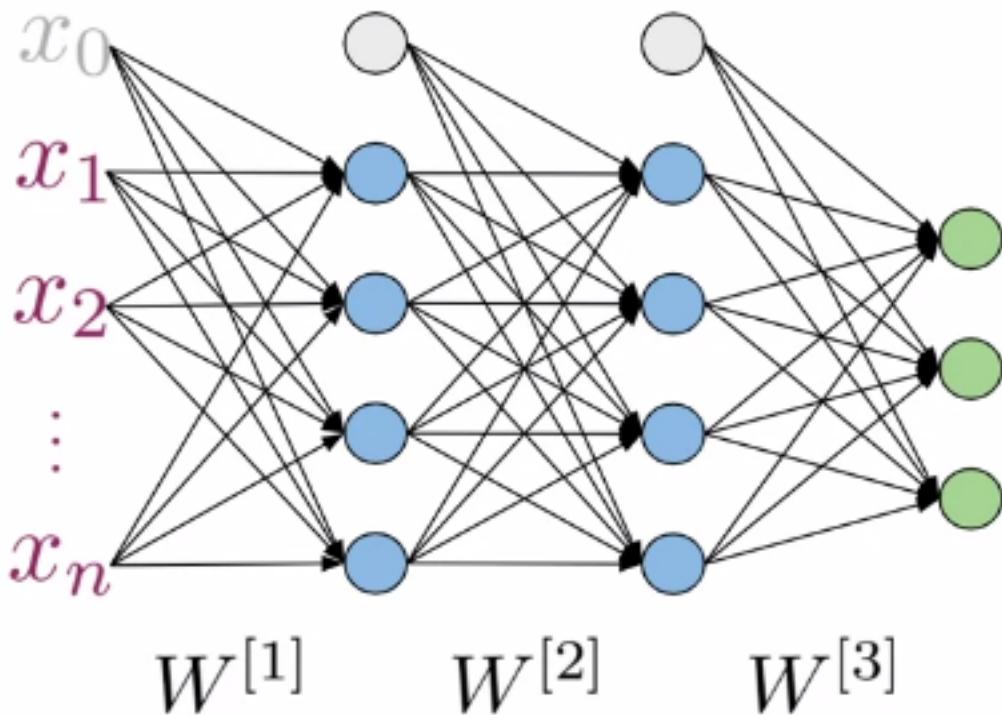


K - classes or possibilities

$$J = - \sum_{j=1}^K y_j \log \hat{y}_j$$

Looking at a single example (x, y)

Cross Entropy Loss



K - classes or possibilities

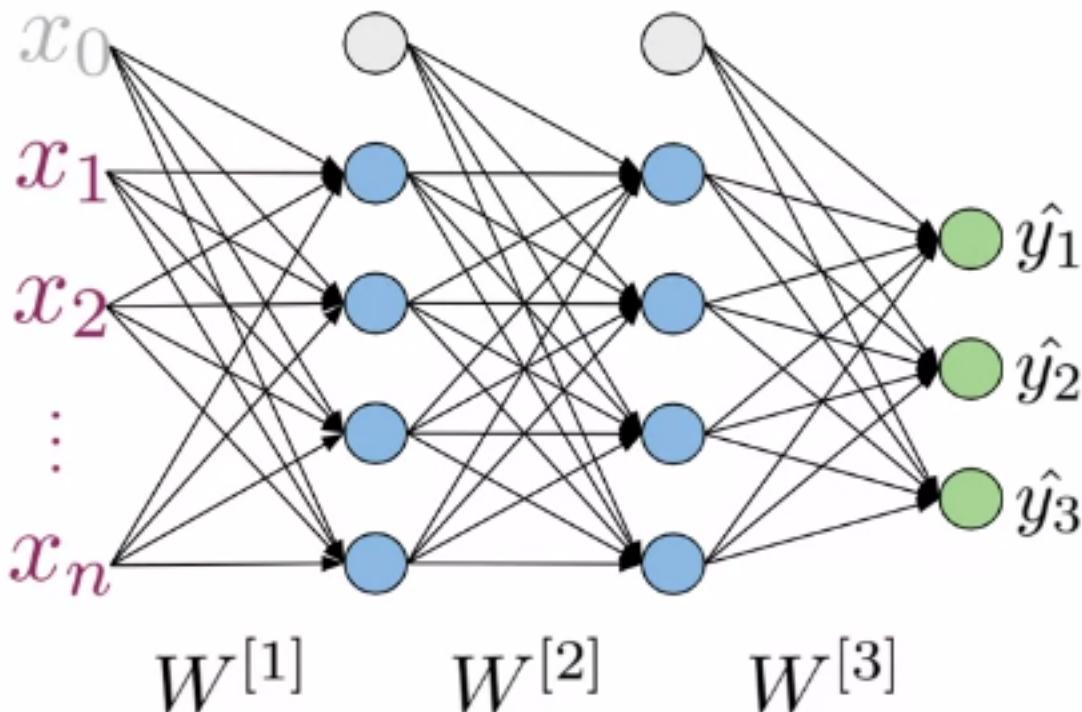
$$J = - \sum_{j=1}^K y_j \log \hat{y}_j$$

Either 0 or 1

An orange arrow points from the term y_j in the equation to a box around it, with the text "Either 0 or 1" above the arrow, indicating that y_j is a binary variable representing the true class of the input example.

Looking at a single example (x, y)

Cross Entropy Loss



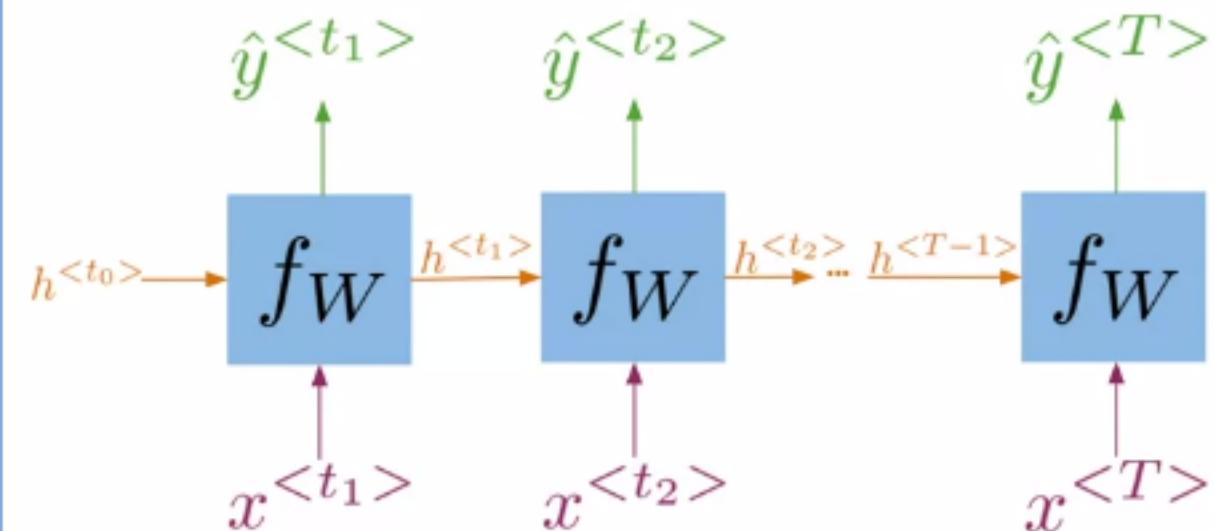
K - classes or possibilities

$$J = - \sum_{j=1}^K y_j \log \hat{y}_j$$

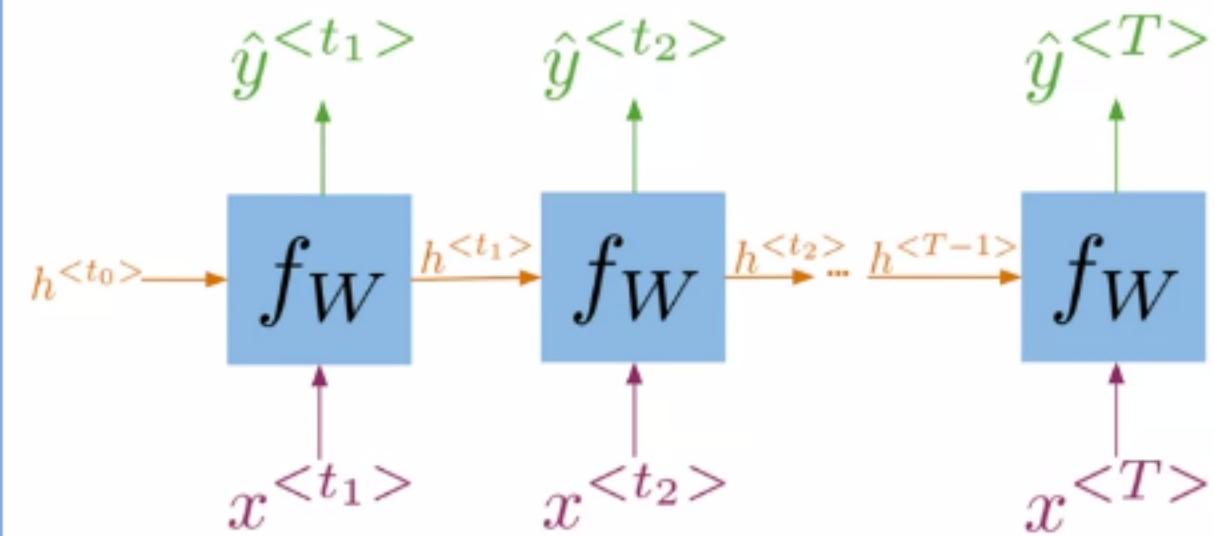
Either 0 or 1

Looking at a single example (x, y)

Cross Entropy Loss



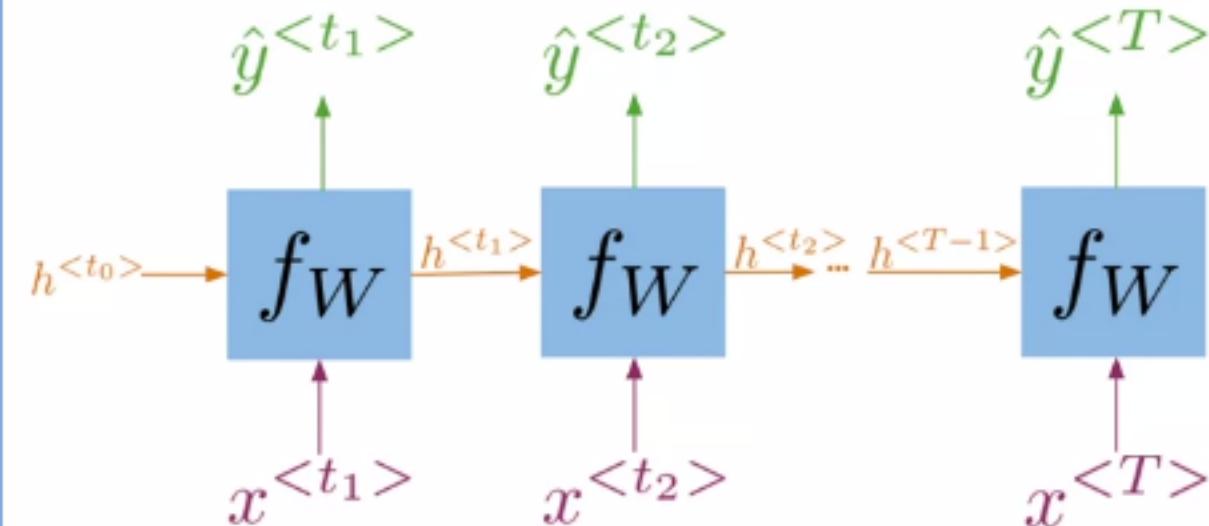
Cross Entropy Loss



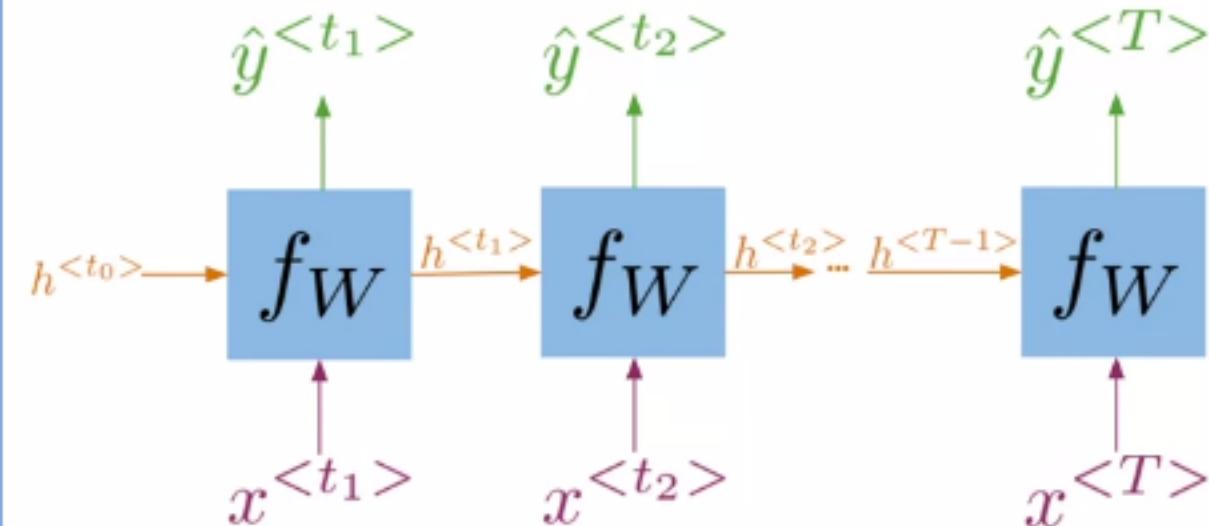
Cross Entropy Loss

$$h^{} = g(W_h[h^{}, x^{}] + b_h)$$

$$\hat{y}^{} = g(W_{yh}h^{} + b_y)$$



Cross Entropy Loss

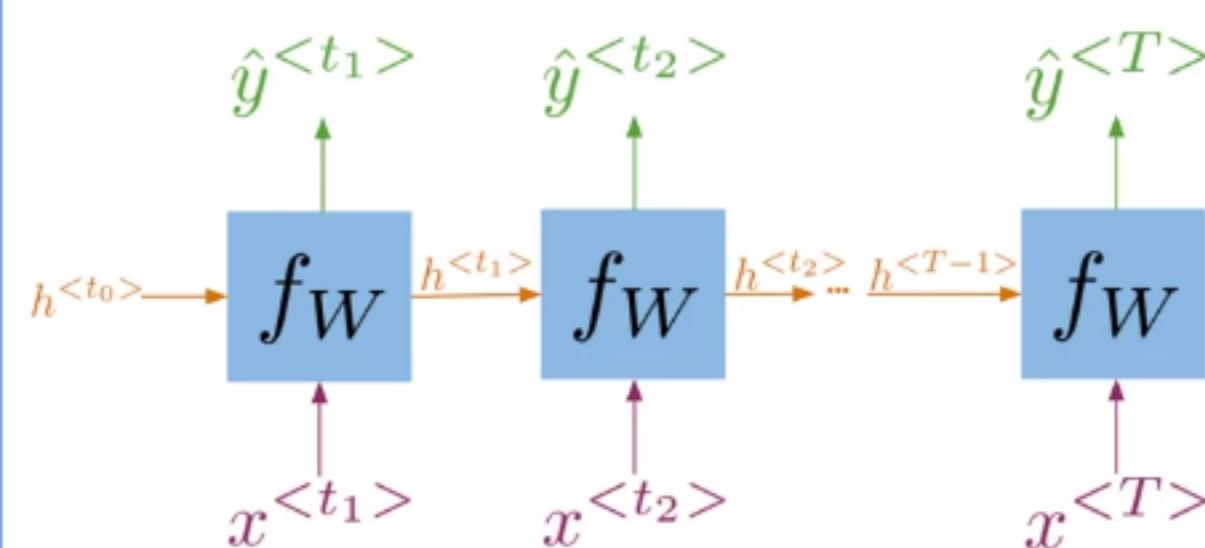


$$h^{<t>} = g(W_h[h^{<t-1>}, x^{<t>}] + b_h)$$

$$\hat{y}^{<t>} = g(W_{yh}h^{<t>} + b_y)$$

$$J = -\frac{1}{T} \sum_{t=1}^T \sum_{j=1}^K y_j^{<t>} \log \hat{y}_j^{<t>}$$

Cross Entropy Loss



$$h^{<t>} = g(W_h[h^{<t-1>}, x^{<t>}] + b_h)$$

$$\hat{y}^{<t>} = g(W_{yh}h^{<t>} + b_y)$$

$$J = -\frac{1}{T} \sum_{t=1}^T \sum_{j=1}^K y_j^{<t>} \log \hat{y}_j^{<t>}$$

Average with respect to time

Summary

For RNNs the loss function is just an average through time!

Outline

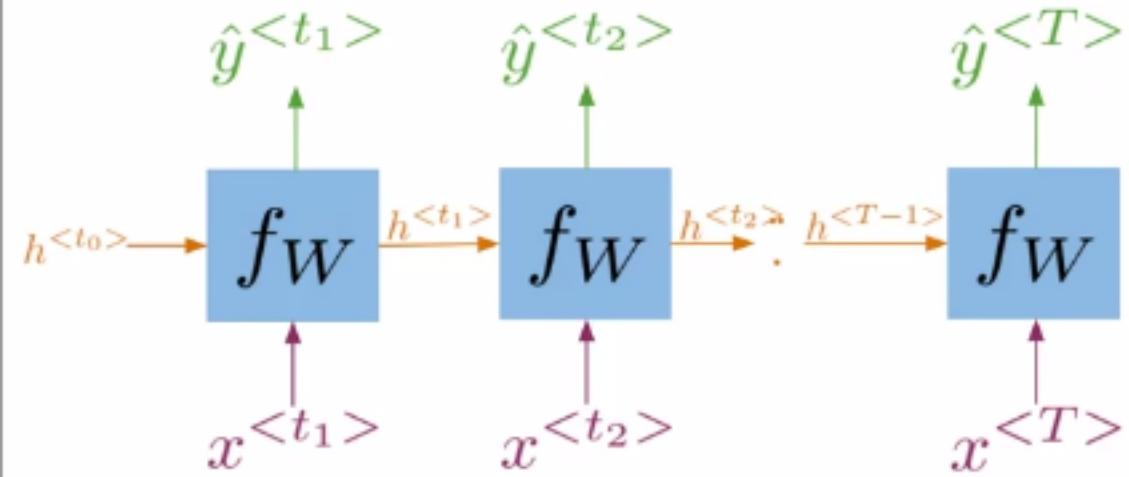
- scan() function in tensorflow
- Computation of forward propagation using abstractions



tf.scan() function

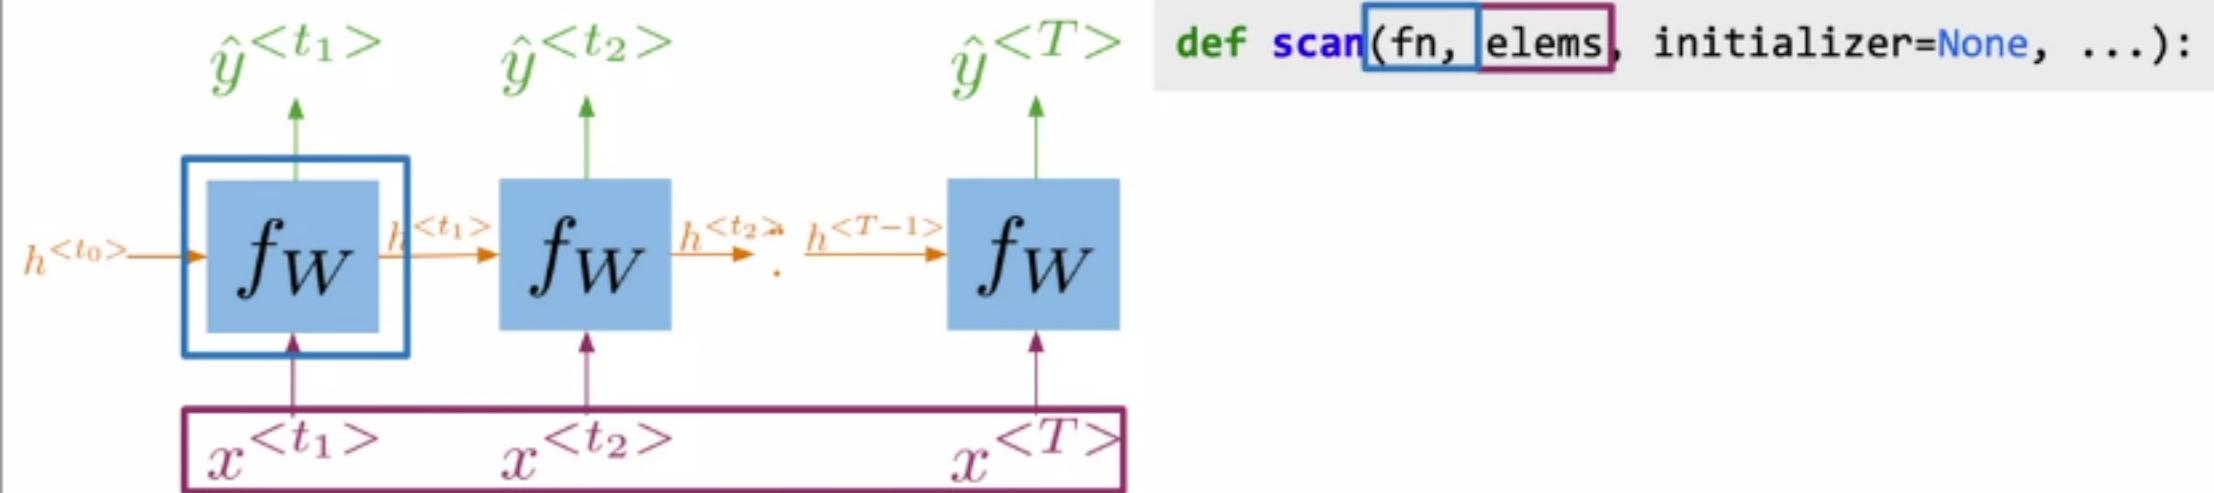
```
def scan(fn, elems, initializer=None, ...):
```

tf.scan() function

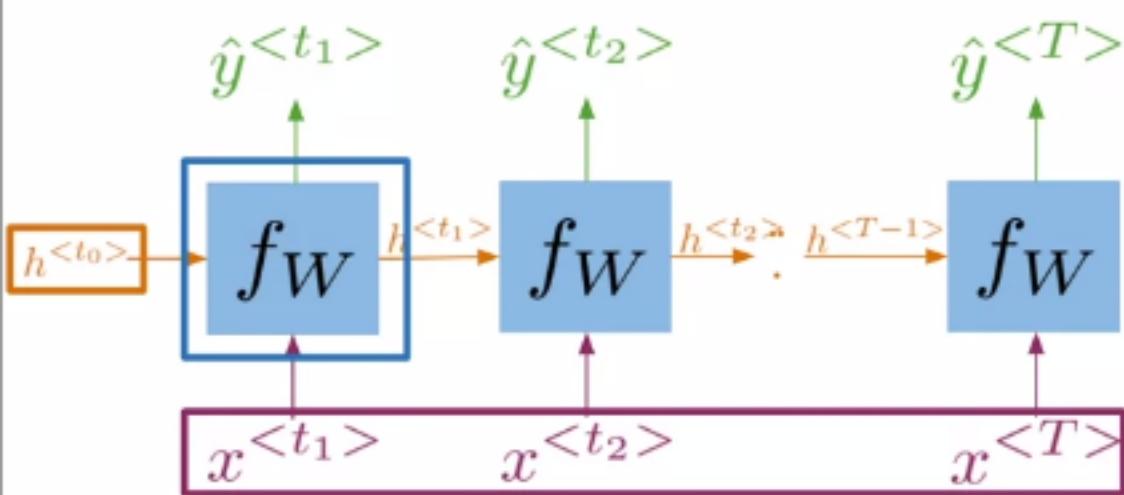


```
def scan(fn, elems, initializer=None, ...):
```

tf.scan() function

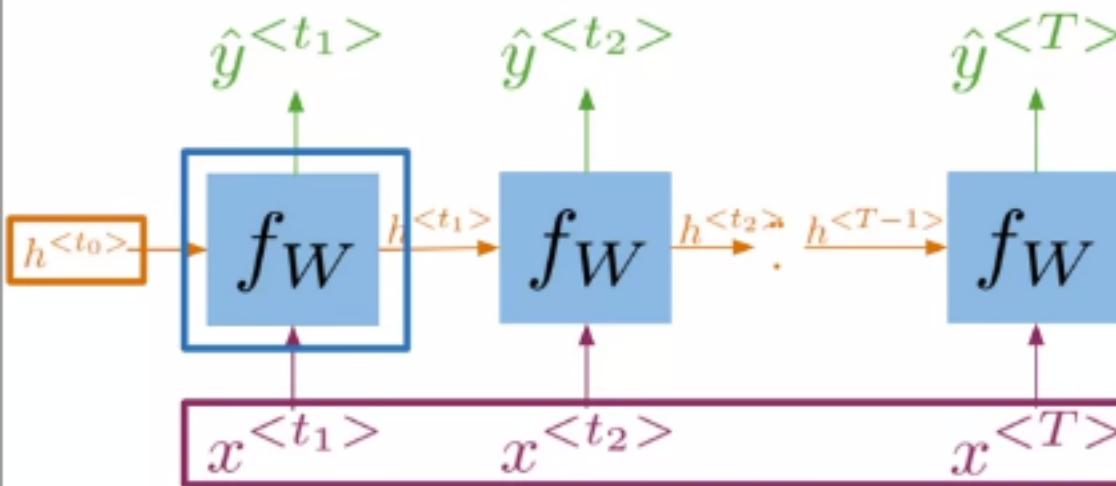


tf.scan() function



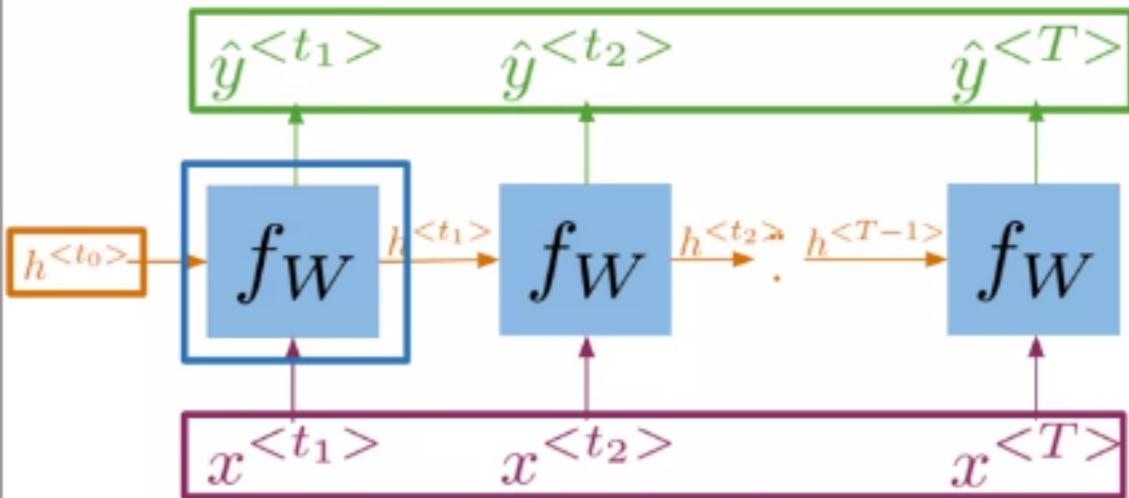
```
def scan(fn, elems, initializer=None, ...):
```

tf.scan() function



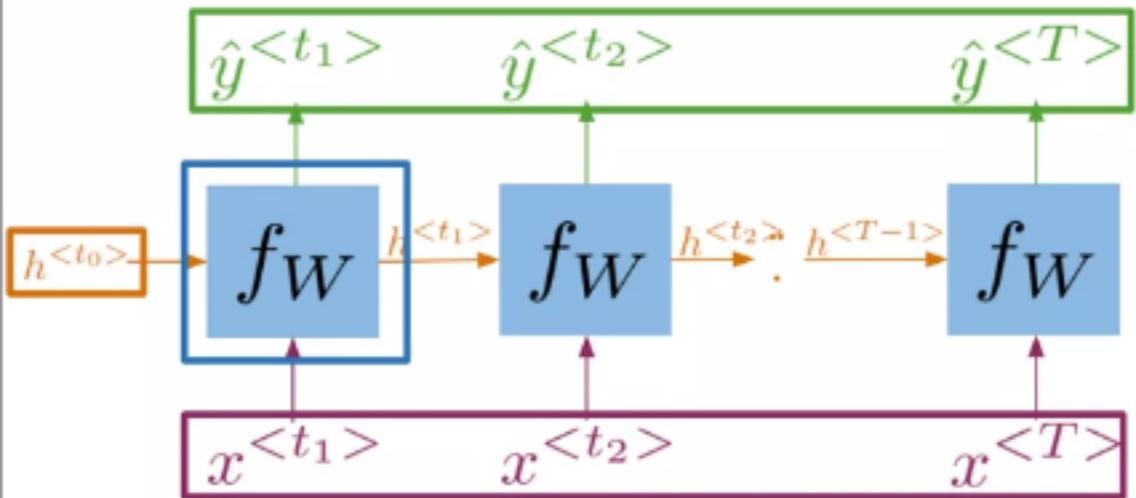
```
def scan(fn, elems, initializer=None, ...):  
    cur_value = initializer  
    ys = []
```

tf.scan() function



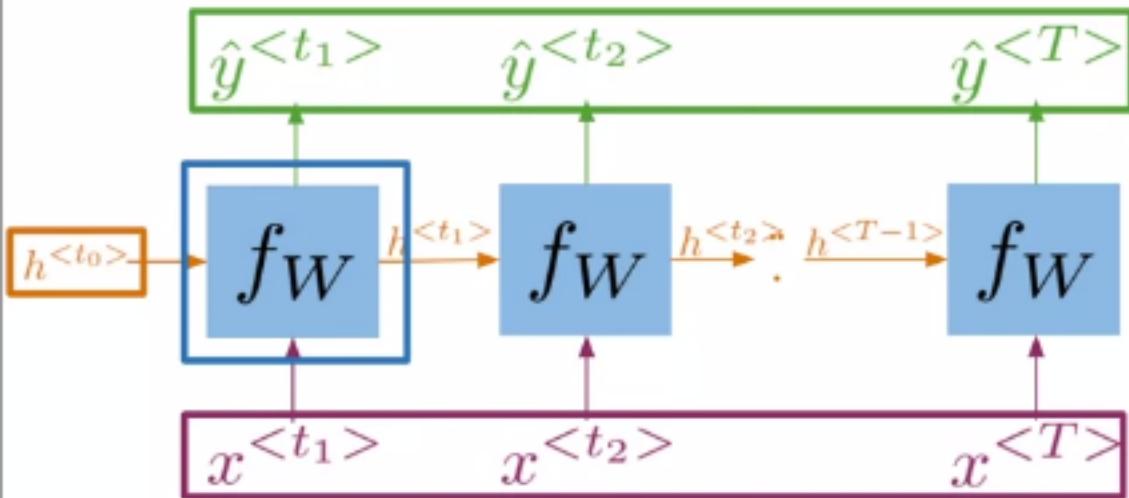
```
def scan(fn, elems, initializer=None, ...):  
    cur_value = initializer  
    ys = []
```

tf.scan() function



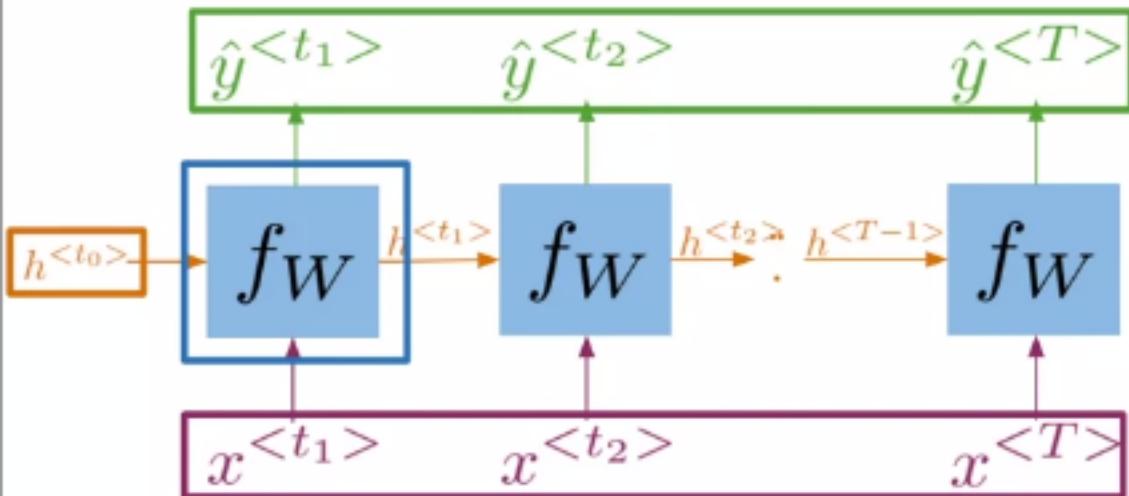
```
def scan(fn, elems, initializer=None, ...):  
    cur_value = initializer  
    ys = []  
    for x in elems:
```

tf.scan() function



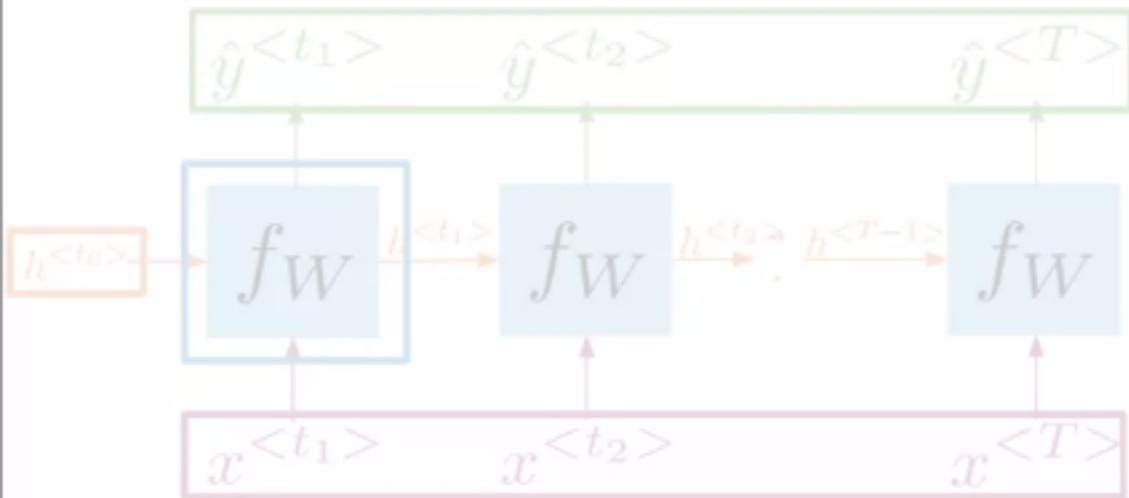
```
def scan(fn, elems, initializer=None, ...):
    cur_value = initializer
    ys = []
    for x in elems:
        y, cur_value = fn(x, cur_value)
        ys.append(y)
```

tf.scan() function



```
def scan(fn, elems, initializer=None, ...):
    cur_value = initializer
    ys = []
    for x in elems:
        y, cur_value = fn(x, cur_value)
        ys.append(y)
    return ys, cur_value
```

tf.scan() function

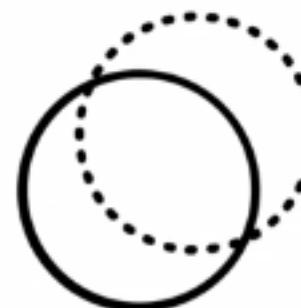


```
def scan(fn, elems, initializer=None, ...):  
    cur_value = initializer  
    ys = []  
    for x in elems:  
        y, cur_value = fn(x, cur_value)  
        ys.append(y)  
    return ys, cur_value
```

Frameworks like Tensorflow need this type of abstraction

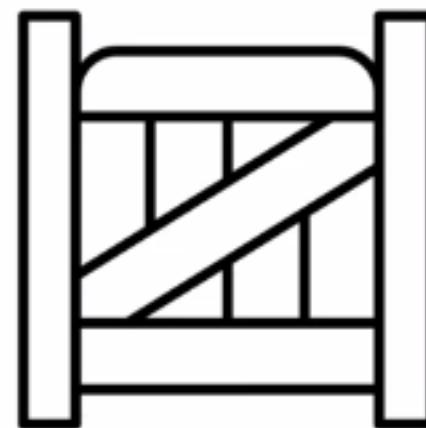
Summary

- Frameworks require abstractions
- `tf.scan()` mimics RNNs



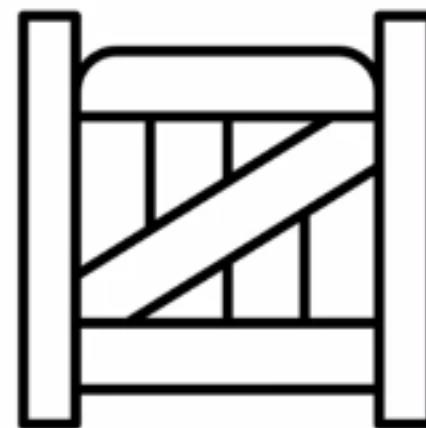
Outline

- Gated recurrent unit (GRU) structure
- Comparison between GRUs and vanilla RNNs



Outline

- Gated recurrent unit (GRU) structure
- Comparison between GRUs and vanilla RNNs



Gated Recurrent Units

“Ants are really interesting. _____ are everywhere.”

Gated Recurrent Units

“Ants are really interesting. _____ are everywhere.”

Gated Recurrent Units

“**Ants** are really interesting. They are everywhere.”

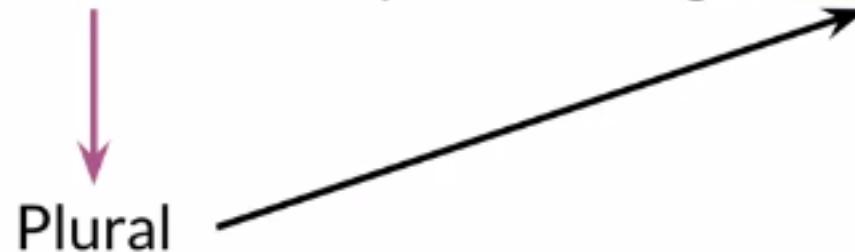
Gated Recurrent Units

“Ants are really interesting. They are everywhere.”



Gated Recurrent Units

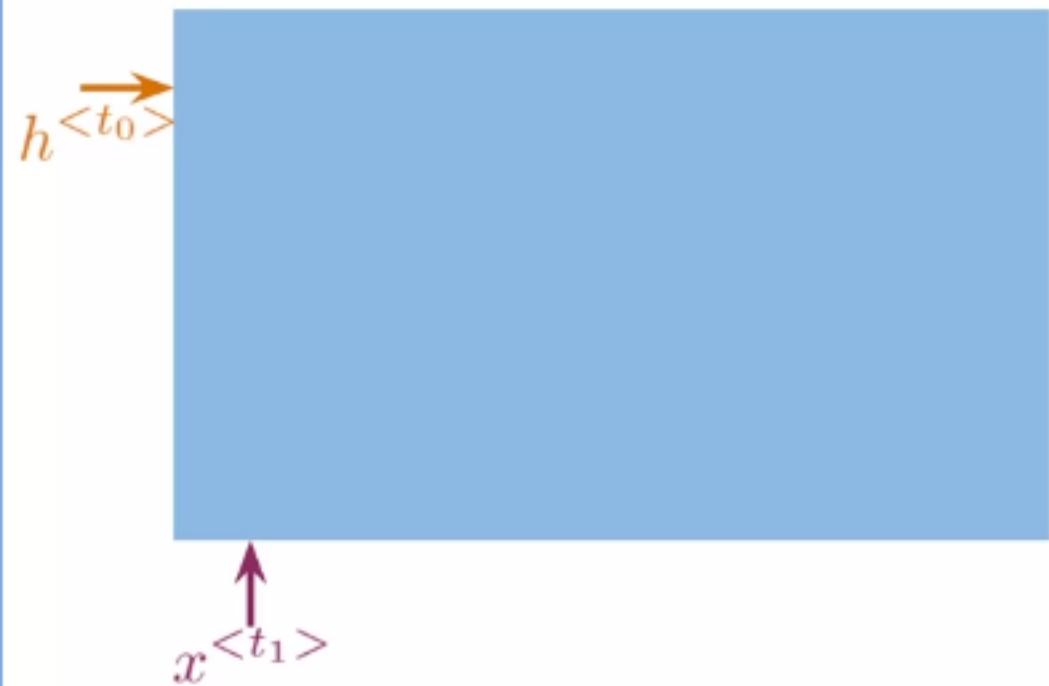
“Ants are really interesting. They are everywhere.”



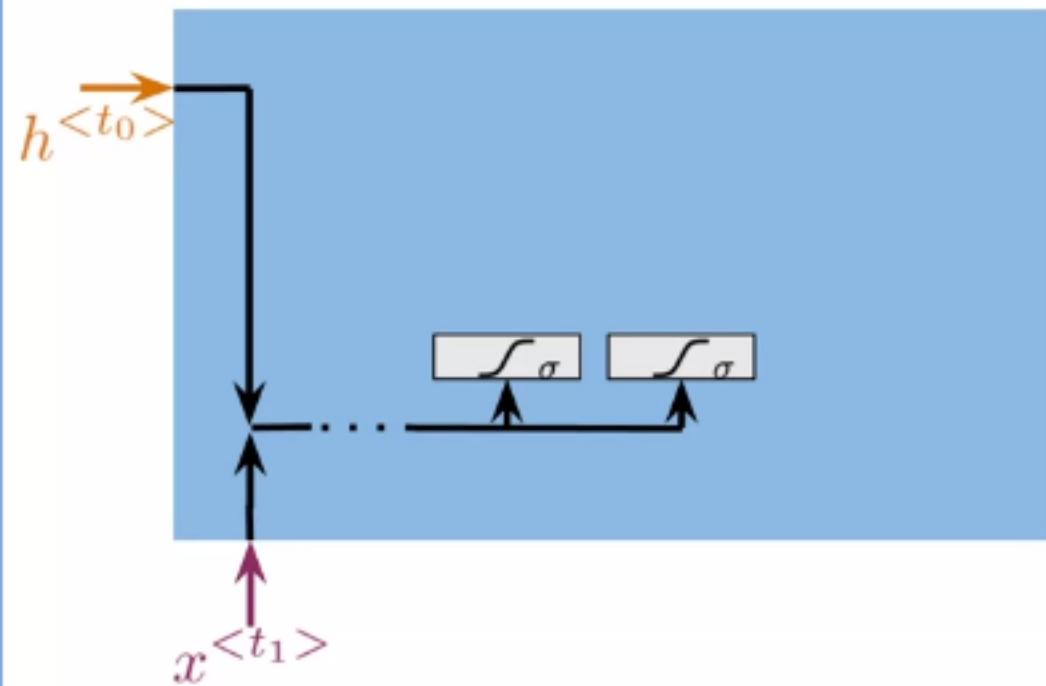
Relevance and update gates to remember important prior information

Gated Recurrent Unit

Gated Recurrent Unit



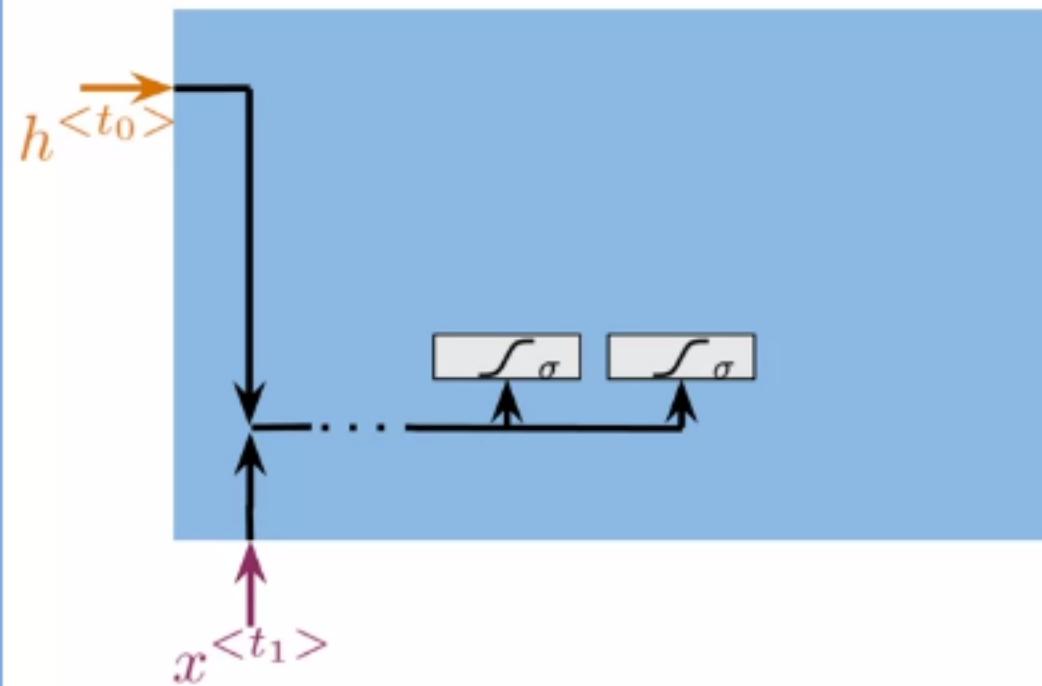
Gated Recurrent Unit



$$\Gamma_r = \sigma(W_r[h^{<t_0>}, x^{<t_1>}] + b_r)$$

$$\Gamma_u = \sigma(W_u[h^{<t_0>}, x^{<t_1>}] + b_u)$$

Gated Recurrent Unit



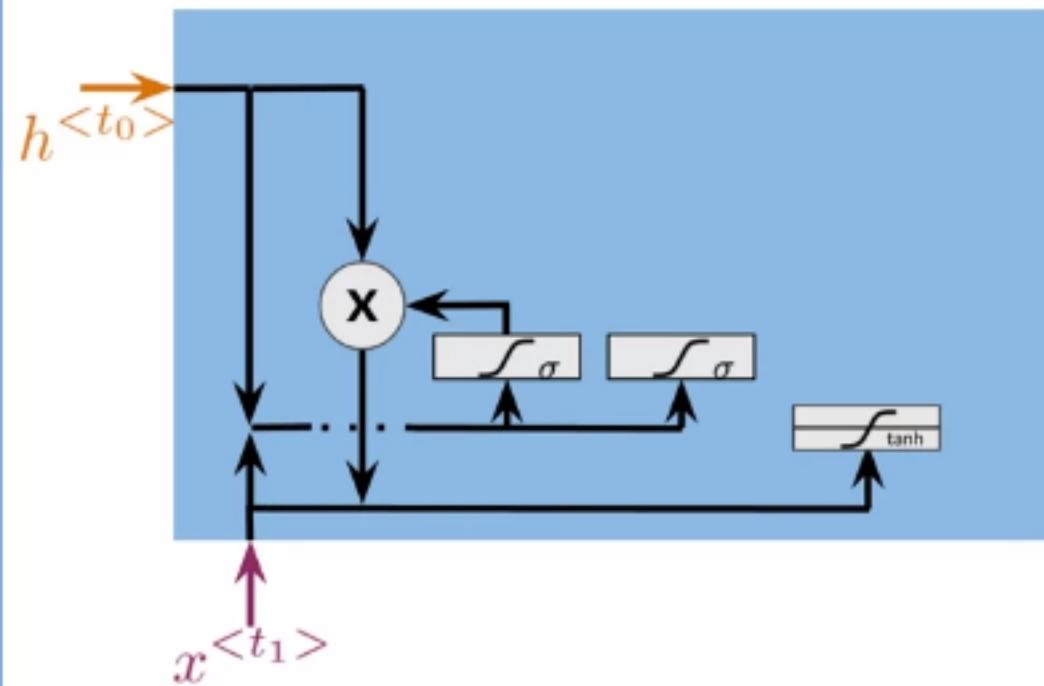
Gates to keep/update relevant
information in the hidden state

$$\Gamma_r = \sigma(W_r[h^{<t_0>}], x^{<t_1>}]) + b_r)$$

$$\Gamma_u = \sigma(W_u[h^{<t_0>}], x^{<t_1>}]) + b_u)$$

Gated Recurrent Unit

Gates to keep/update relevant information in the hidden state



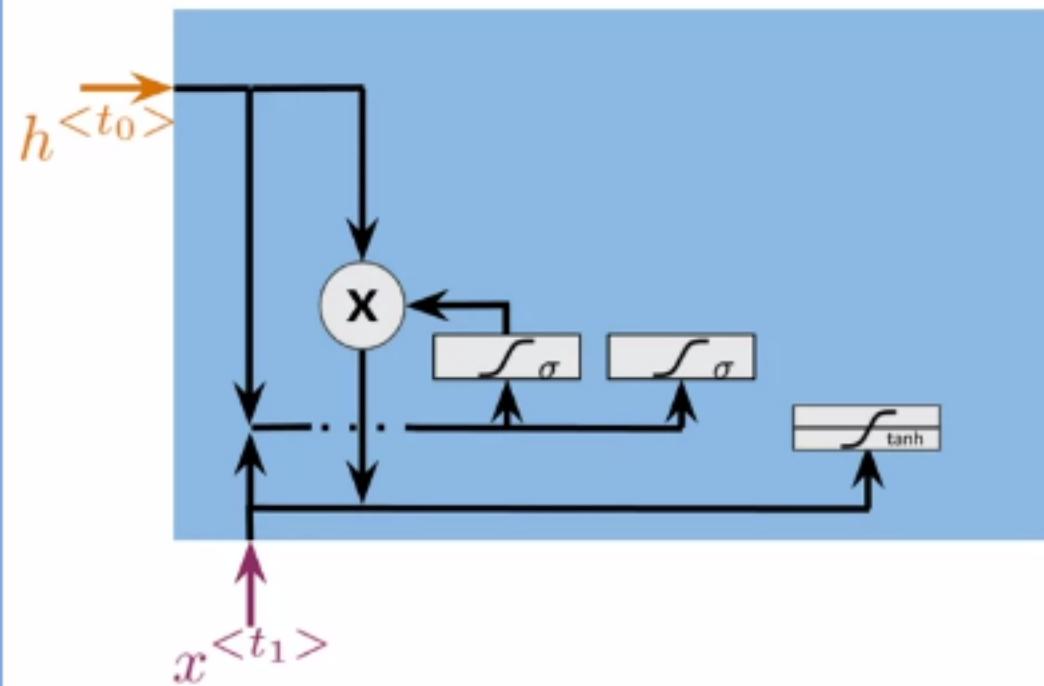
$$\Gamma_r = \sigma(W_r[h^{t_0}, x^{t_1}] + b_r)$$

$$\Gamma_u = \sigma(W_u[h^{t_0}, x^{t_1}] + b_u)$$

$$h'^{t_1} = \tanh(W_h[\Gamma_r * h^{t_0}, x^{t_1}] + b_h)$$

Gated Recurrent Unit

Gates to keep/update relevant information in the hidden state

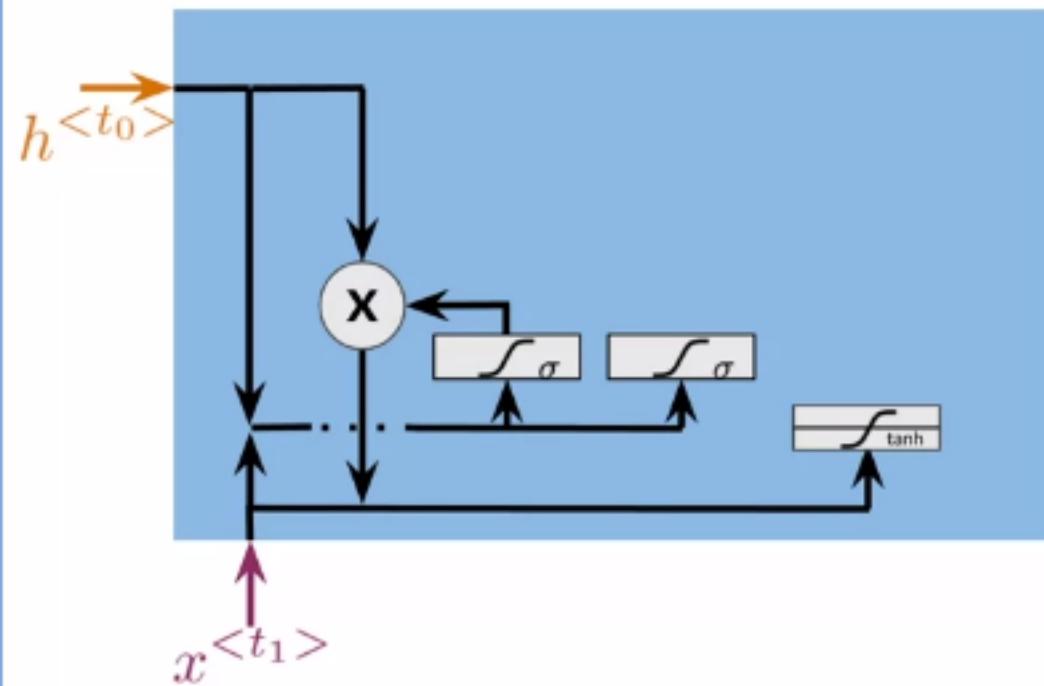


$$\Gamma_r = \sigma(W_r[h^{t_0}, x^{t_1}] + b_r)$$

$$\Gamma_u = \sigma(W_u[h^{t_0}, x^{t_1}] + b_u)$$

$$h'^{t_1} = \tanh(W_h[\Gamma_r * h^{t_0}, x^{t_1}] + b_h)$$

Gated Recurrent Unit



Gates to keep/update relevant information in the hidden state

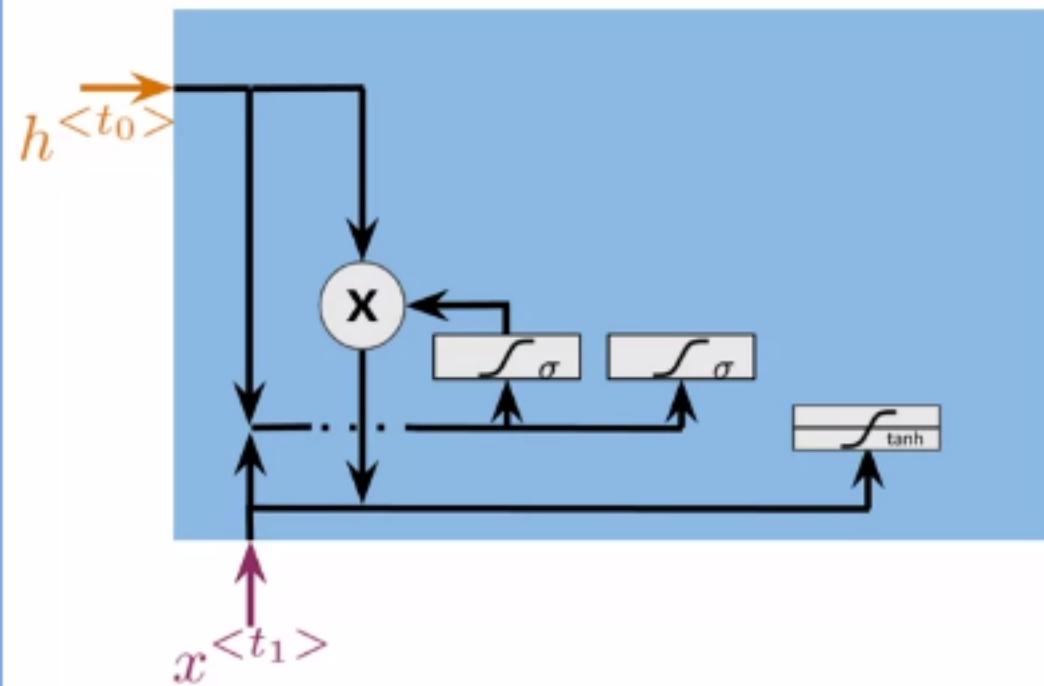
$$\Gamma_r = \sigma(W_r[h^{<t_0>}, x^{<t_1>}] + b_r)$$

$$\Gamma_u = \sigma(W_u[h^{<t_0>}, x^{<t_1>}] + b_u)$$

$$h'^{<t_1>} = \tanh(W_h[\Gamma_r * h^{<t_0>}, x^{<t_1>}] + b_h)$$

Hidden state candidate

Gated Recurrent Unit



Gates to keep/update relevant information in the hidden state

$$\Gamma_r = \sigma(W_r[h^{<t_0>}, x^{<t_1>}] + b_r)$$

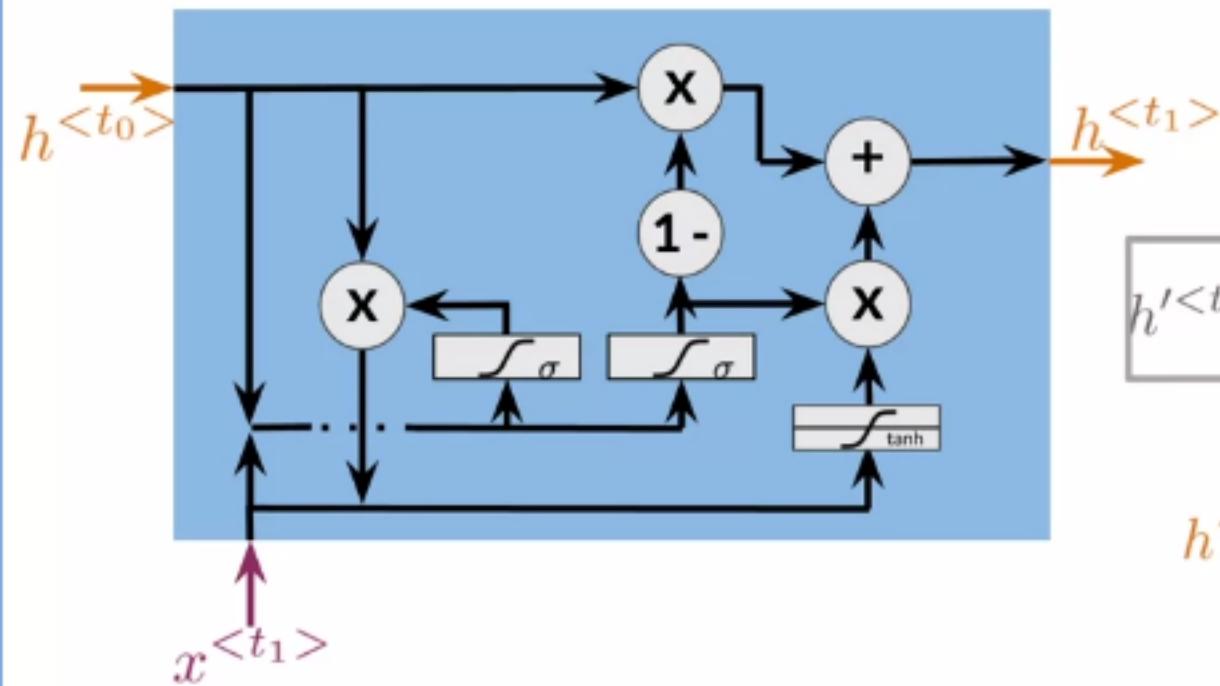
$$\Gamma_u = \sigma(W_u[h^{<t_0>}, x^{<t_1>}] + b_u)$$

$$h'^{<t_1>} = \tanh(W_h[\Gamma_r * h^{<t_0>}, x^{<t_1>}] + b_h)$$

Hidden state candidate

Gated Recurrent Unit

Gates to keep/update relevant information in the hidden state



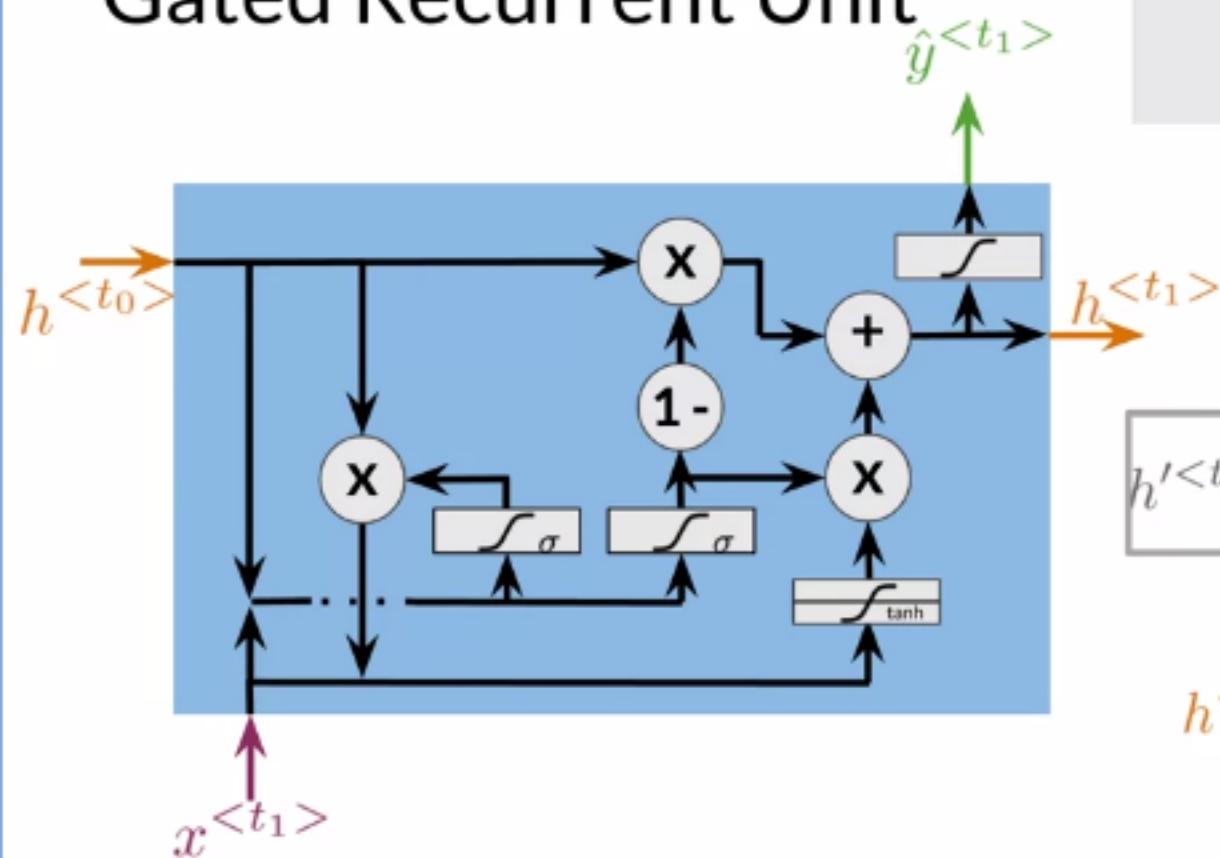
$$\Gamma_r = \sigma(W_r[h^{<t_0>}, x^{<t_1>}] + b_r)$$
$$\Gamma_u = \sigma(W_u[h^{<t_0>}, x^{<t_1>}] + b_u)$$

$$h'^{<t_1>} = \tanh(W_h[\Gamma_r * h^{<t_0>}, x^{<t_1>}] + b_h)$$

Hidden state candidate

$$h^{<t_1>} = \Gamma_u * h^{<t_0>} + (1 - \Gamma_u) * h'^{<t_1>}$$

Gated Recurrent Unit



Gates to keep/update relevant information in the hidden state

$$\Gamma_r = \sigma(W_r[h^{t_0}, x^{t_1}] + b_r)$$

$$\Gamma_u = \sigma(W_u[h^{t_0}, x^{t_1}] + b_u)$$

$$h'^{t_1} = \tanh(W_h[\Gamma_r * h^{t_0}, x^{t_1}] + b_h)$$

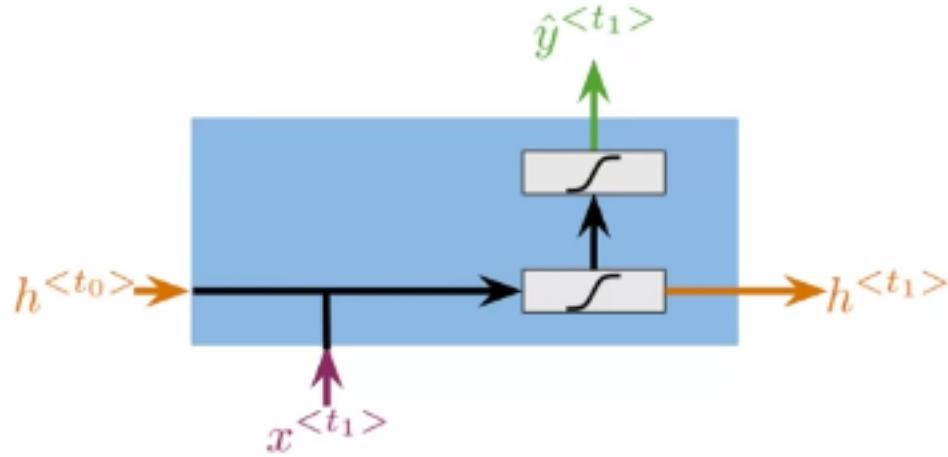
Hidden state candidate

$$h^{t_1} = \Gamma_u * h^{t_0} + (1 - \Gamma_u) * h'^{t_1}$$

$$\hat{y}^{t_1} = g(W_y h^{t_1} + b_y)$$

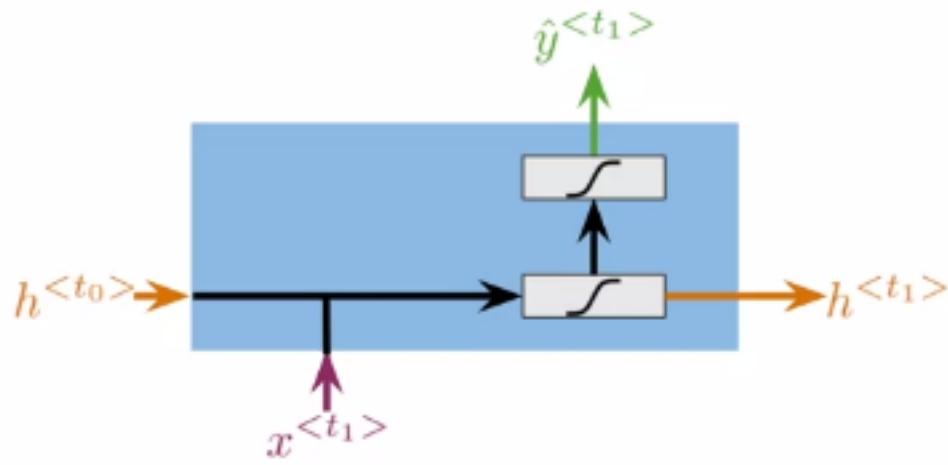
Vanilla RNN vs GRUs

Vanilla RNN vs GRUs



$$h^{<t>} = g(W_h[h^{<t-1>}, x^{<t>}] + b_h)$$

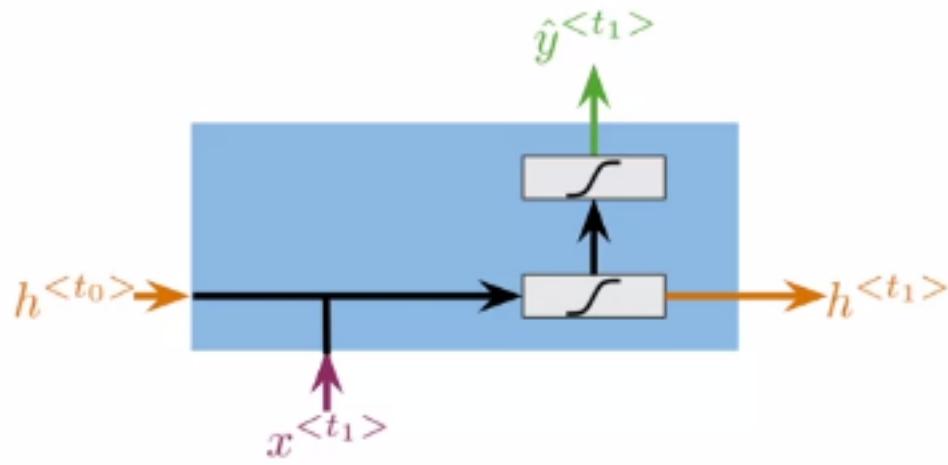
Vanilla RNN vs GRUs



$$h^{<t>} = g(W_h[h^{<t-1>}, x^{<t>}] + b_h)$$

$$\hat{y}^{<t>} = g(W_{yh}h^{<t>} + b_y)$$

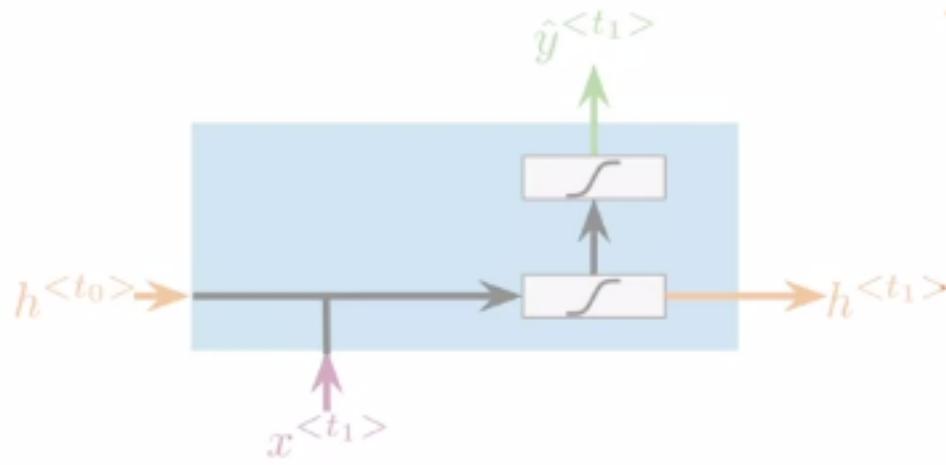
Vanilla RNN vs GRUs



$$h^{<t>} = g(W_h[h^{<t-1>}, x^{<t>}] + b_h)$$

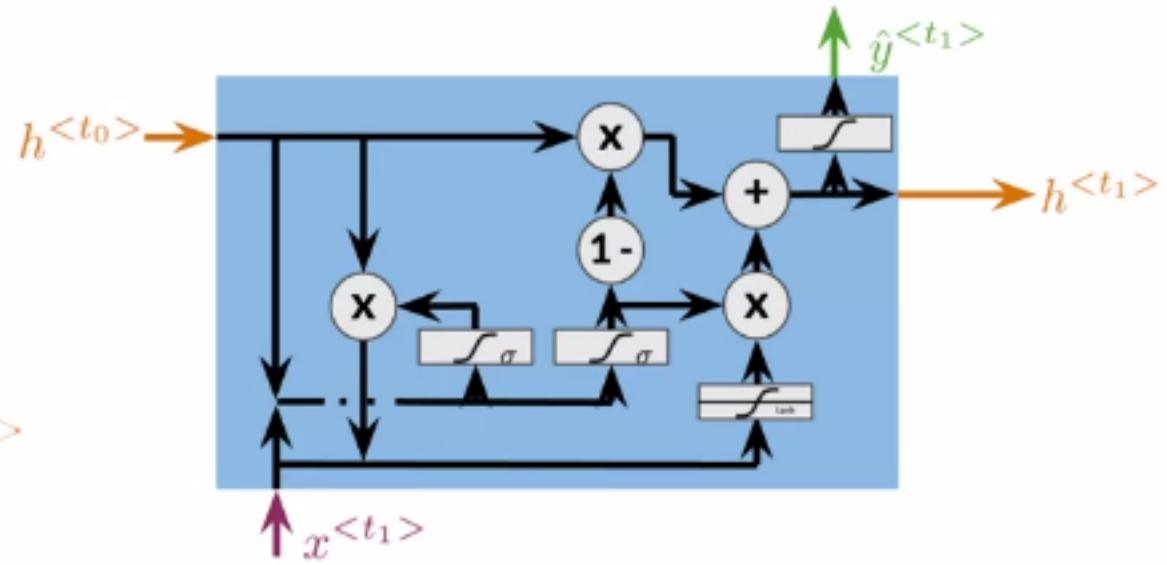
$$\hat{y}^{<t>} = g(W_{yh}h^{<t>} + b_y)$$

Vanilla RNN vs GRUs

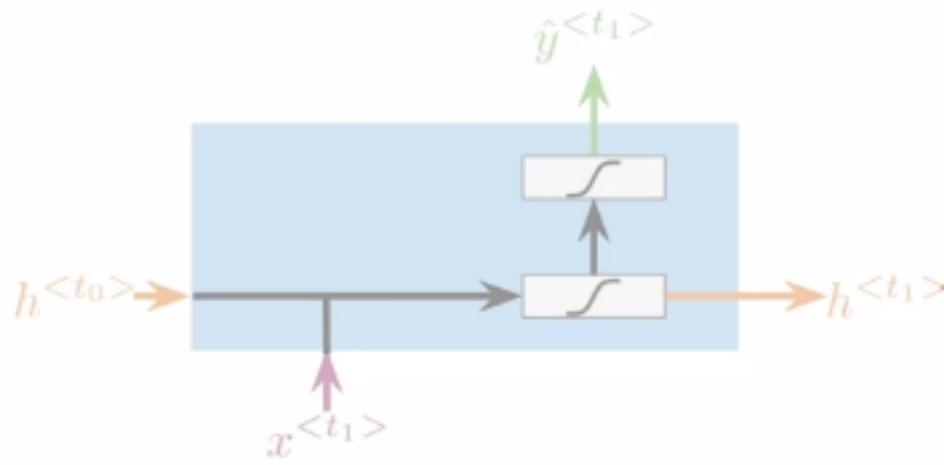


$$h^{<t>} = g(W_h[h^{<t-1>}, x^{<t>}] + b_h)$$

$$\hat{y}^{<t>} = g(W_{yh}h^{<t>} + b_y)$$

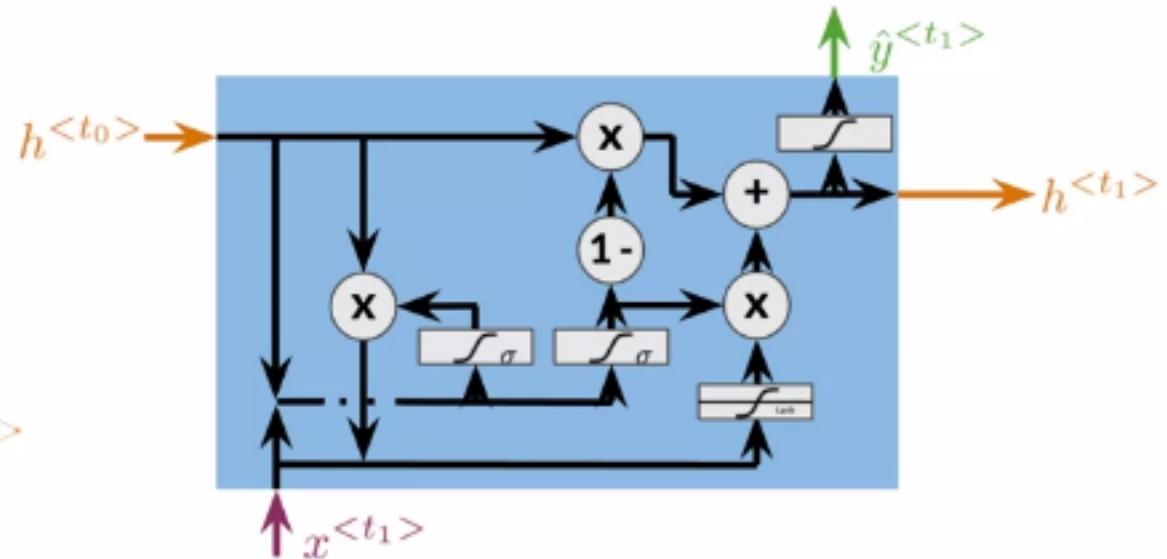


Vanilla RNN vs GRUs



$$h^{<t>} = g(W_h[h^{<t-1>}, x^{<t>}] + b_h)$$

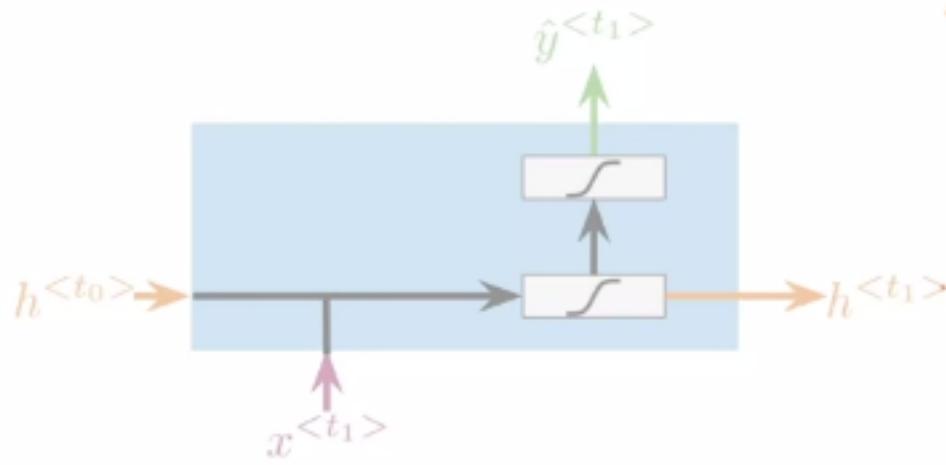
$$\hat{y}^{<t>} = g(W_{yh}h^{<t>} + b_y)$$



$$\Gamma_u = \sigma(W_u[h^{<t_0>}, x^{<t_1>}] + b_u)$$

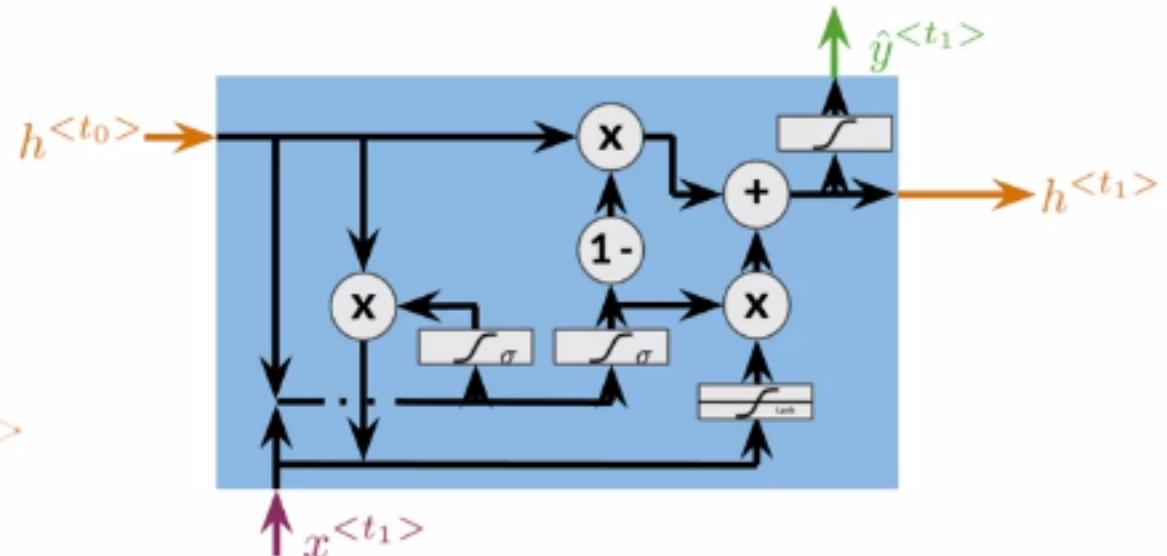
$$\Gamma_r = \sigma(W_r[h^{<t_0>}, x^{<t_1>}] + b_r)$$

Vanilla RNN vs GRUs



$$h^{<t>} = g(W_h[h^{<t-1>}, x^{<t>}] + b_h)$$

$$\hat{y}^{<t>} = g(W_{yh}h^{<t>} + b_y)$$

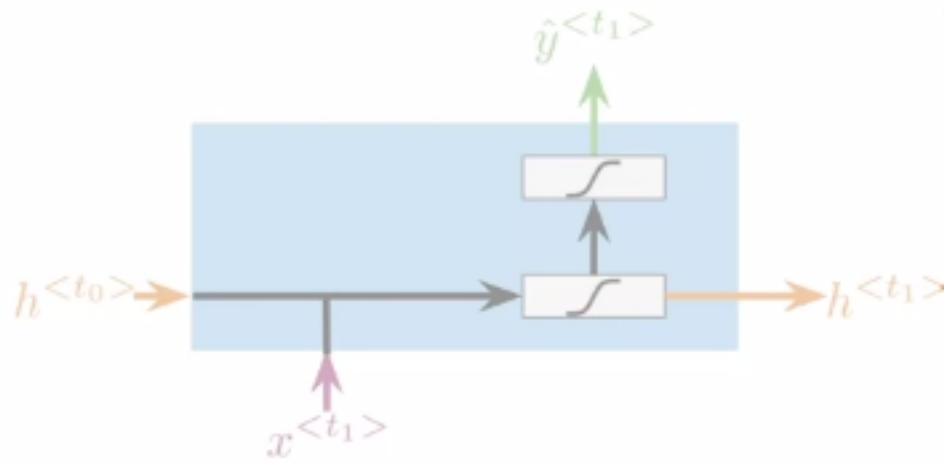


$$\Gamma_u = \sigma(W_u[h^{<t_0>}, x^{<t_1>}] + b_u)$$

$$\Gamma_r = \sigma(W_r[h^{<t_0>}, x^{<t_1>}] + b_r)$$

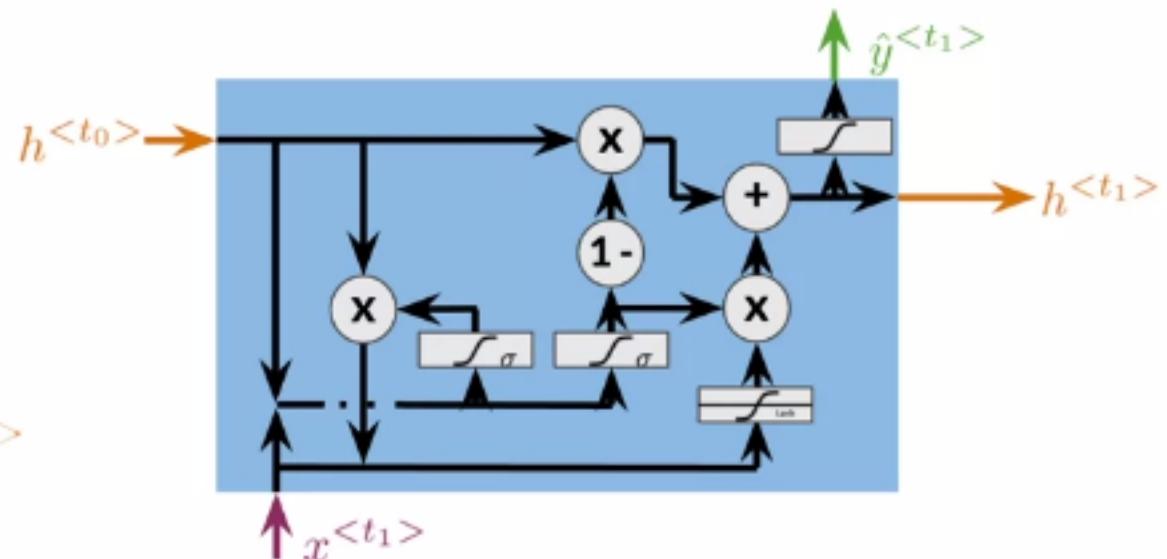
$$h'^{<t_1>} = \tanh(W_h[\Gamma_r * h^{<t_0>}, x^{<t_1>}] + b_h)$$

Vanilla RNN vs GRUs



$$h^{<t>} = g(W_h[h^{<t-1>}, x^{<t>}] + b_h)$$

$$\hat{y}^{<t>} = g(W_{yh}h^{<t>} + b_y)$$



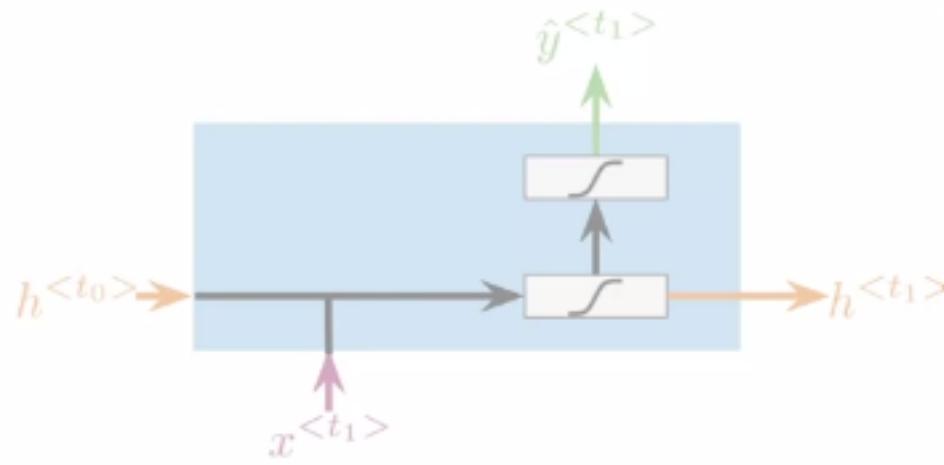
$$\Gamma_u = \sigma(W_u[h^{<t_0>}, x^{<t_1>}] + b_u)$$

$$\Gamma_r = \sigma(W_r[h^{<t_0>}, x^{<t_1>}] + b_r)$$

$$h'^{<t_1>} = \tanh (W_h[\Gamma_r * h^{<t_0>}, x^{<t_1>}] + b_h)$$

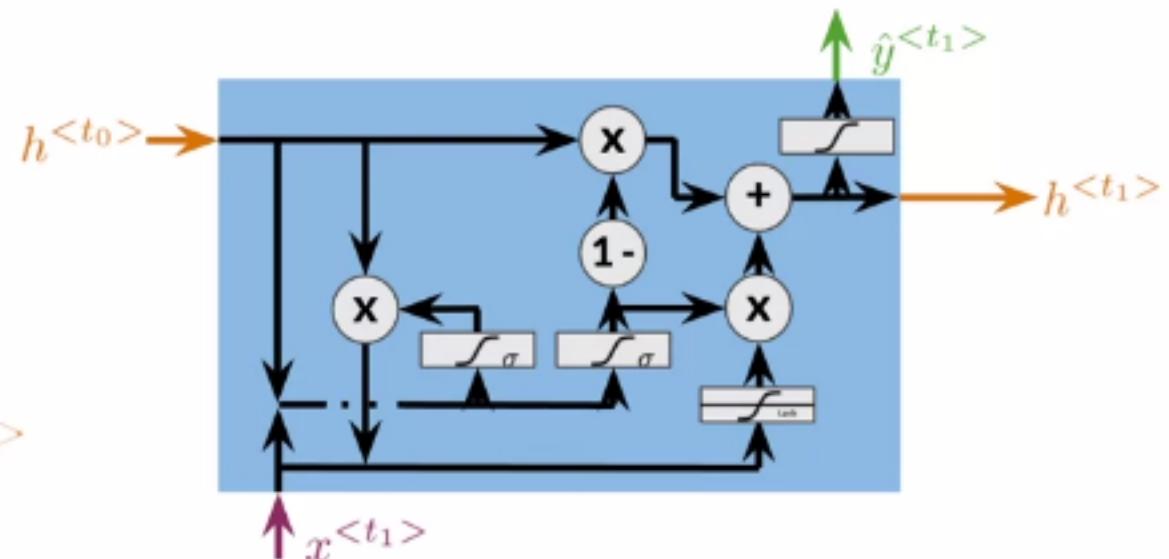
$$h^{<t_1>} = \Gamma_u * h^{<t_0>} + (1 - \Gamma_u) * h'^{<t_1>}$$

Vanilla RNN vs GRUs



$$h^{<t>} = g(W_h[h^{<t-1>}, x^{<t>}] + b_h)$$

$$\hat{y}^{<t>} = g(W_{yh}h^{<t>} + b_y)$$



$$\Gamma_u = \sigma(W_u[h^{<t_0>}, x^{<t_1>}] + b_u)$$

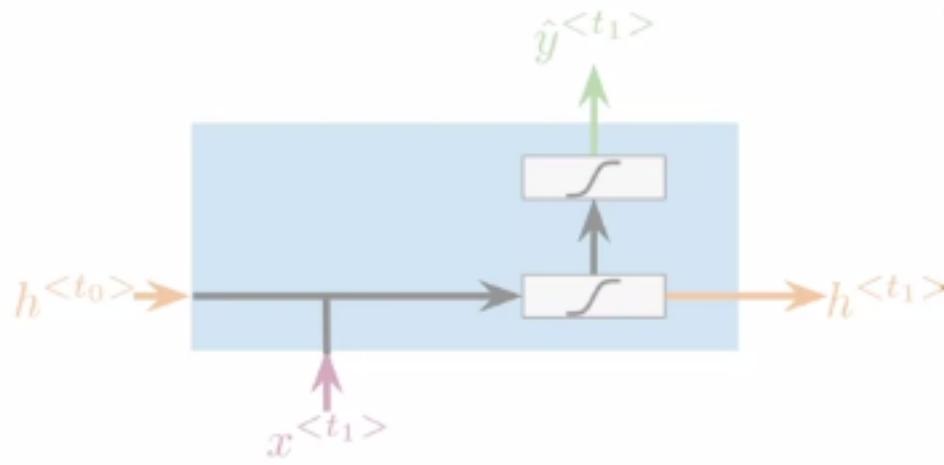
$$\Gamma_r = \sigma(W_r[h^{<t_0>}, x^{<t_1>}] + b_r)$$

$$h'^{<t_1>} = \tanh (W_h[\Gamma_r * h^{<t_0>}, x^{<t_1>}] + b_h)$$

$$h^{<t_1>} = \Gamma_u * h^{<t_0>} + (1 - \Gamma_u) * h'^{<t_1>}$$

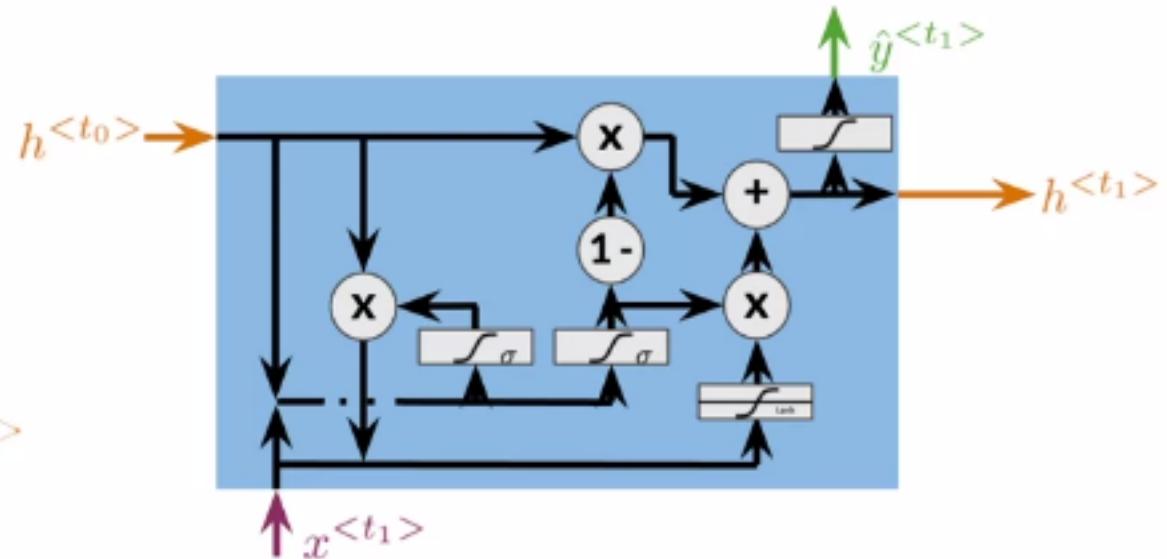
$$\hat{y}^{<t_1>} = g(W_yh^{<t_1>} + b_y)$$

Vanilla RNN vs GRUs



$$h^{<t>} = g(W_h[h^{<t-1>}, x^{<t>}] + b_h)$$

$$\hat{y}^{<t>} = g(W_y h^{<t>} + b_y)$$



$$\Gamma_u = \sigma(W_u[h^{<t_0>}, x^{<t_1>}] + b_u)$$

$$\Gamma_r = \sigma(W_r[h^{<t_0>}, x^{<t_1>}] + b_r)$$

$$h'^{<t_1>} = \tanh(W_h[\Gamma_r * h^{<t_0>}, x^{<t_1>}] + b_h)$$

$$h^{<t_1>} = \Gamma_u * h^{<t_0>} + (1 - \Gamma_u) * h'^{<t_1>}$$

$$\hat{y}^{<t_1>} = g(W_y h^{<t_1>} + b_y)$$

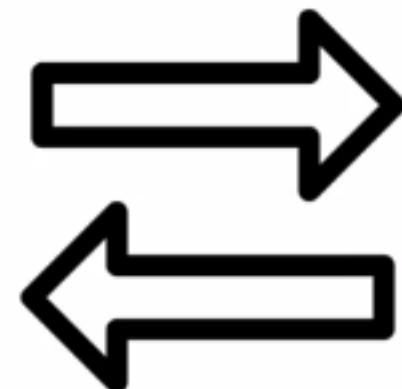
Summary

- GRUs “decide” how to update the hidden state
- GRUs help preserve important information



Outline

- How bidirectional RNNs propagate information
- Forward propagation in deep RNNs



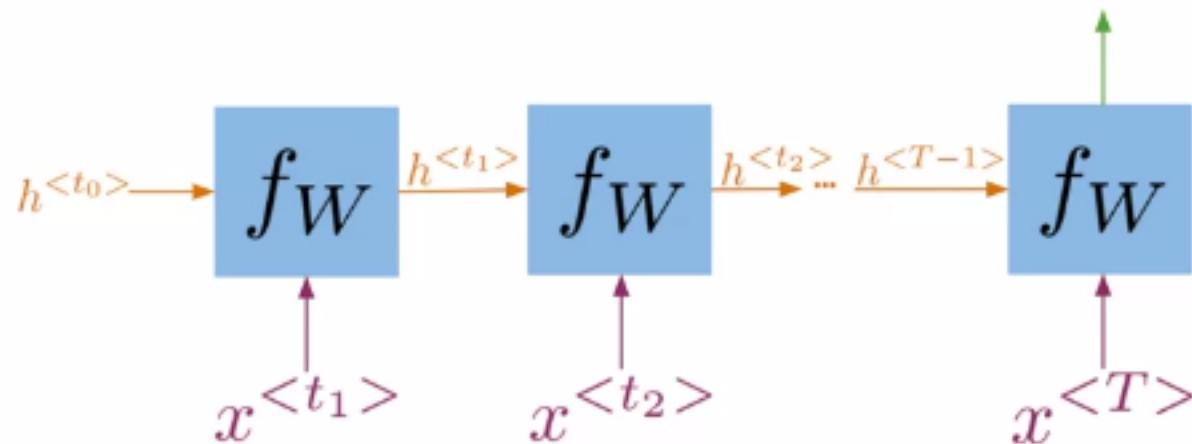
Bi-directional RNNs

Bi-directional RNNs

I was trying really hard to get a hold of _____. **Louise**, finally answered when I was about to give up.

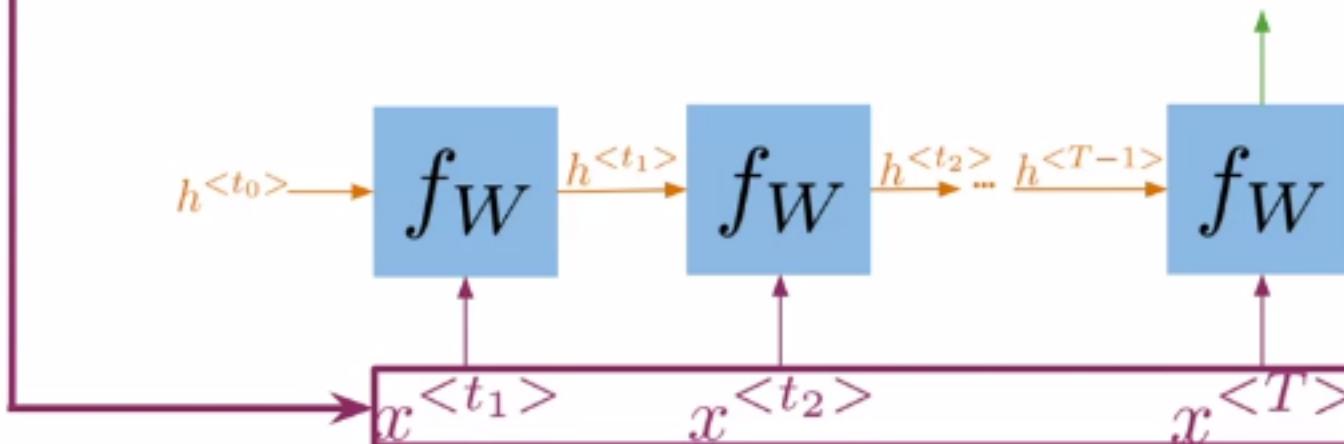
Bi-directional RNNs

I was trying really hard to get a hold of _____. **Louise**, finally answered when I was about to give up.



Bi-directional RNNs

I was trying really hard to get a hold of _____.
Louise, finally
answered when I was about to give up.



Bi-directional RNNs

Bi-directional RNNs

$x^{}$

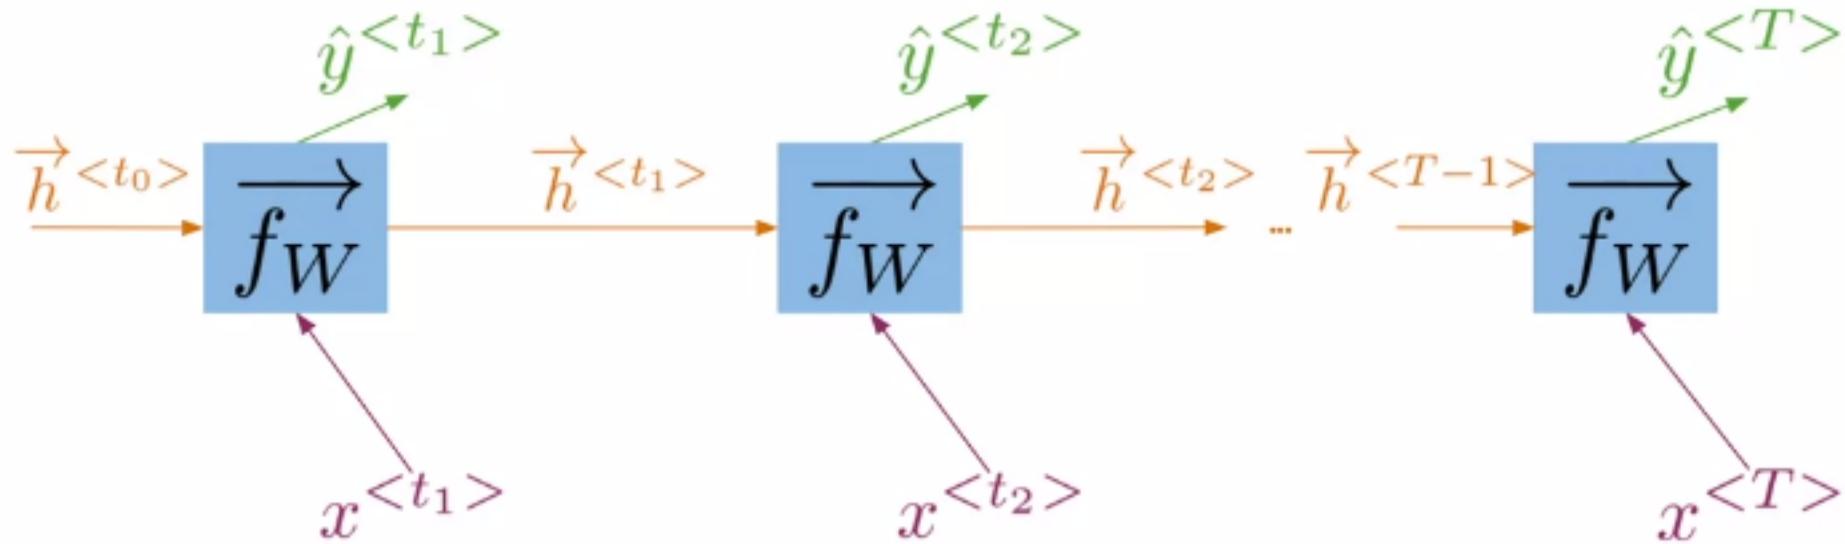
$x^{}$

$x^{}$

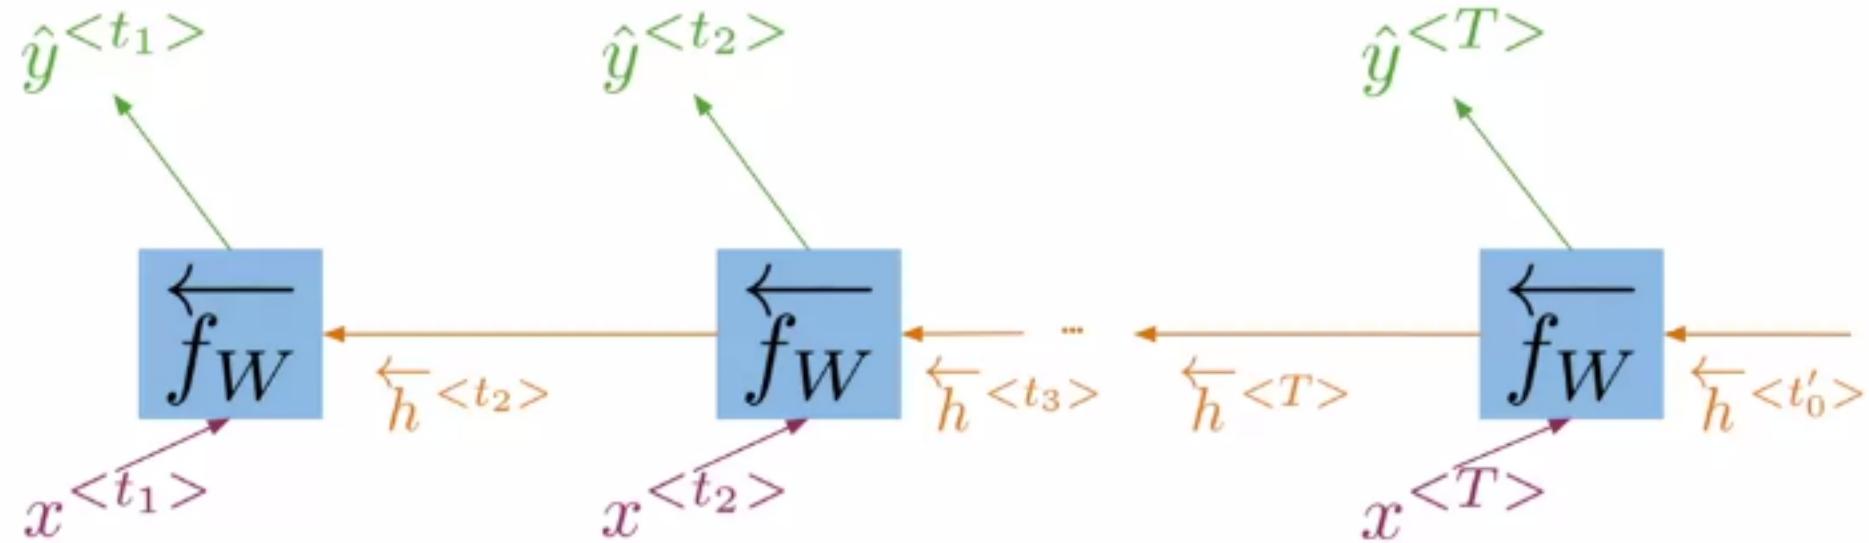
Bi-directional RNNs

 $\hat{y}^{}$ $\hat{y}^{}$ $\hat{y}^{}$ $x^{}$ $x^{}$ $x^{}$

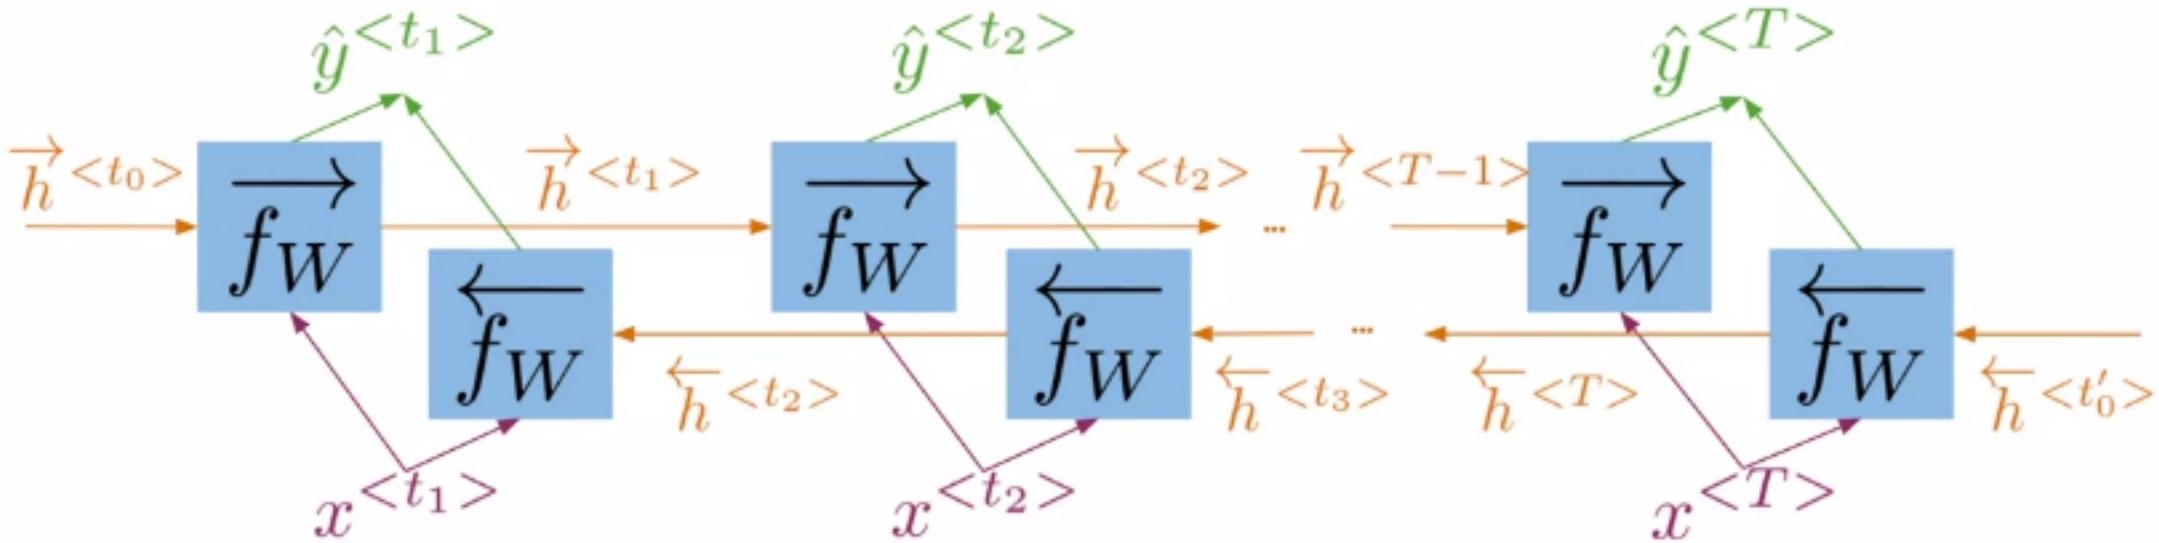
Bi-directional RNNs



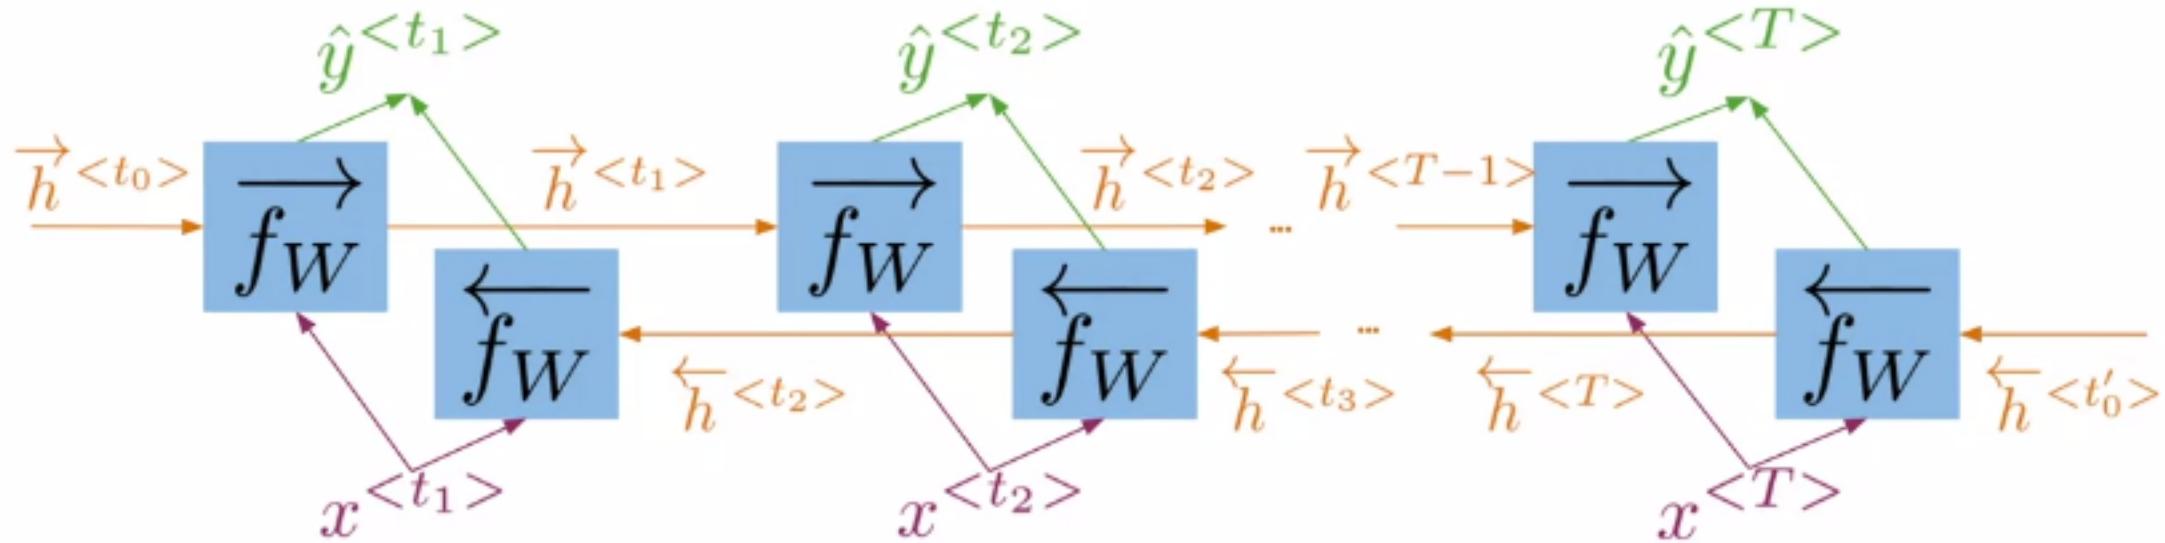
Bi-directional RNNs



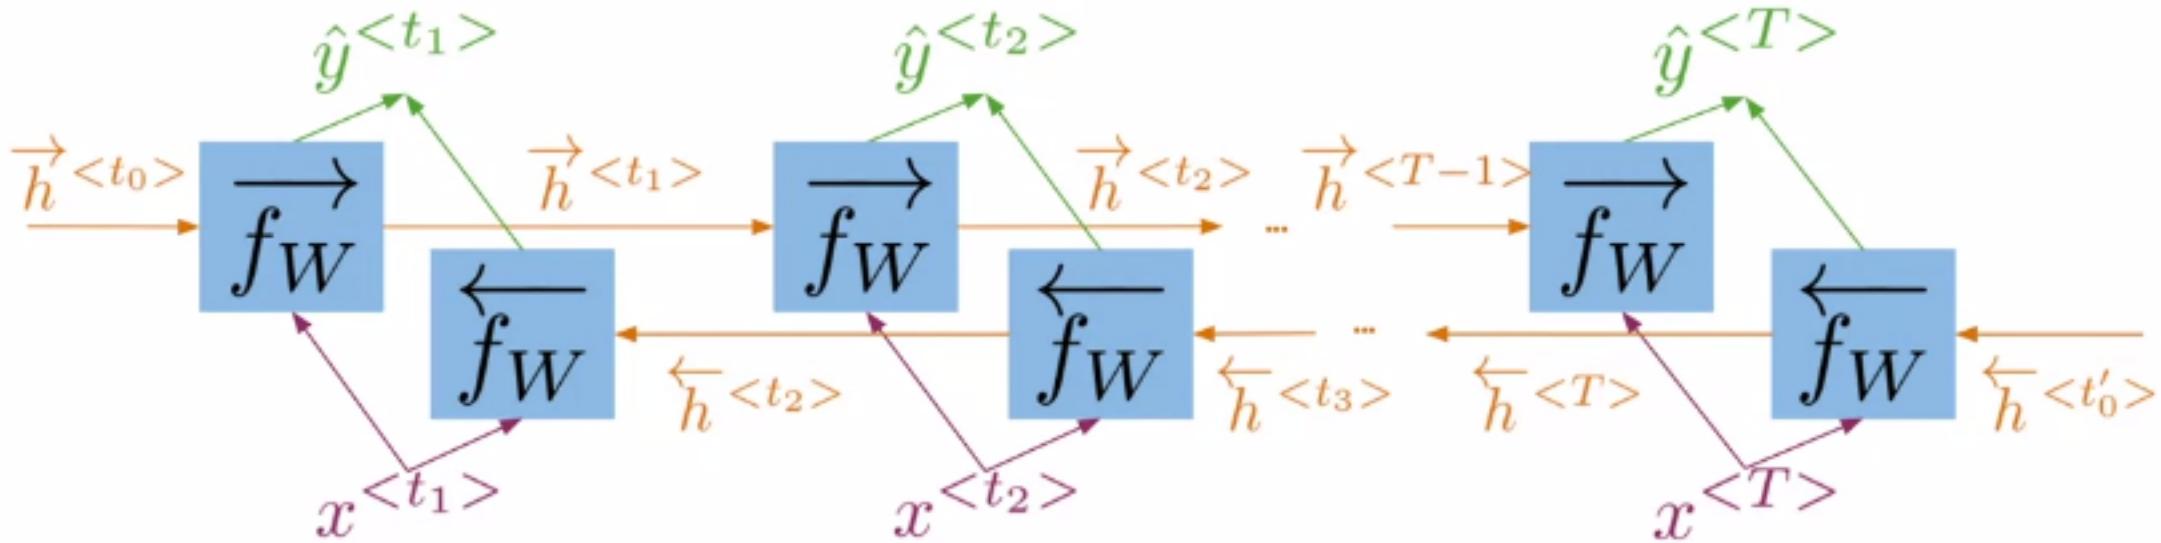
Bi-directional RNNs



Bi-directional RNNs

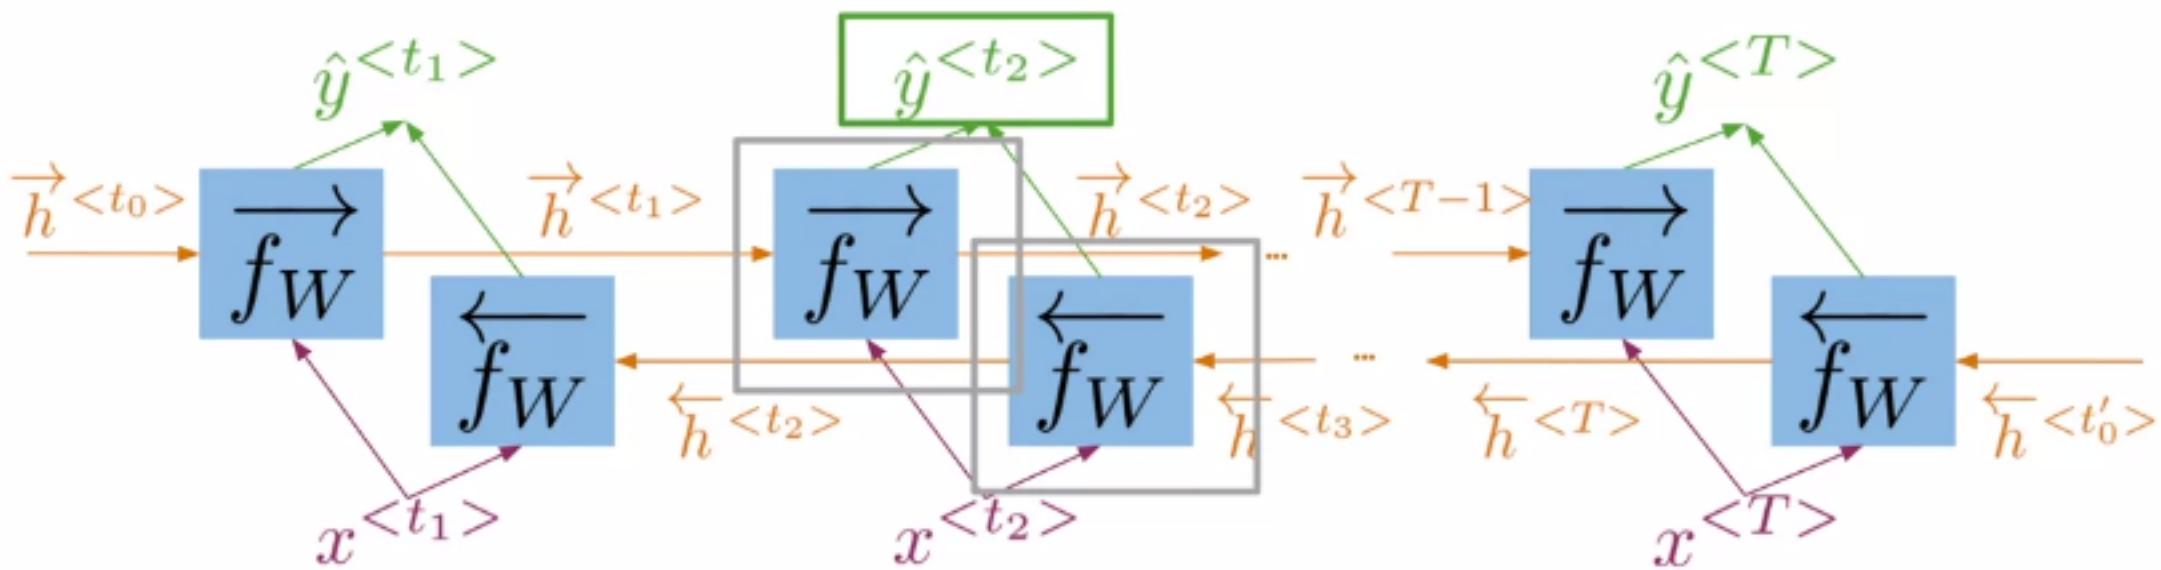


Bi-directional RNNs



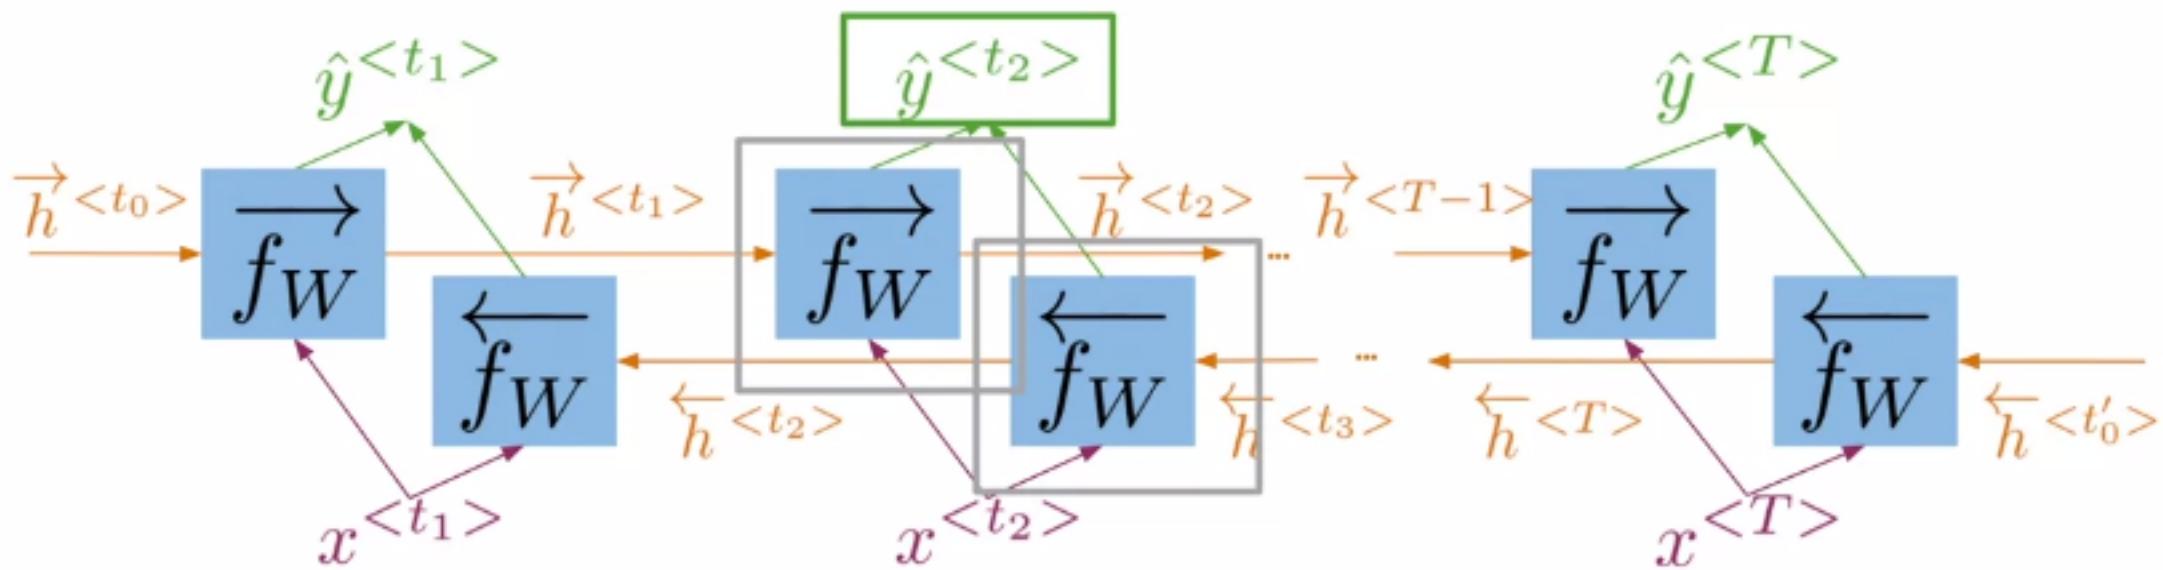
Information flows from the past and from the future
independently

Bi-directional RNNs



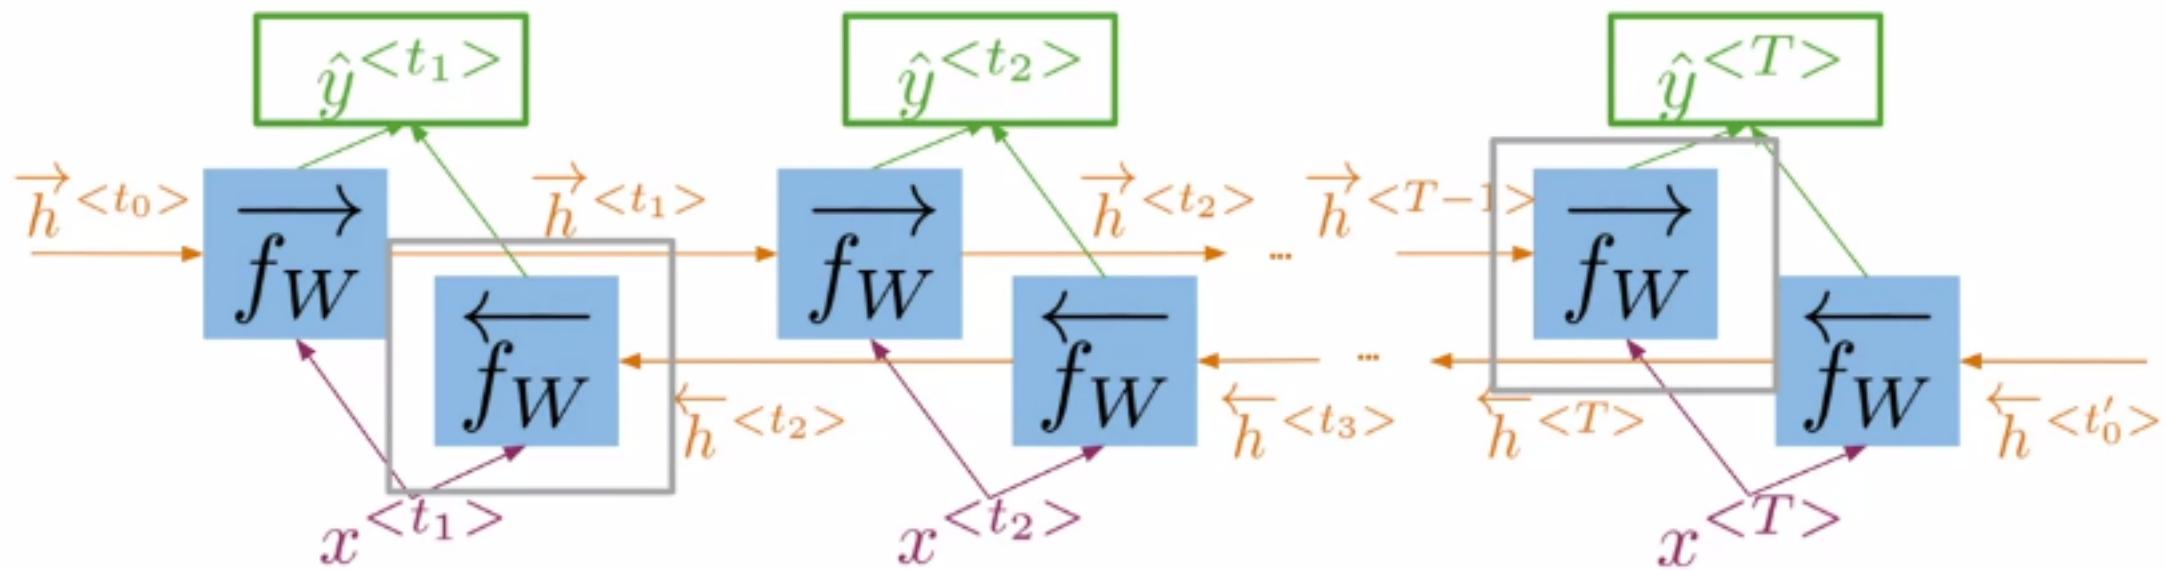
Information flows from the past and from the future
independently

Bi-directional RNNs



$$\hat{y}^{<t>} = g(W_y[\vec{h}^{<t>}, \overleftarrow{h}^{<t>}] + b_y)$$

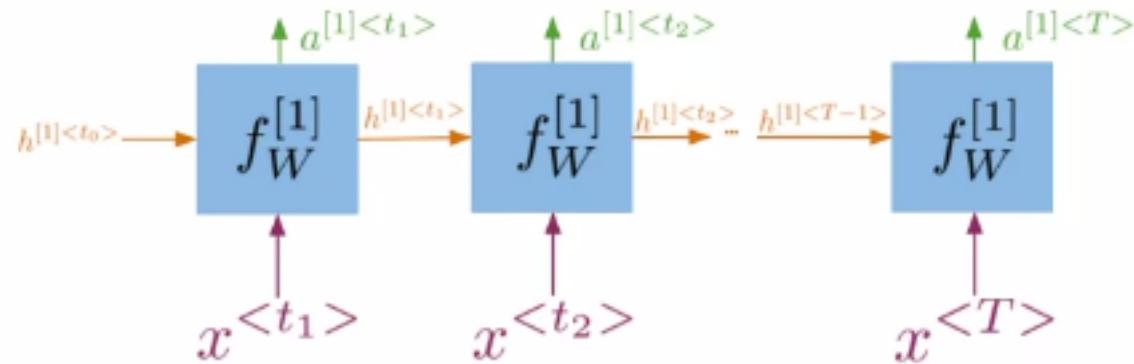
Bi-directional RNNs



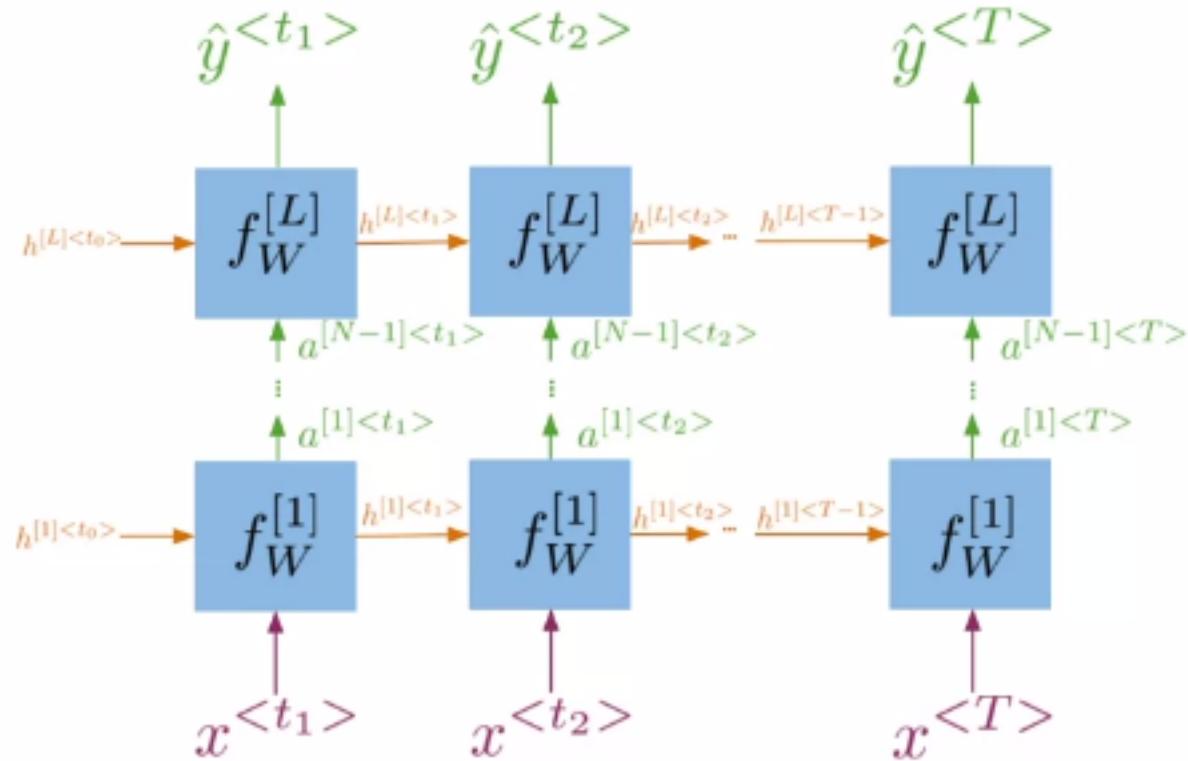
$$\hat{y}^{<t>} = g(W_y[\vec{h}^{<t>}, \overleftarrow{h}^{<t>}] + b_y)$$

Deep RNNs

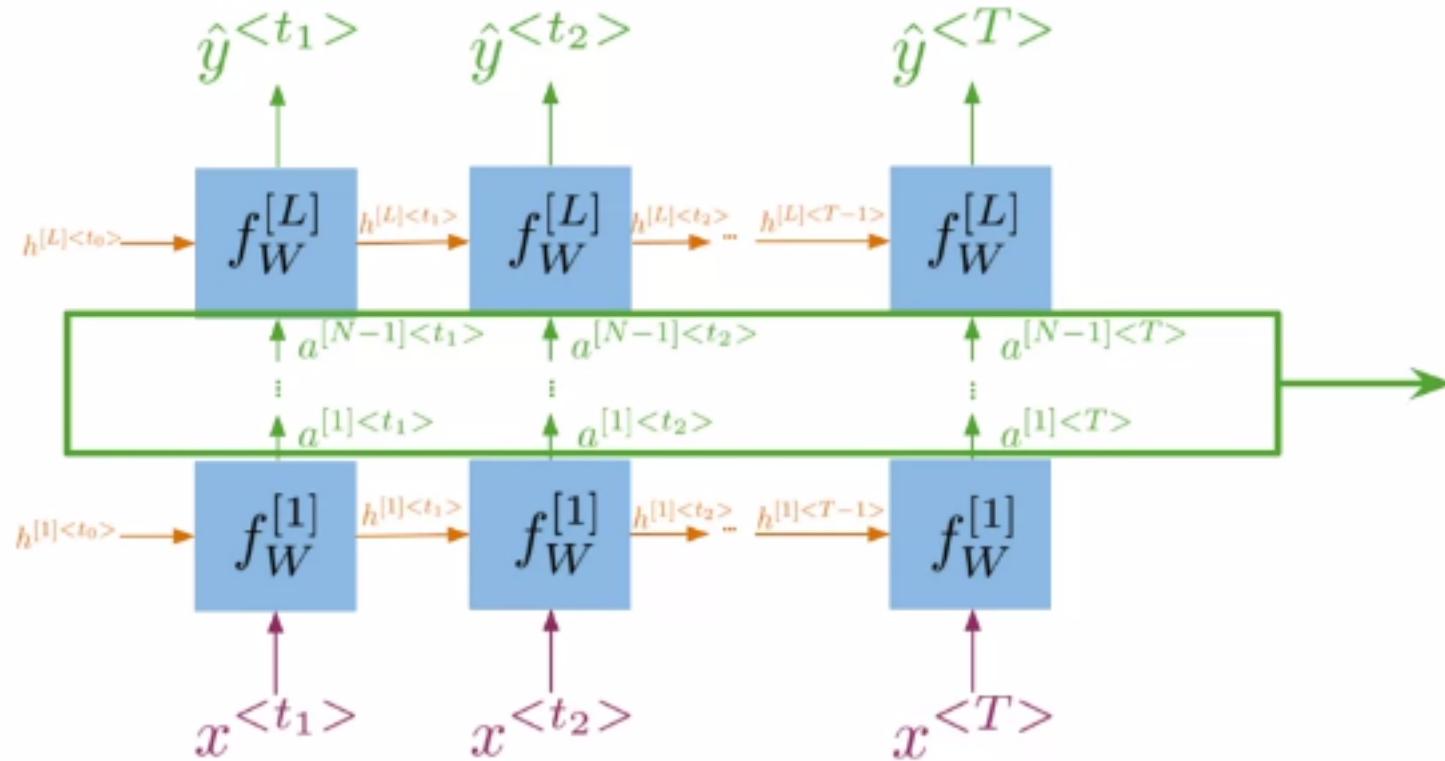
Deep RNNs



Deep RNNs

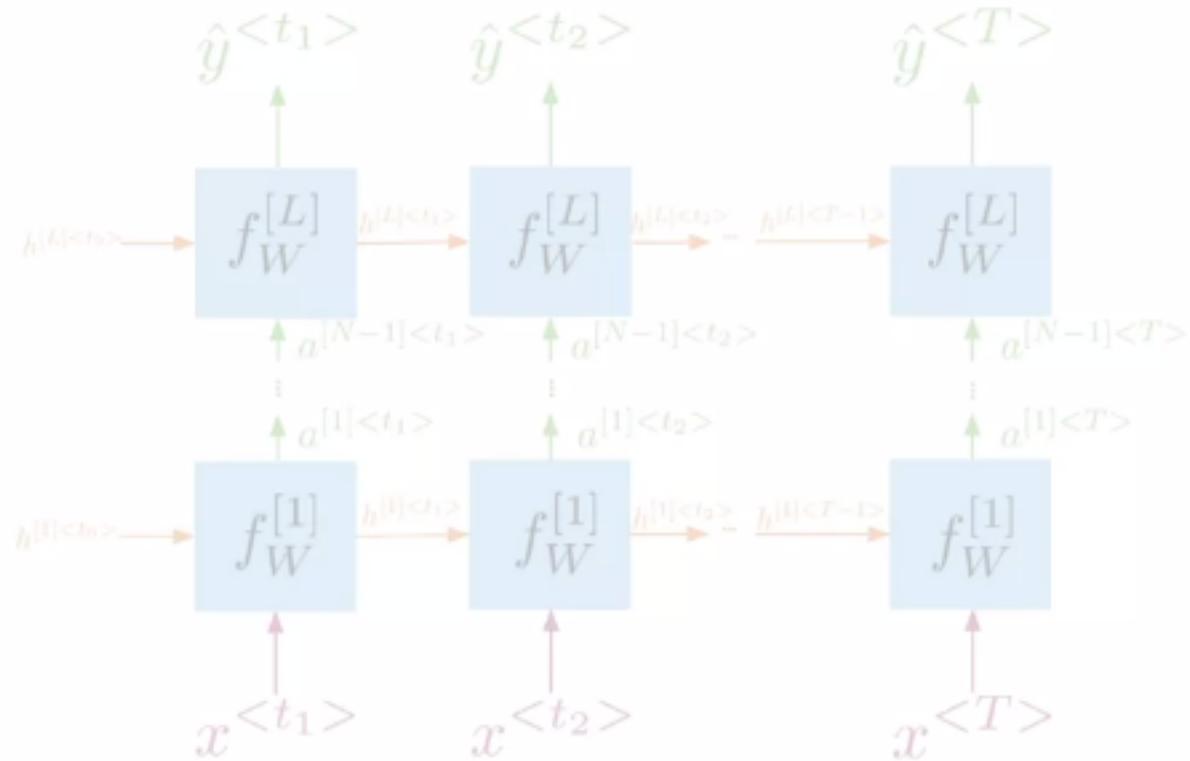


Deep RNNs



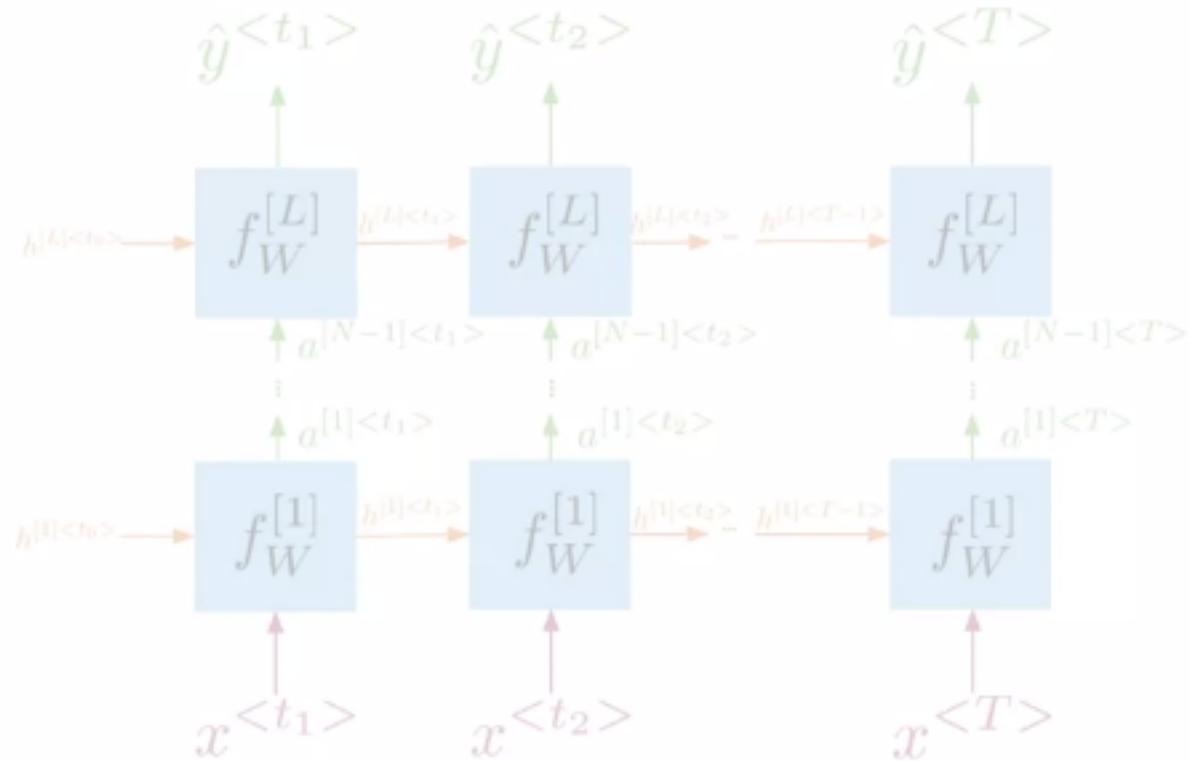
Intermediate
layers and
activations

Deep RNNs



$$h^{[l]<t>} = f^{[l]}(W_h^{[l]}[h^{[l]<t-1>}, a^{[l-1]<t>}] + b_h^{[l]})$$
$$a^{[l]<t>} = f^{[l]}(W_a^{[l]} h^{[l]<t>} + b_a^{[l]})$$

Deep RNNs



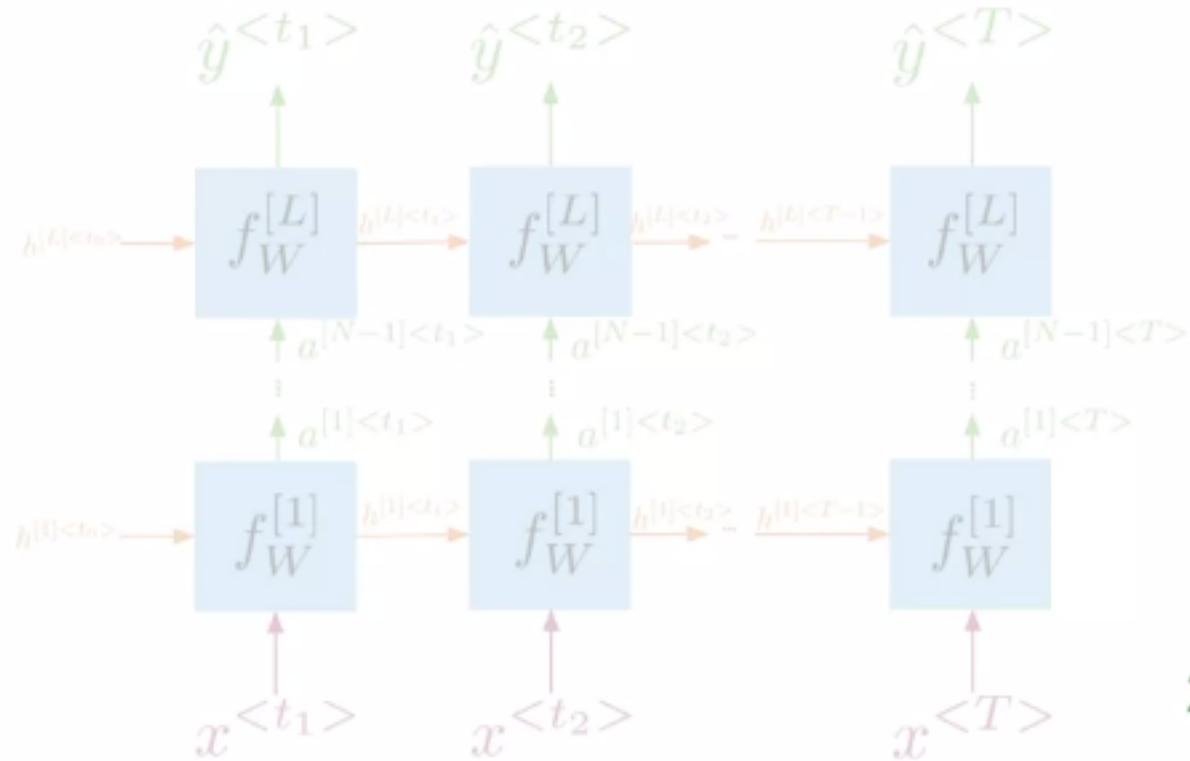
$$h^{[l]<t>} = f^{[l]}(W_h^{[l]}[h^{[l]<t-1>}, a^{[l-1]<t>}] + b_h^{[l]})$$

$$a^{[l]<t>} = f^{[l]}(W_a^{[l]} h^{[l]<t>} + b_a^{[l]})$$



1. Get hidden states for current layer

Deep RNNs



$$h^{[l]} < t > = f^{[l]}(W_h^{[l]} h^{[l]} < t-1 >, a^{[l-1]} < t > + b_h^{[l]})$$
$$a^{[l]} < t > = f^{[l]}(W_a^{[l]} h^{[l]} < t > + b_a^{[l]})$$



1. Get hidden states for current layer
2. Pass the activations to the next layer

Summary

- In bidirectional RNNs, the outputs take information from the past and the future
- Deep RNNs have more than one layer, which helps in complex tasks

